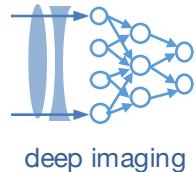


Lecture 5: A gentle introduction to optimization

Machine Learning and Imaging

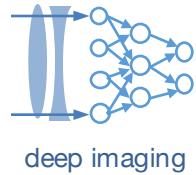
BME 590L
Roarke Horstmeyer



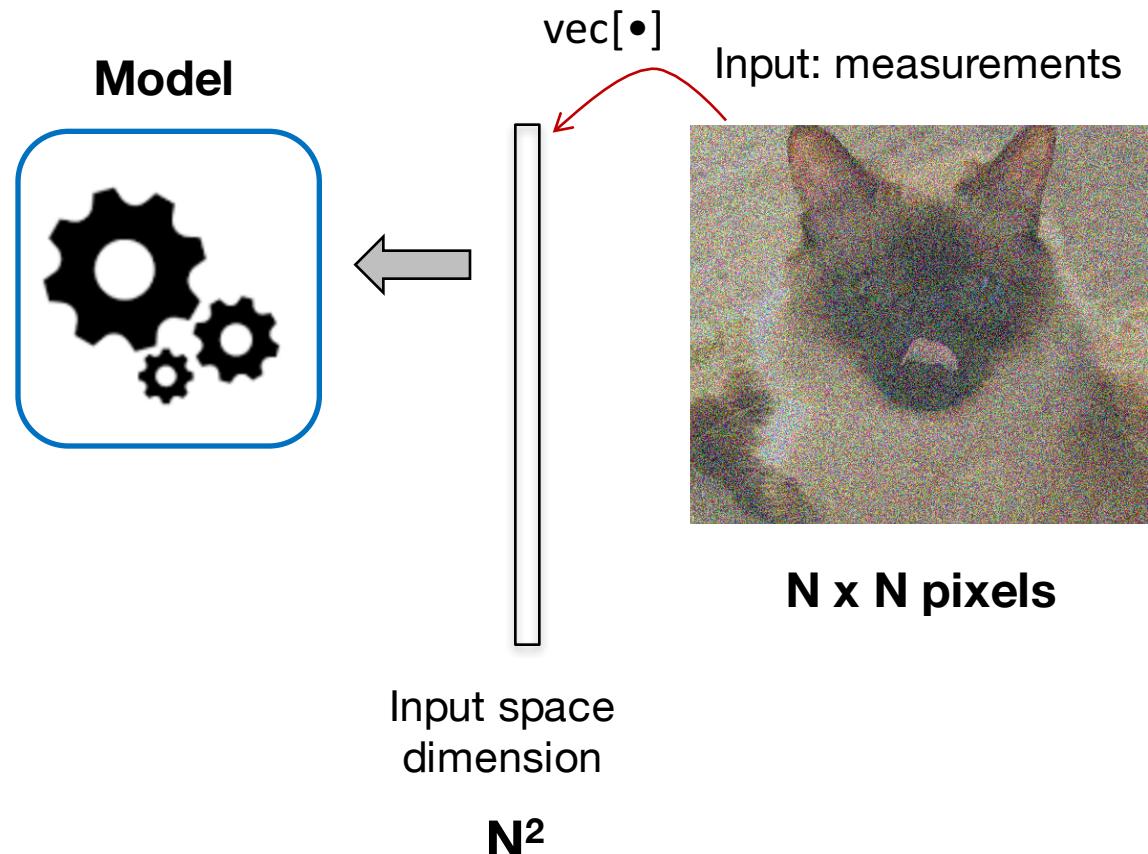
Mathematical Optimization: "Selection of a **best element** (with regard to some **criterion**) from **a set of available alternatives**"

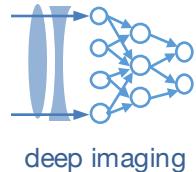
3 elements:

- 1) Your desired output (a better image, a clean signal, a classification of "cat" or "dog", etc.)
- 2) A model of what you are looking for - how you form the desired output from your measured data
- 3) A cost function, to measure how close you're getting to the answer (the cost function minimum)

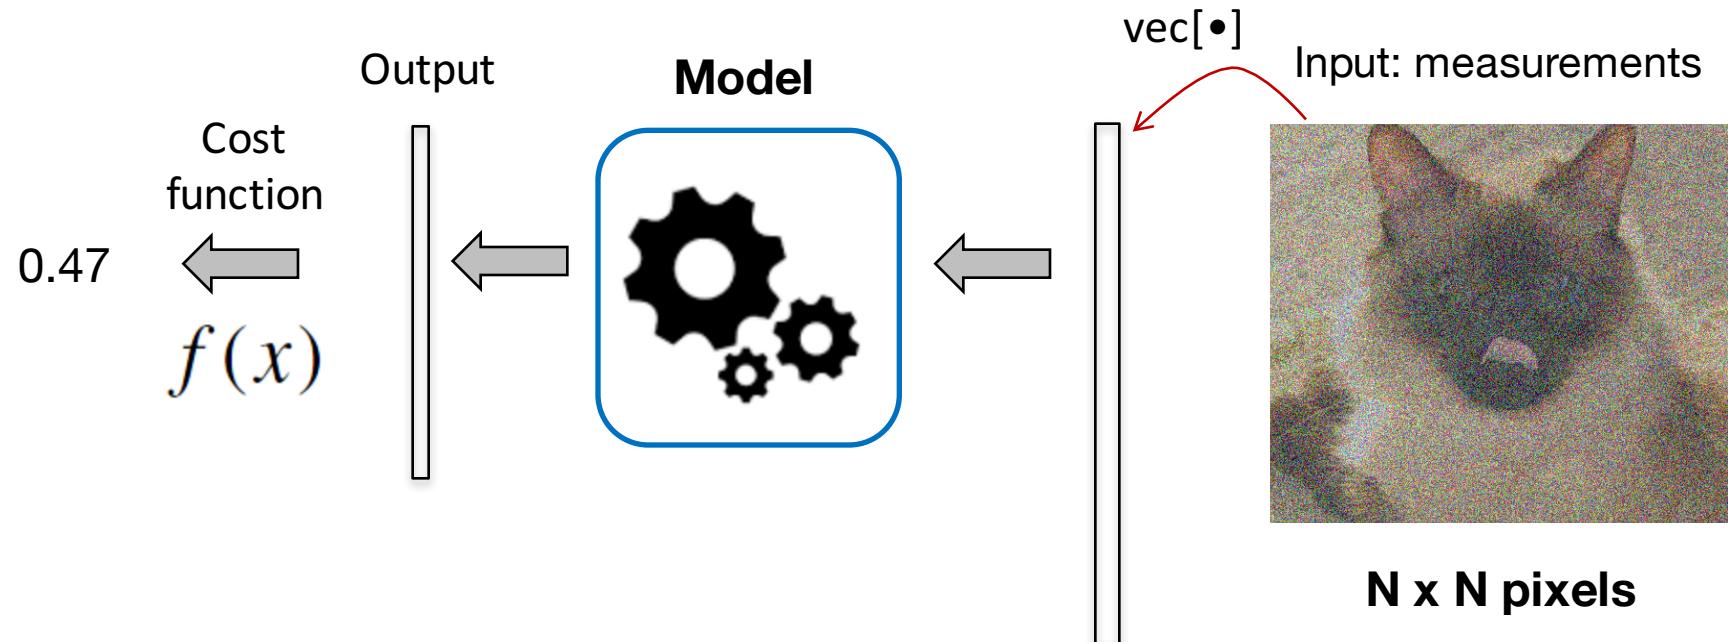


Generalized optimization pipeline





Generalized optimization pipeline



Performance measure

1 number!

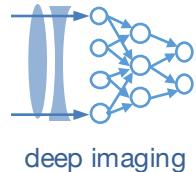
Output space dimension

M

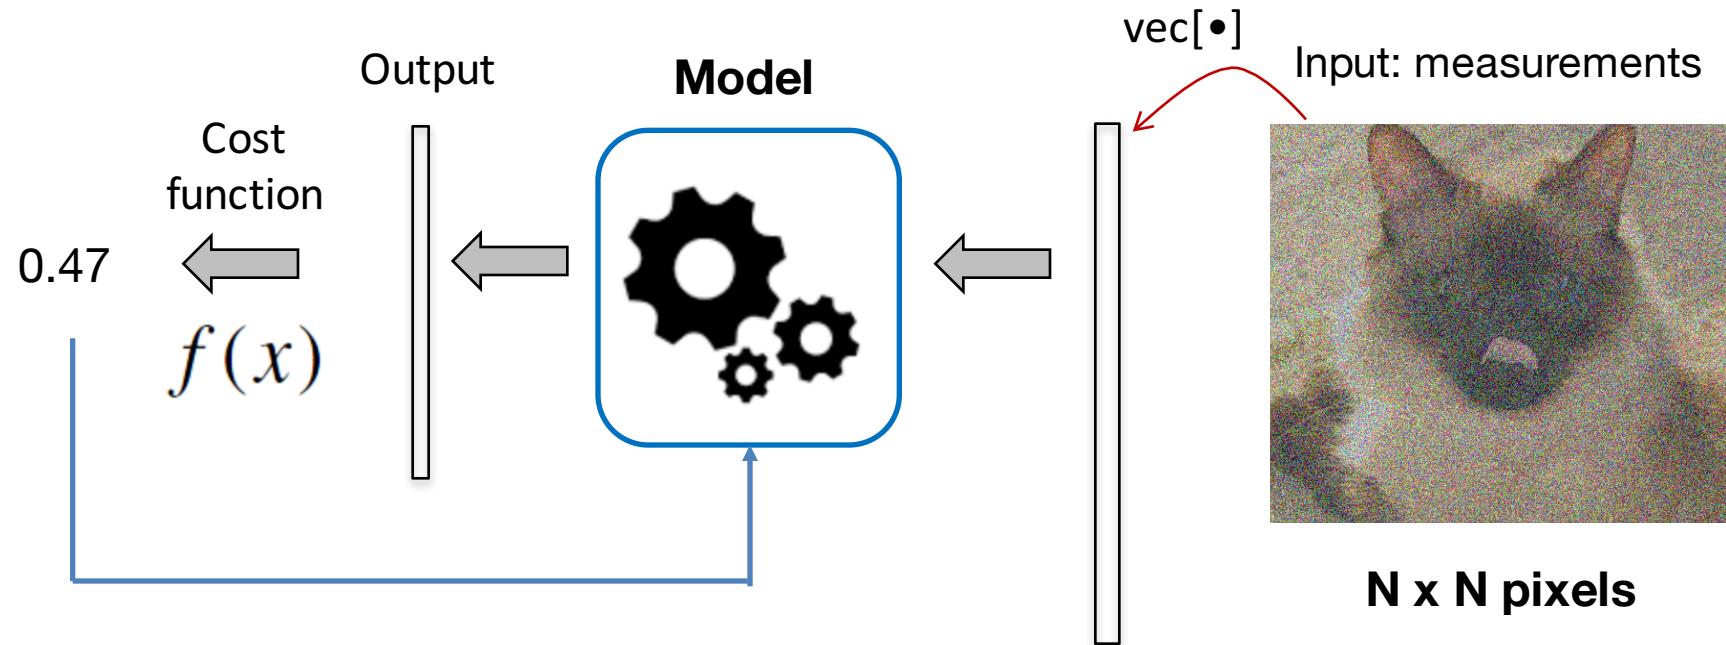
Input space dimension

N^2

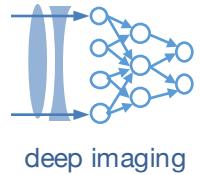
How well did we do?



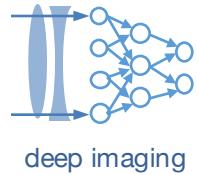
Machine learning: update model to decrease error



How well did
we do?



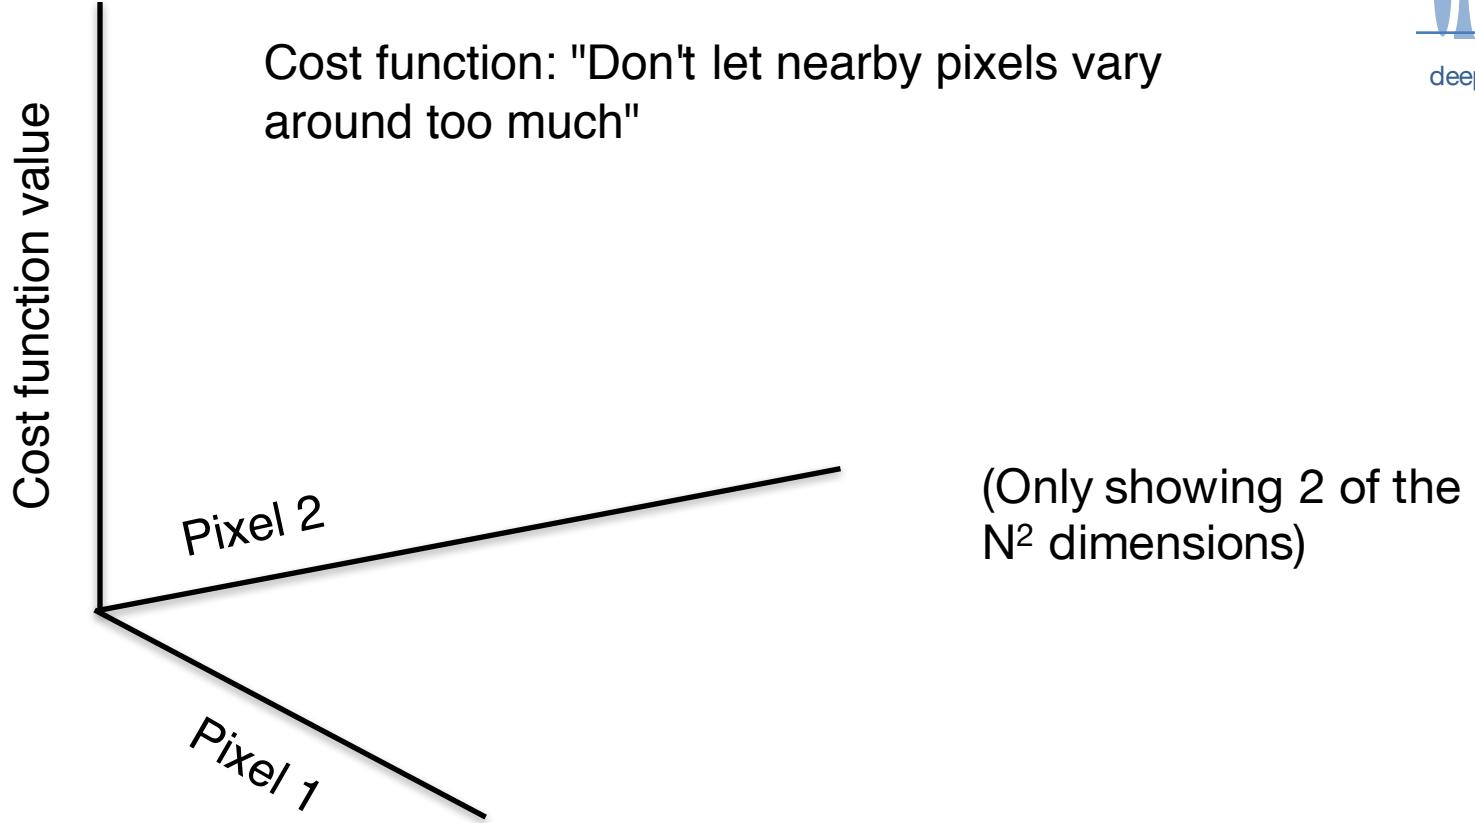
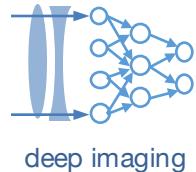
De-noising: "What is the closest image to what I detected, except without so many fluctuations"?



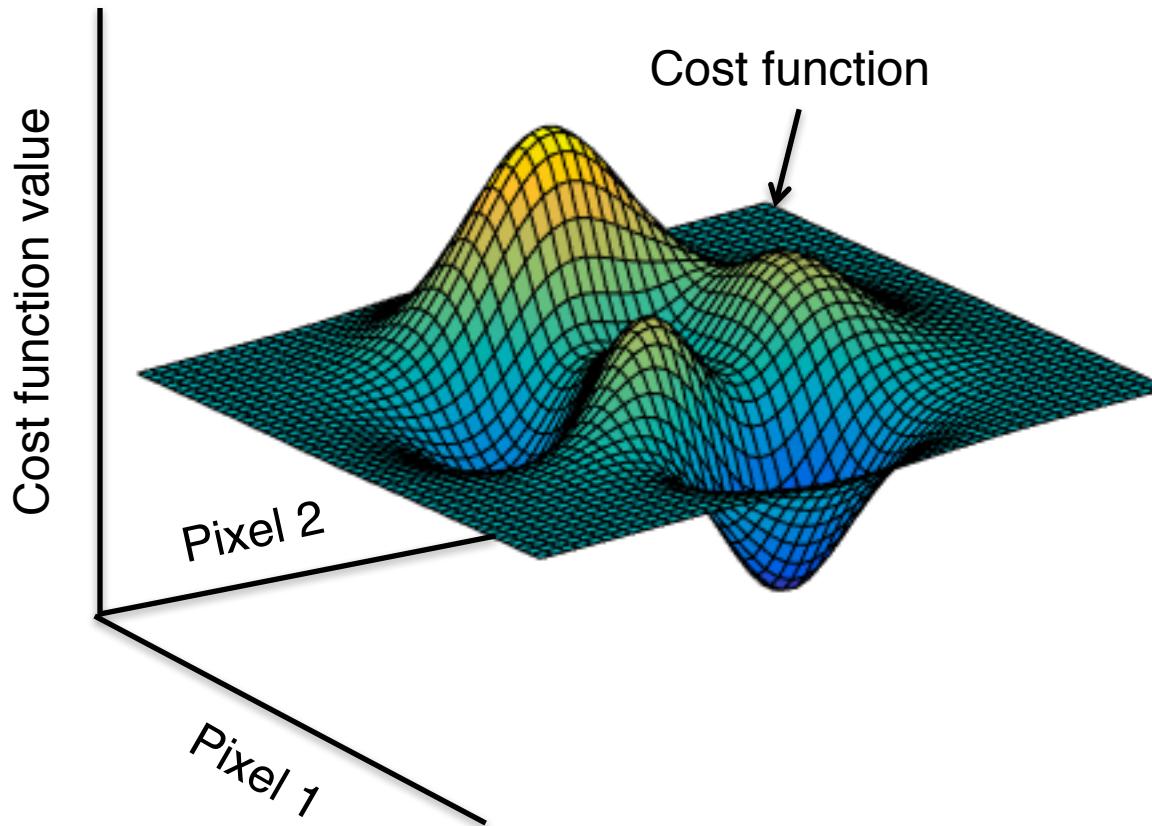
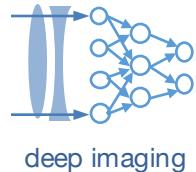
Input dimension: $N \times N$ image

Output dimension: $N \times N$ image

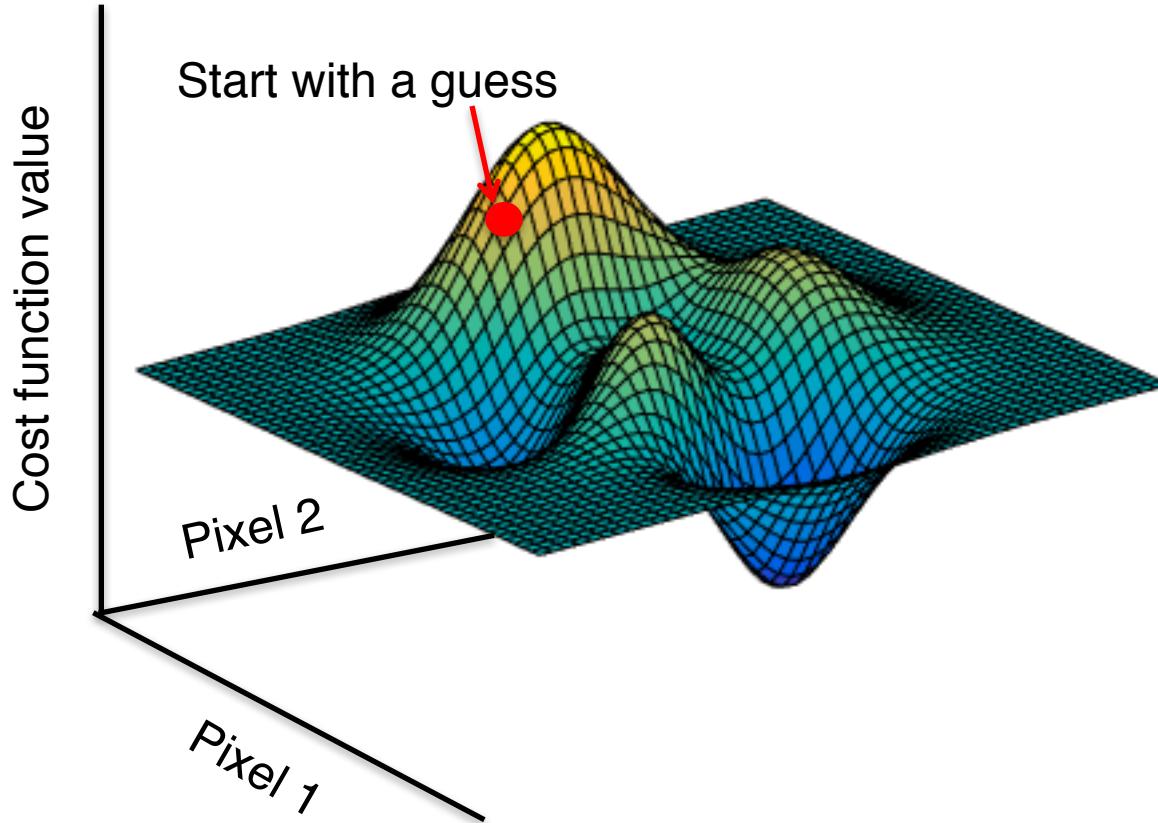
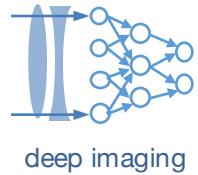
De-noising: "What is the closest image to what I detected, except without so many fluctuations"?



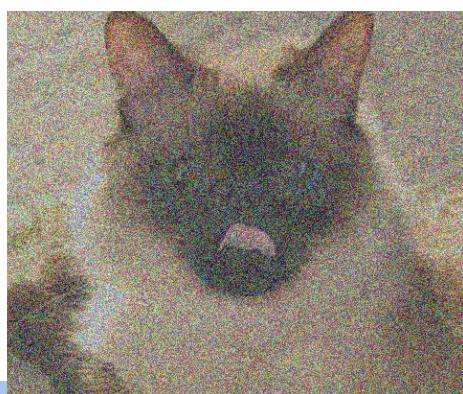
De-noising: "What is the closest image to what I detected, except without so many fluctuations"?



De-noising: "What is the closest image to what I detected, except without so many fluctuations"?

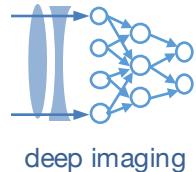


Noisy image

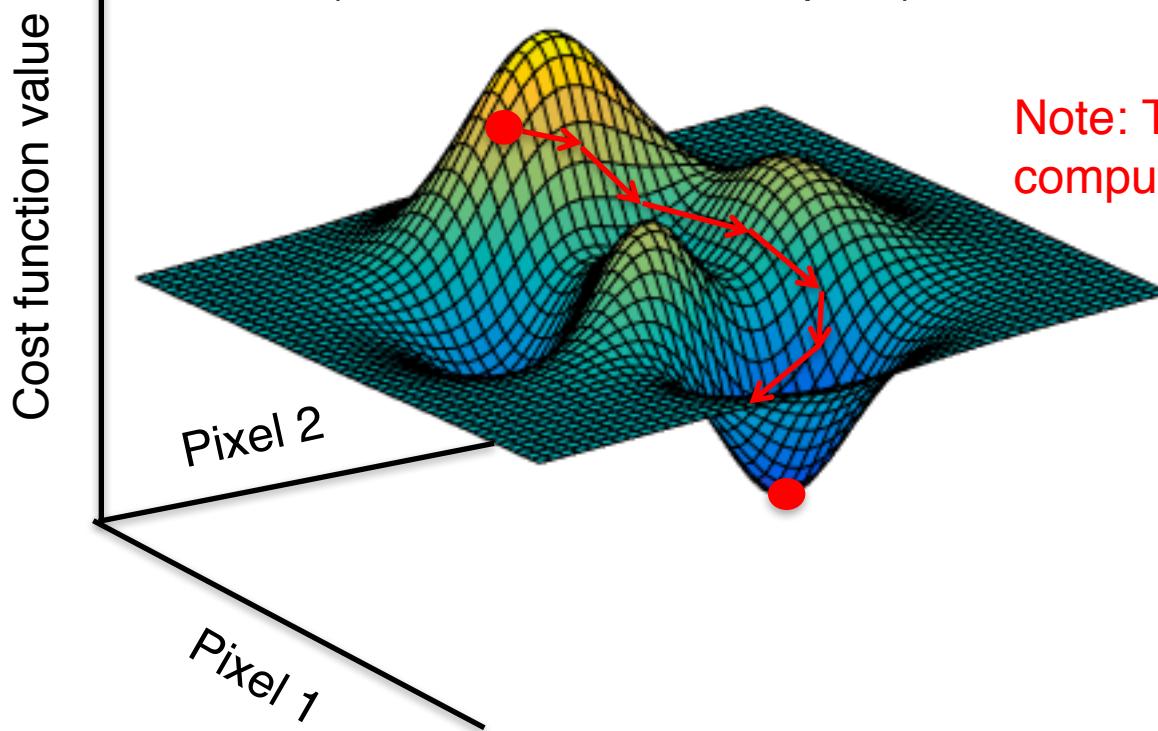


Pixel value 1
Pixel value 2

:

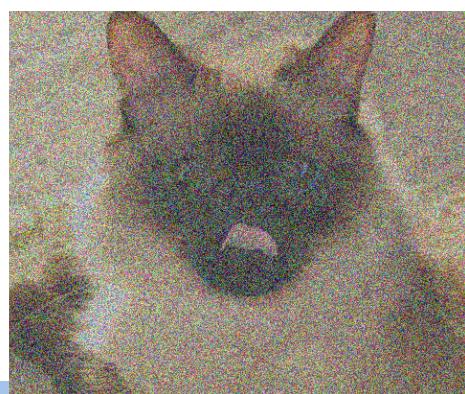


"Descend" to minimize cost function
(tweak values of each pixel)



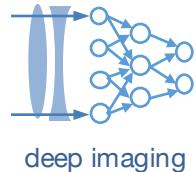
Note: This math part
computers are good at

Noisy image

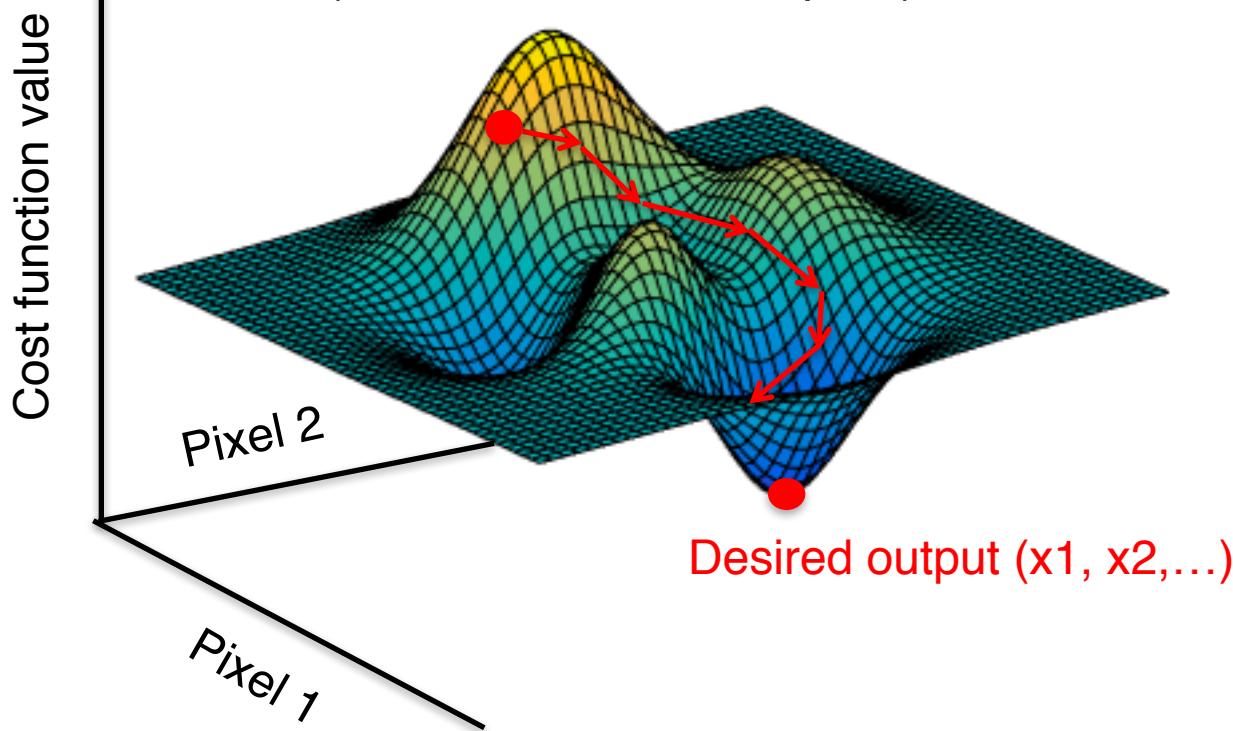


Pixel value 1
Pixel value 2

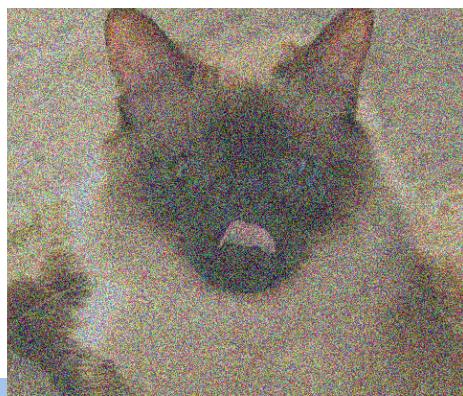
:



"Descend" to minimize cost function
(tweak values of each pixel)



Noisy image



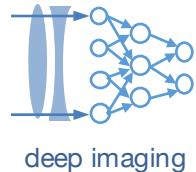
Pixel value 1
Pixel value 2
⋮



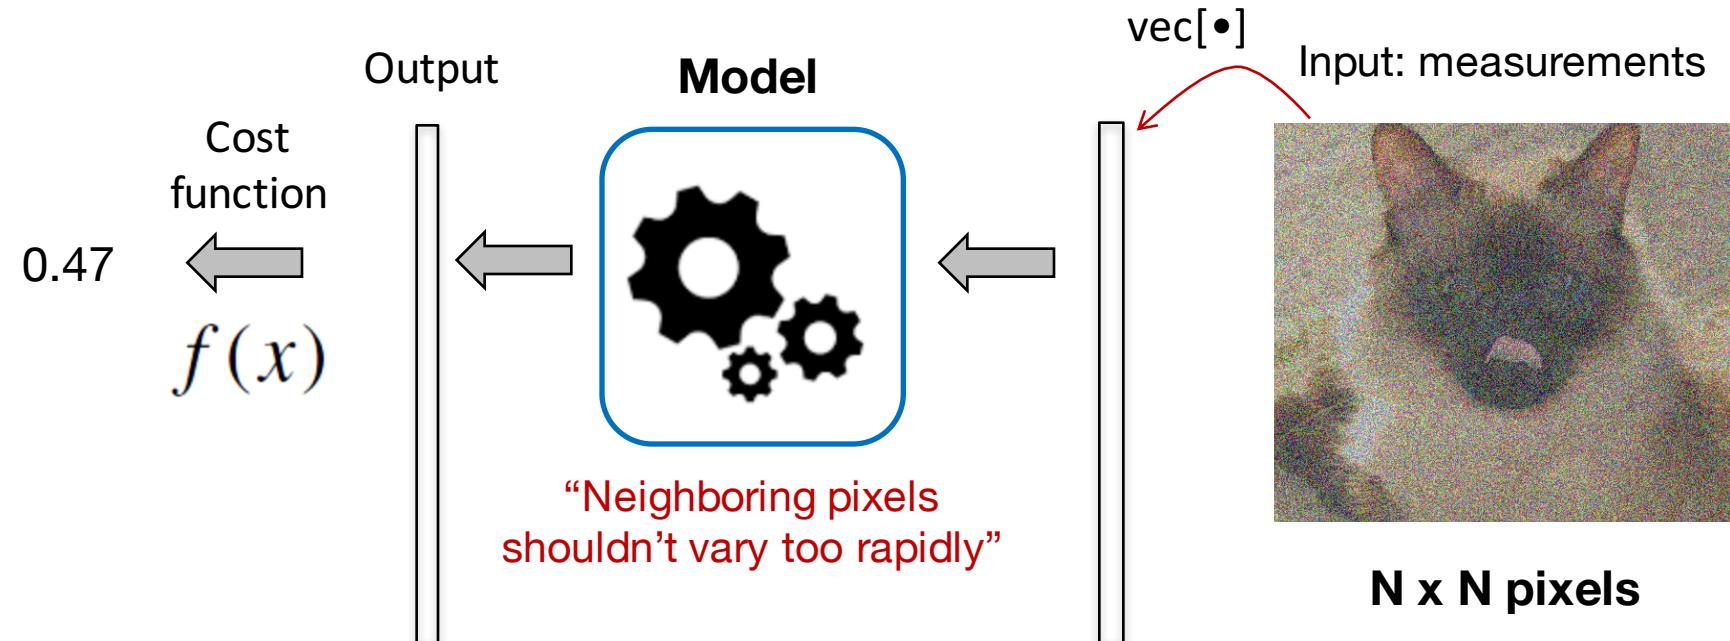
Desired output



x_1
 x_2
⋮



Optimization pipeline for denoising



Performance measure

Mean-squared error

Output space dimension N^2



Input space dimension N^2

$N \times N \text{ pixels}$

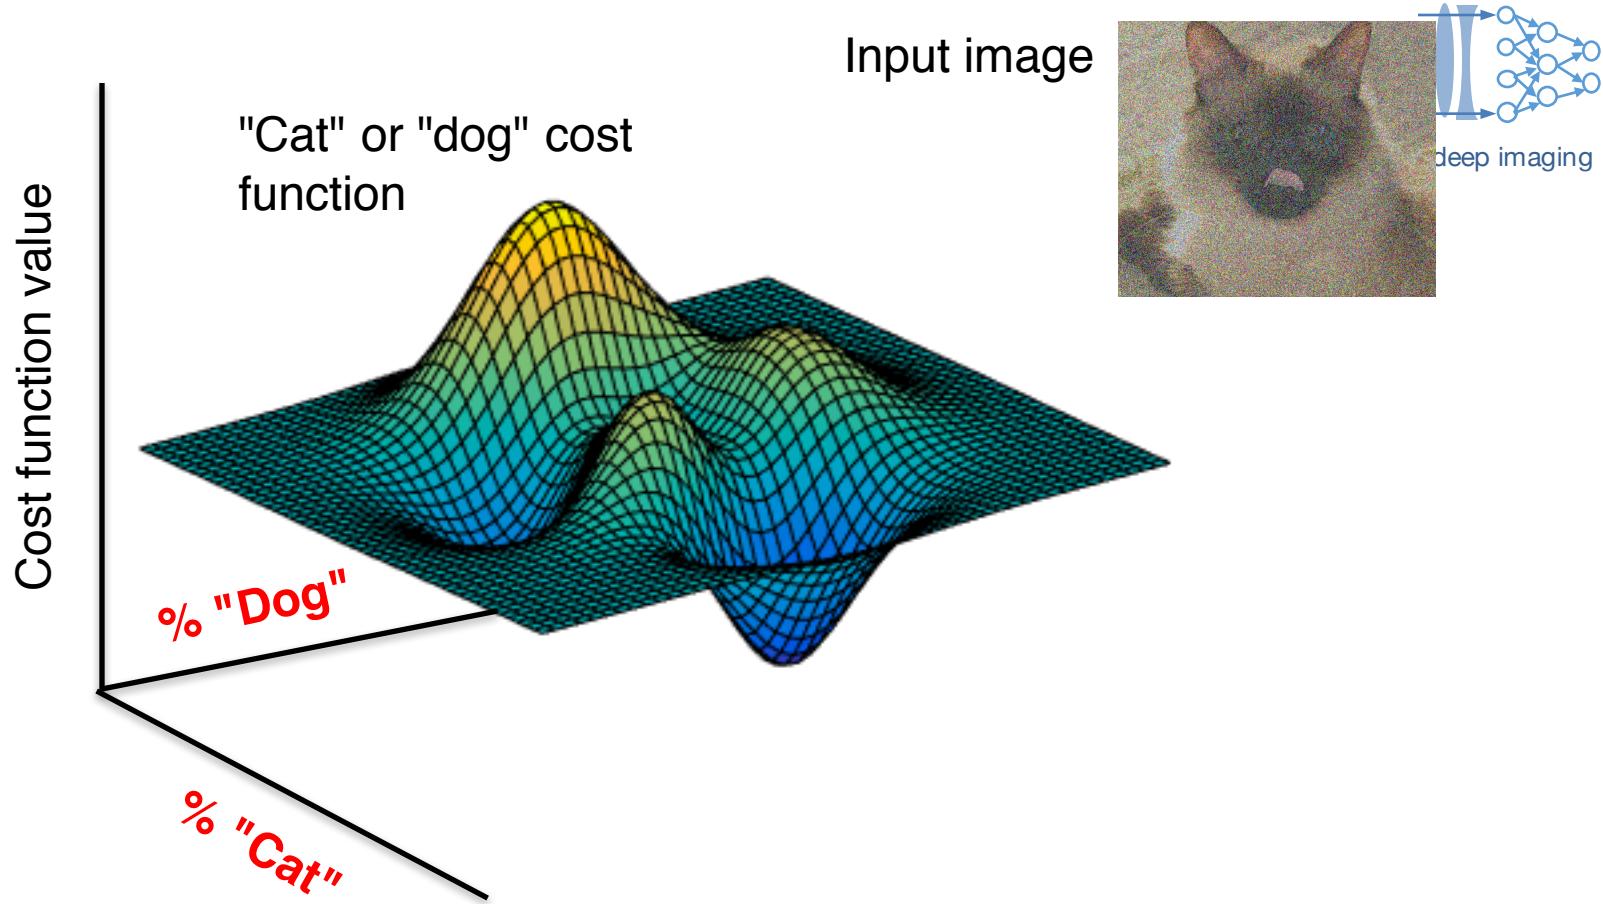
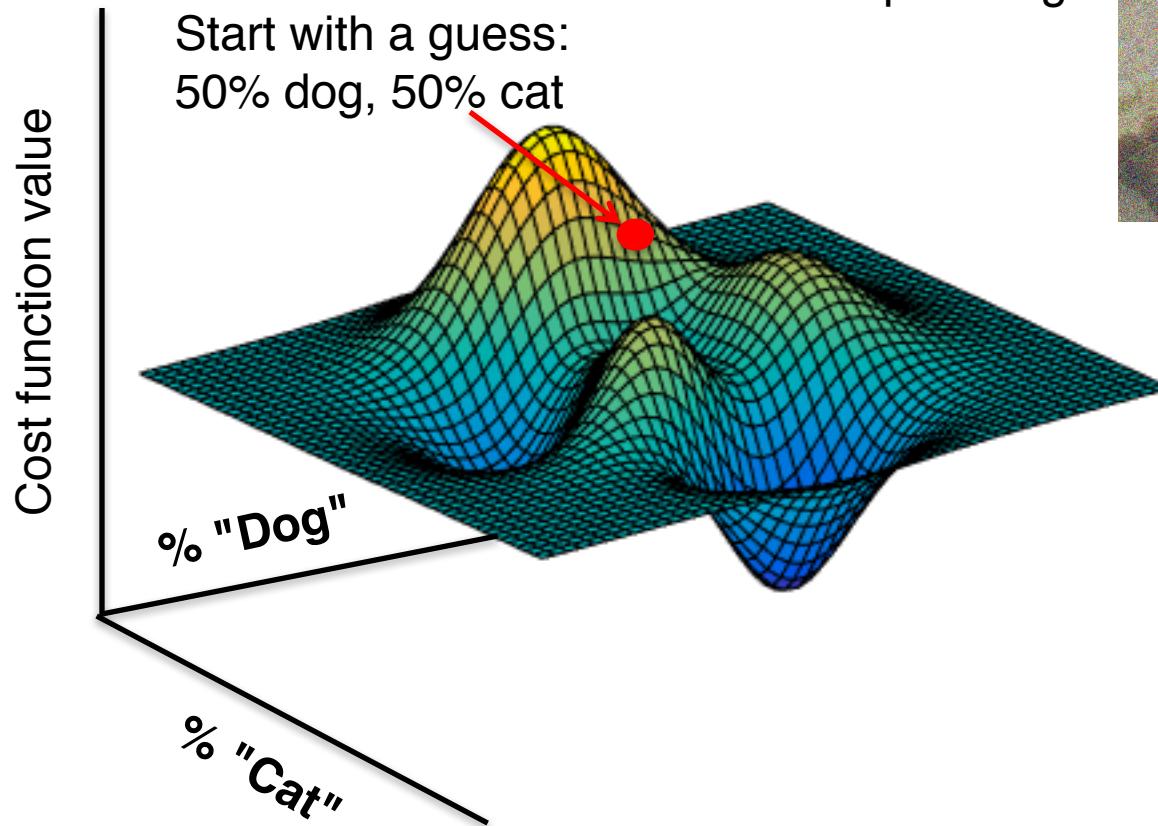
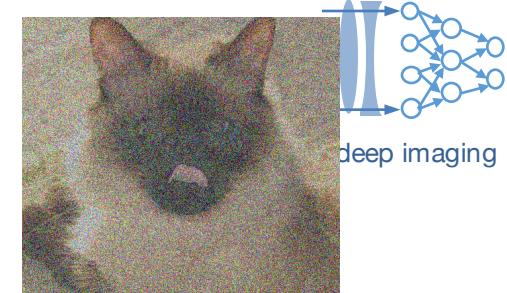
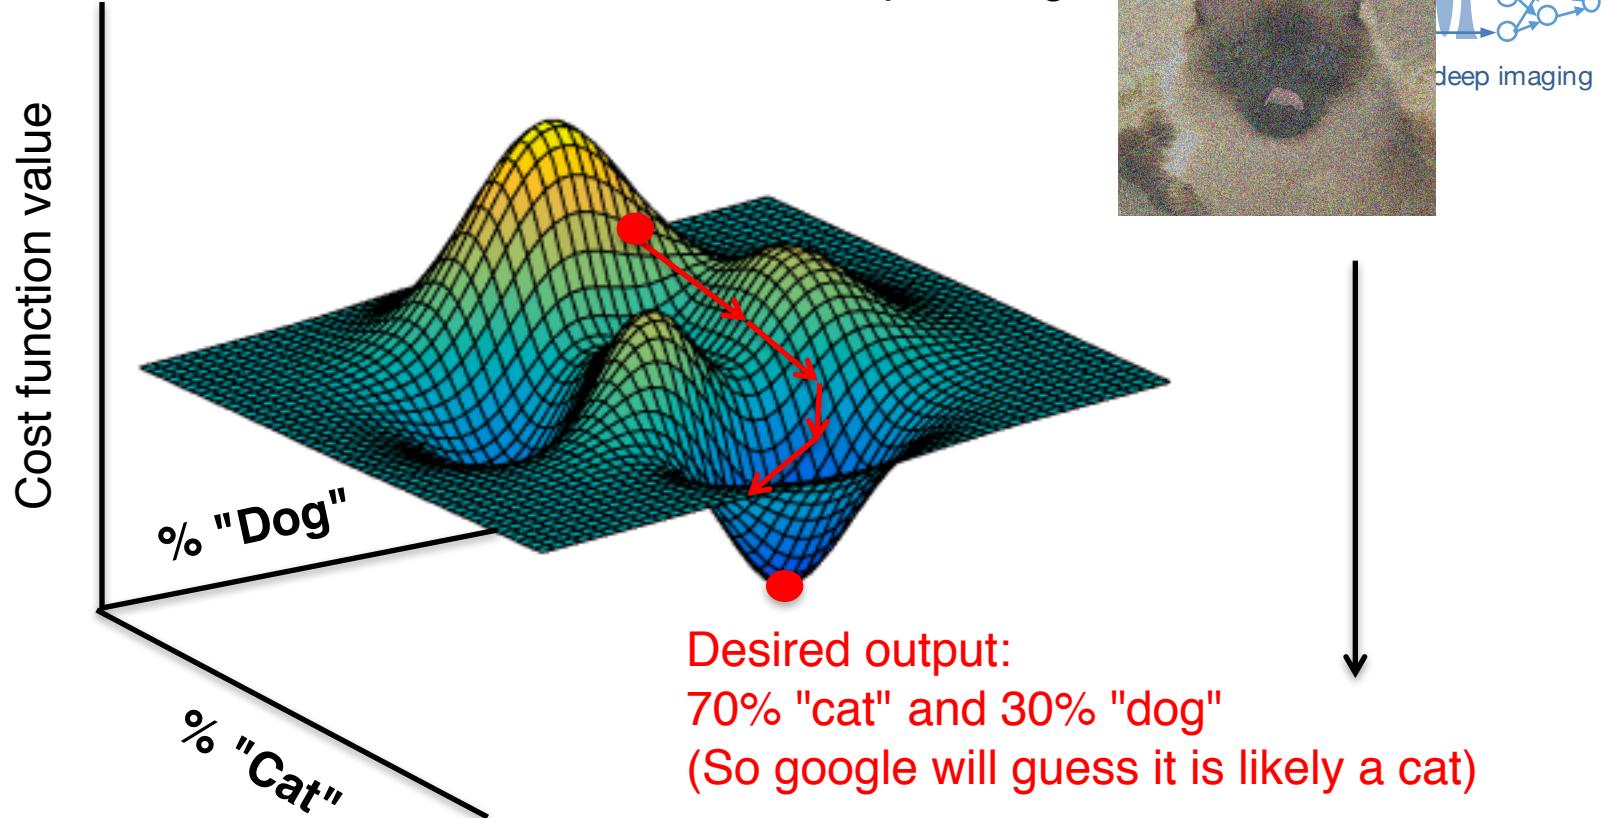


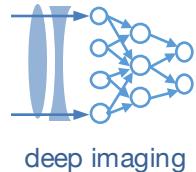
Image Classification: "Is the image of a dog or a cat"?



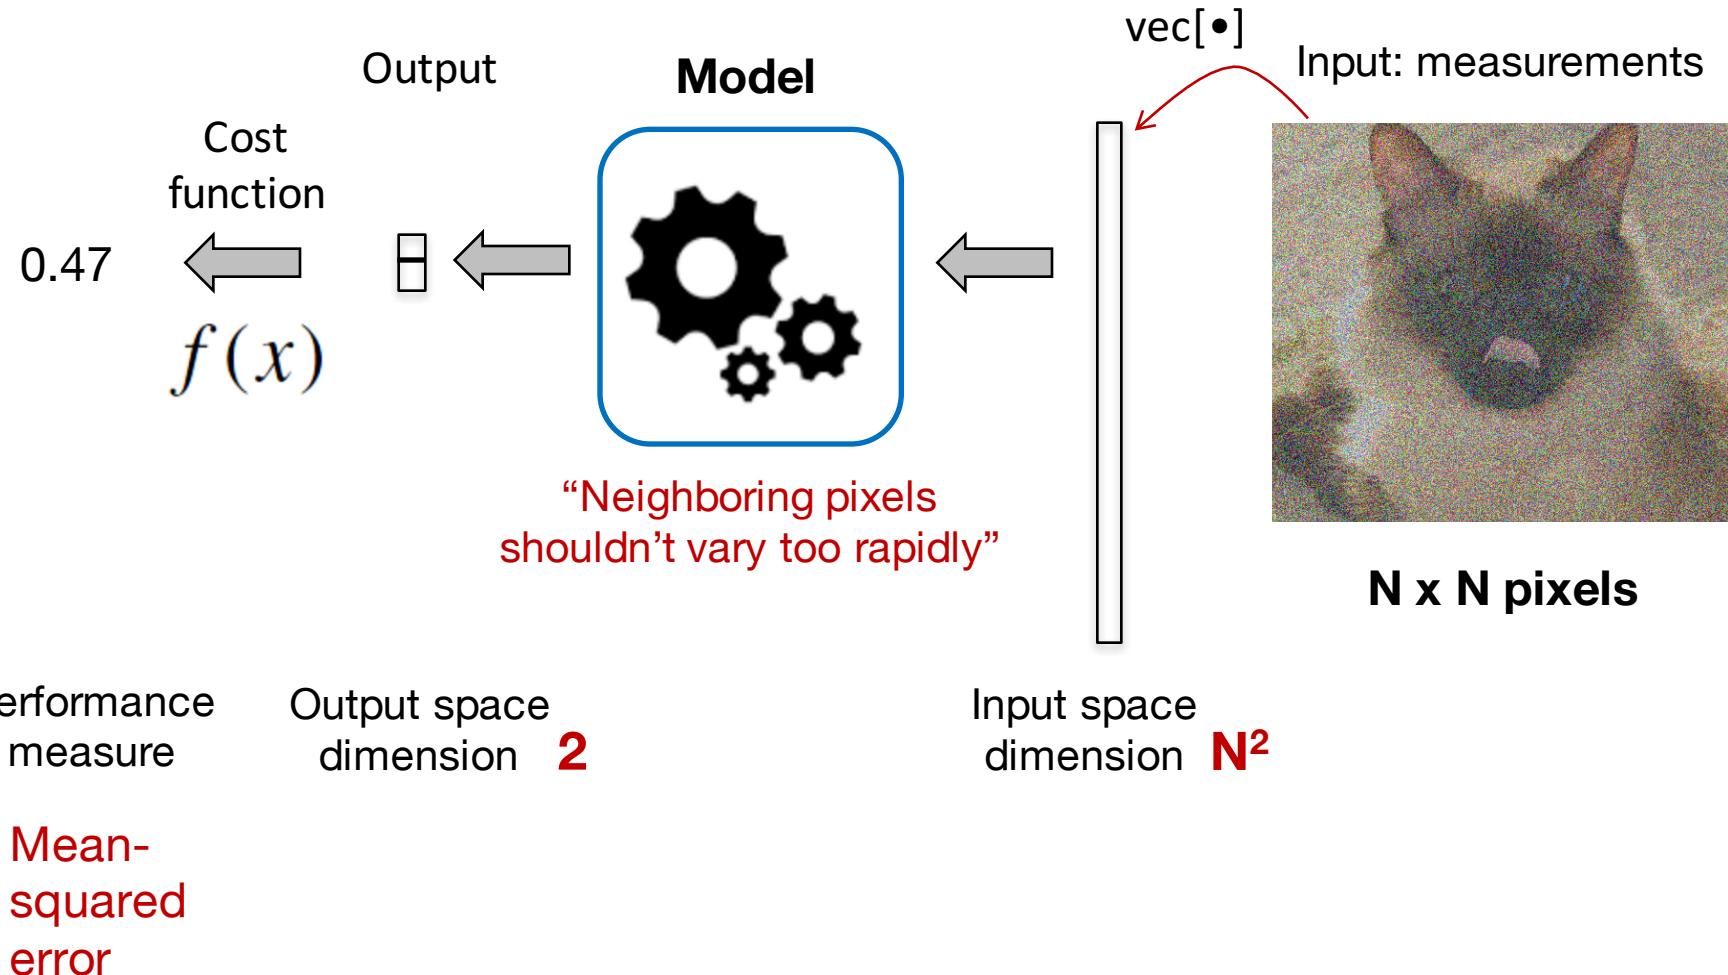
Input image

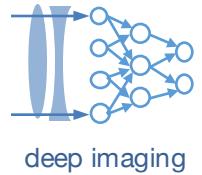






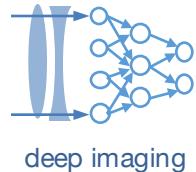
Optimization pipeline for classification





A simple example: spectral unmixing

(For whatever reason, whenever I get confused about optimization, I think about this example)

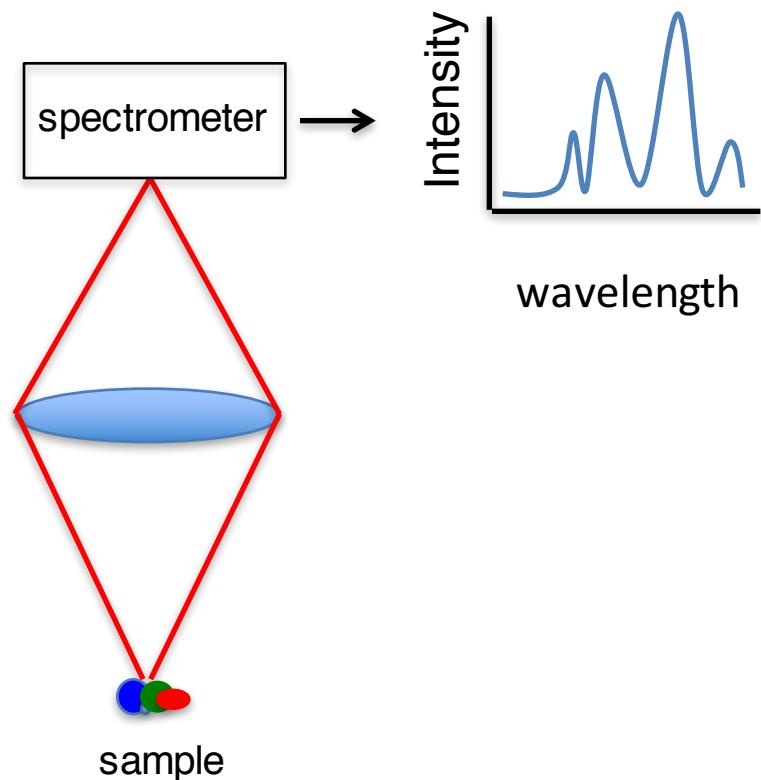


A simple example: spectral unmixing

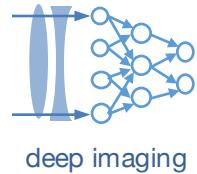
(For whatever reason, whenever I get confused about optimization, I think about this example)

The setup:

- measure the color (spectral) response of a sample (e.g., how much red, green and blue there is, or several hundred measurements of its different colors).
- You know that the sample can only contain 9 different fluorophores.
- What % of each fluorophores is in your sample?

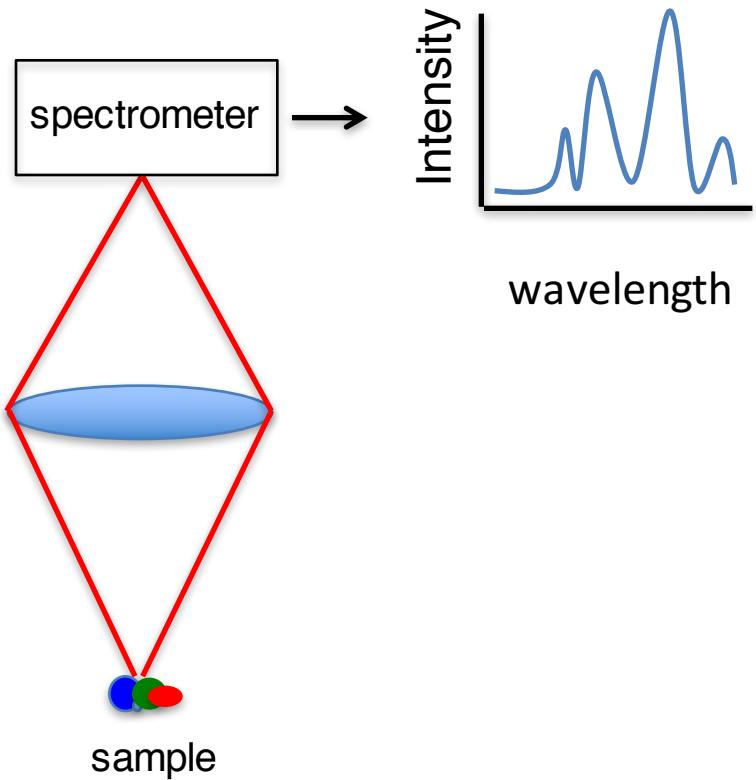


A simple example: spectral unmixing



3 elements of optimization:

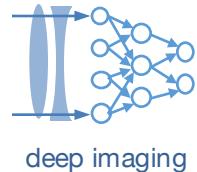
1) Desired output



2) The model

3) The cost function

A simple example: spectral unmixing



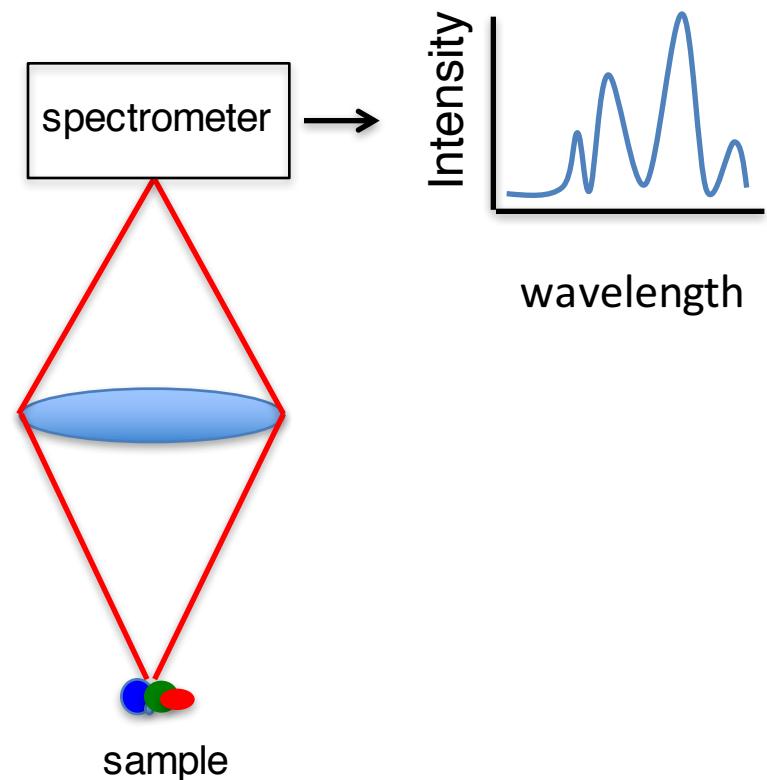
3 elements of optimization:

1) Desired output

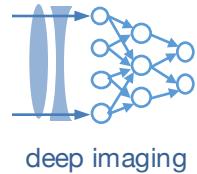
What % of each of the 9 fluorophores

2) The model

3) The cost function



A simple example: spectral unmixing



3 elements of optimization:

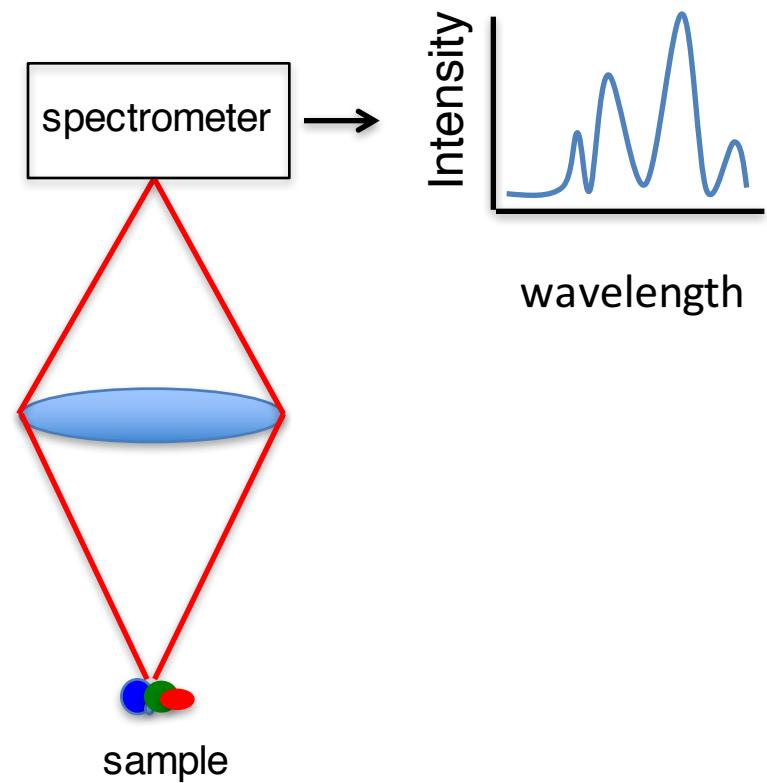
1) Desired output

What % of each of the 9 fluorophores

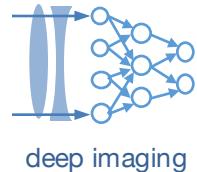
2) The model

"Dictionary" of the 9 different spectra

3) The cost function



A simple example: spectral unmixing



3 elements of optimization:

1) Desired output

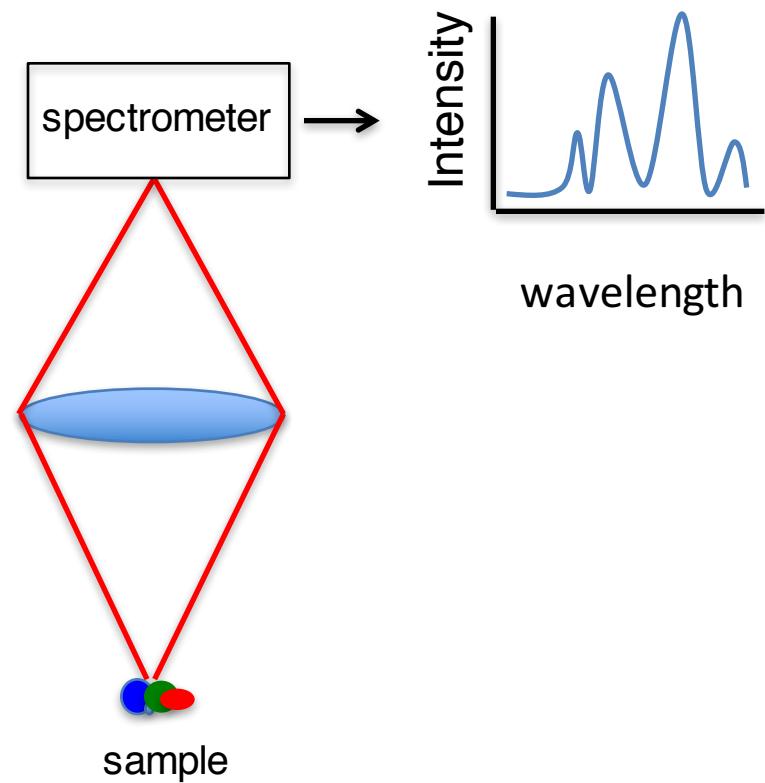
What % of each of the 9 fluorophores

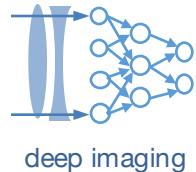
2) The model

"Dictionary" of the 9 different spectra

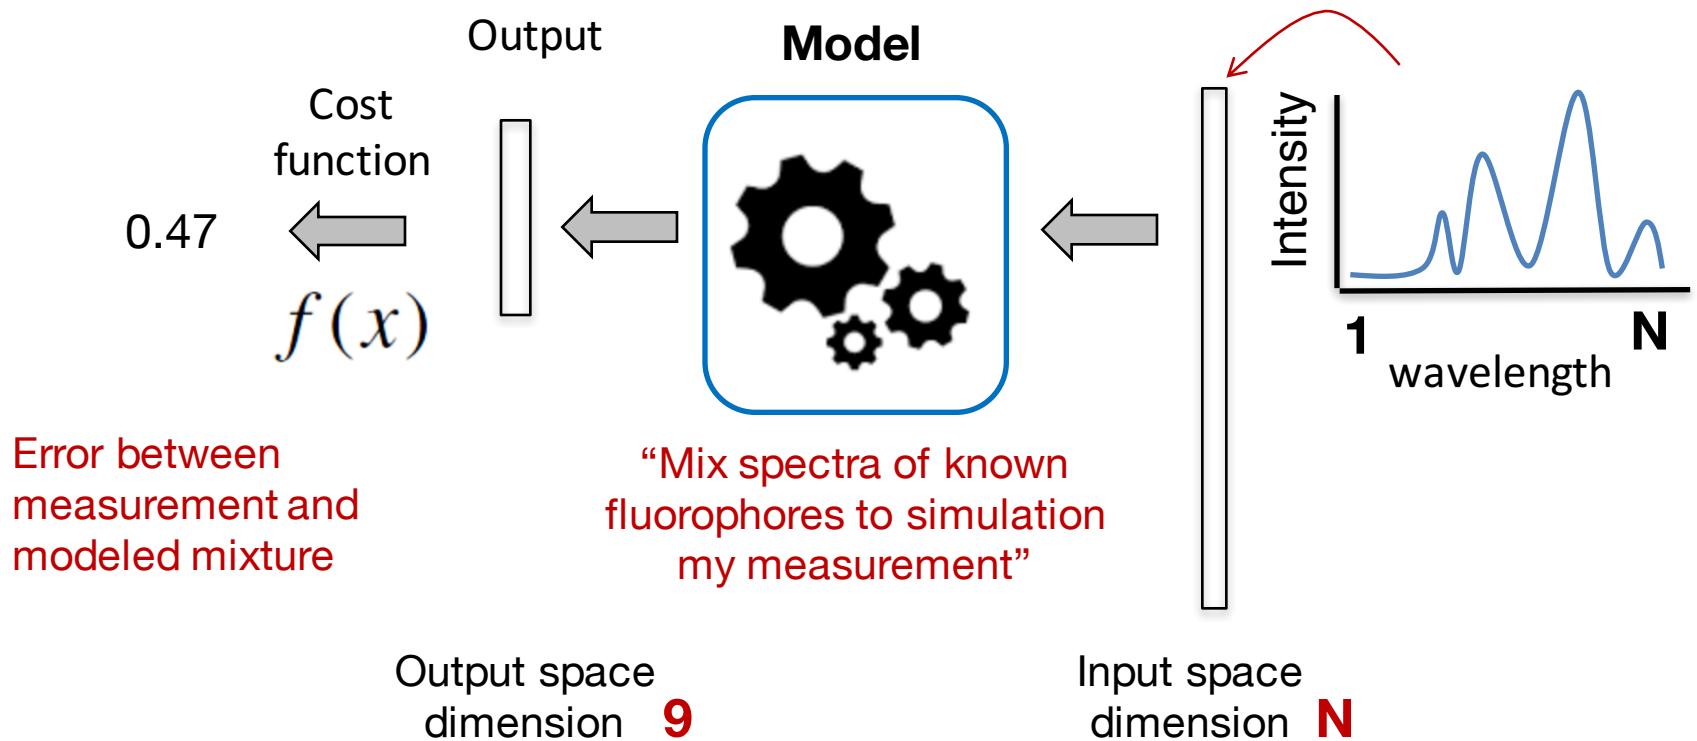
3) The cost function

Minimum mean squared error (to start)



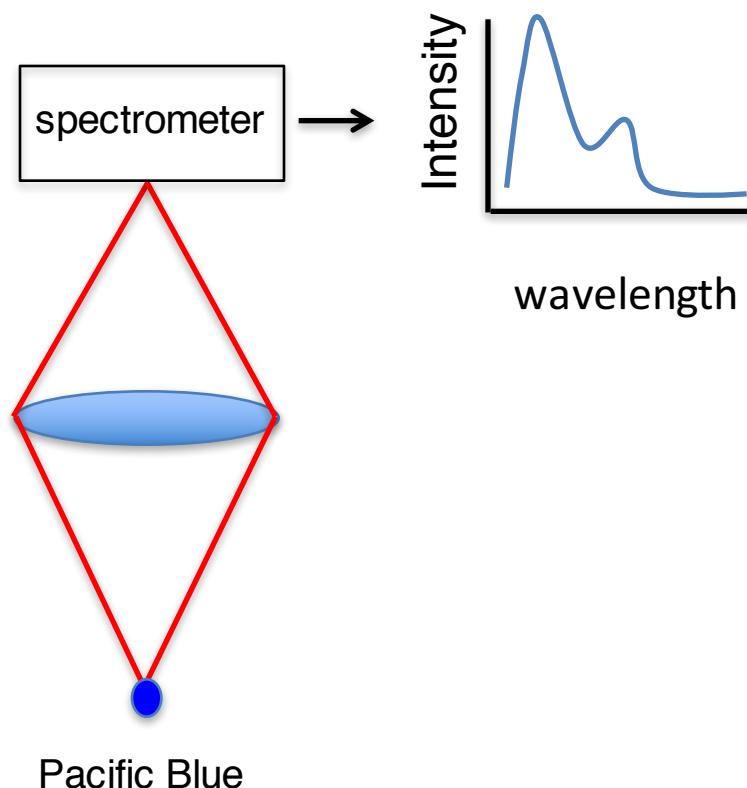


Optimization pipeline for spectral unmixing



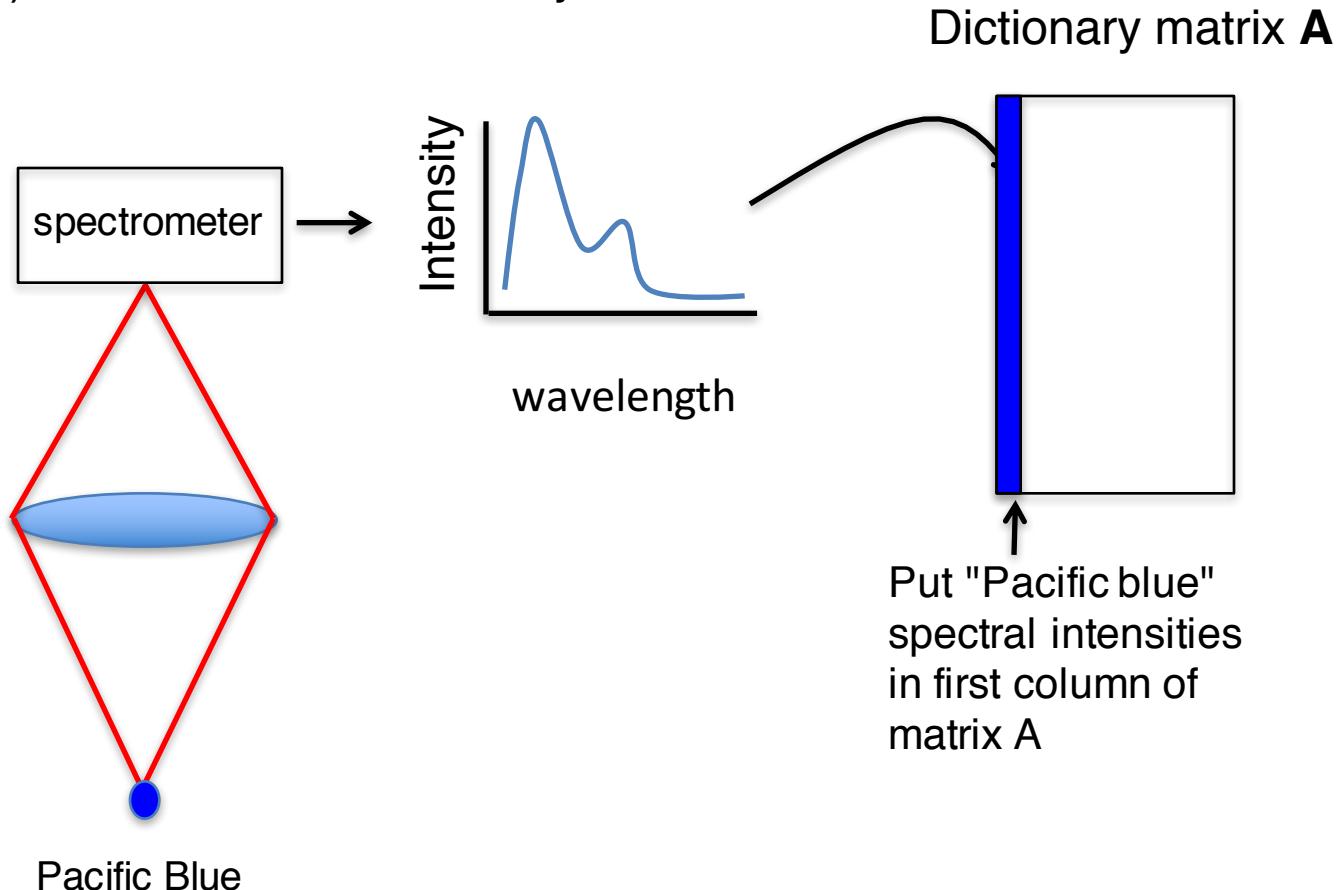
Mathematical model for spectral unmixing

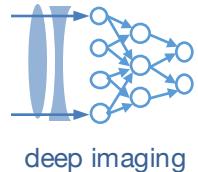
a) First make the "dictionary":



Mathematical model for spectral unmixing

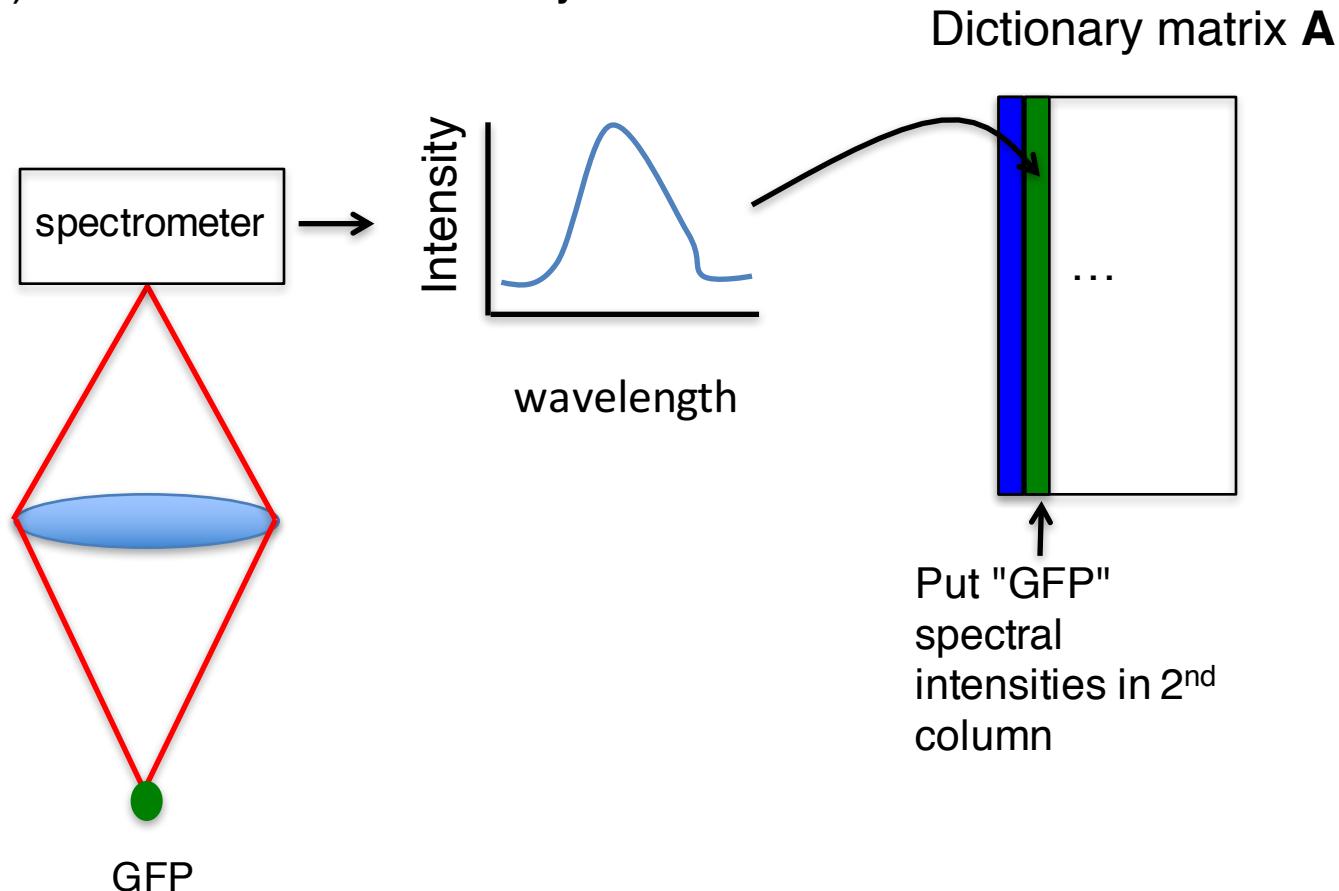
a) First make the "dictionary":

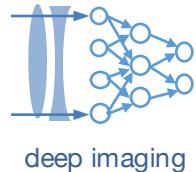




Mathematical model for spectral unmixing

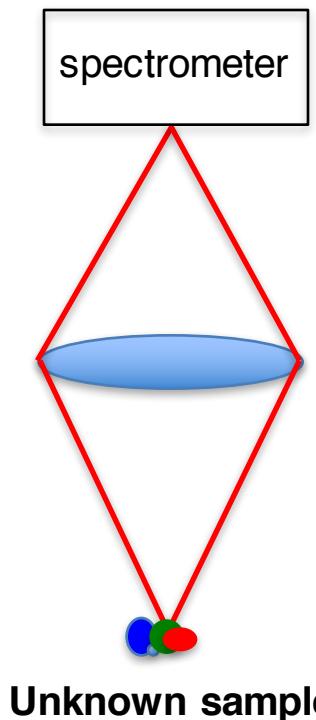
a) First make the "dictionary":



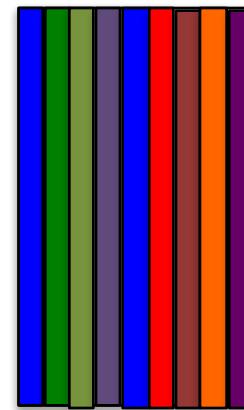


Mathematical model for spectral unmixing

b) Model the unknown sample %'s
(the desired output)

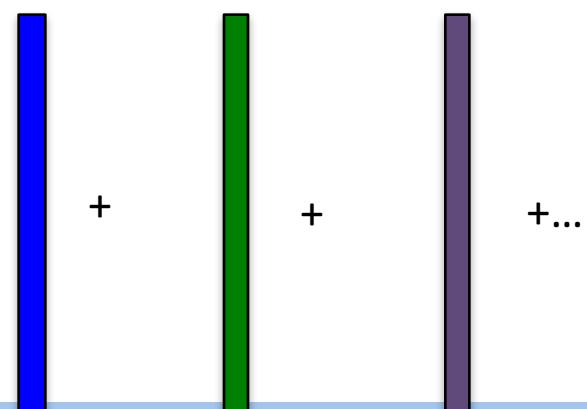


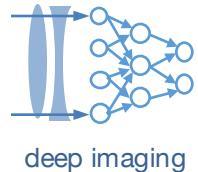
Dictionary matrix \mathbf{A}



9 possible spectra

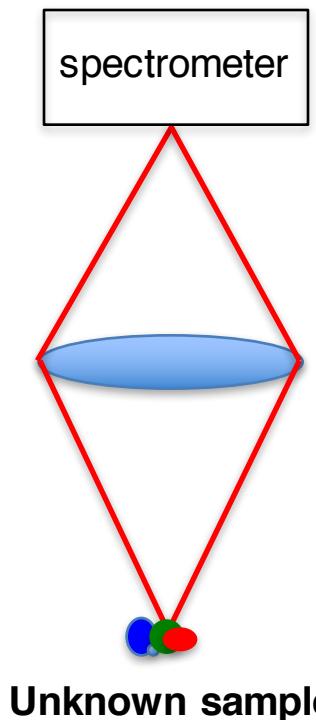
Some mixture...





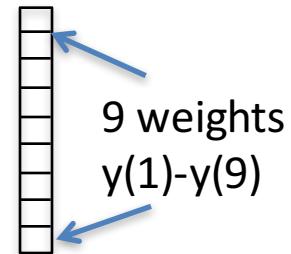
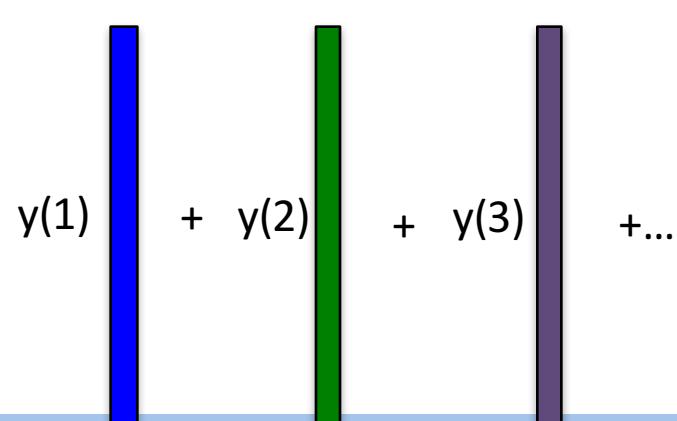
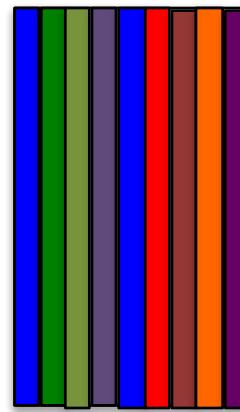
Mathematical model for spectral unmixing

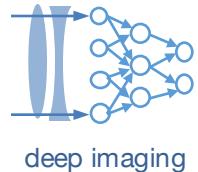
b) Model the unknown sample %'s
(the desired output)



Each weight
in x is
percentage:

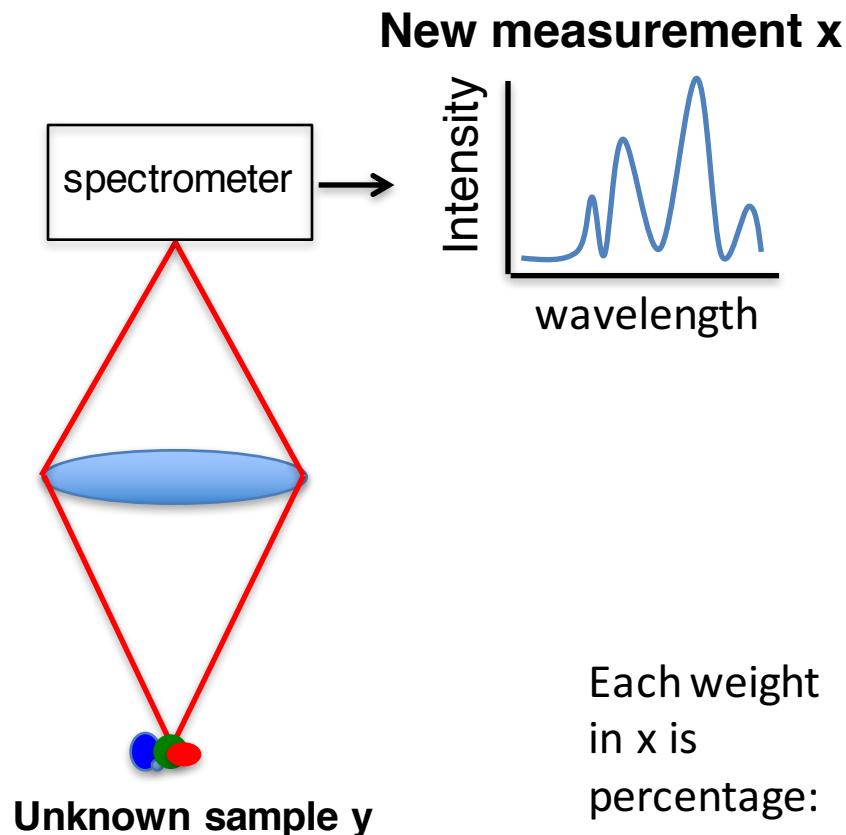
Dictionary matrix **A**
Unknown
sample %'s y





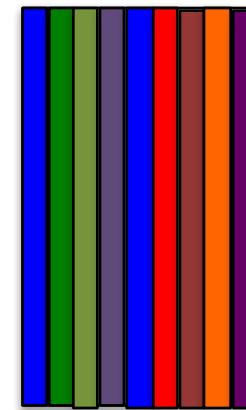
Mathematical model for spectral unmixing

b) Model the unknown sample %'s

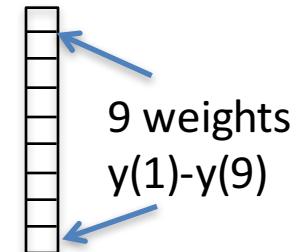


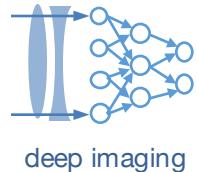
Dictionary matrix **A**

Unknown sample %'s **y**

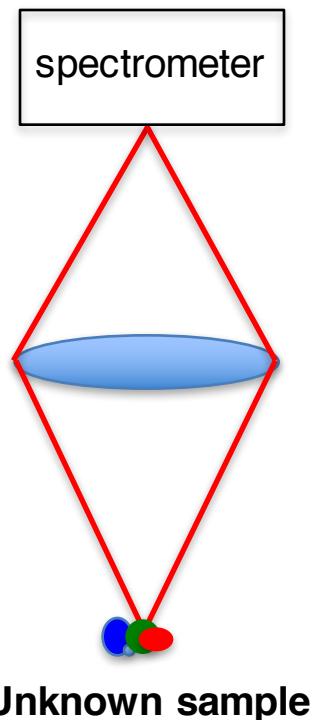


$$y(1) + y(2) + y(3) + \dots$$





Mathematical model for spectral unmixing



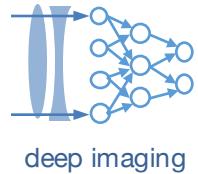
$$x = \text{Dictionary matrix } A \text{ Unknown sample } y$$

The Dictionary matrix A is represented as a vertical stack of nine colored bars: blue, green, olive green, purple, blue, red, orange, and purple. To the right, a vertical vector of nine squares is labeled "9 weights". Two blue arrows point from the text "9 weights" to the second and eighth squares of the vector.

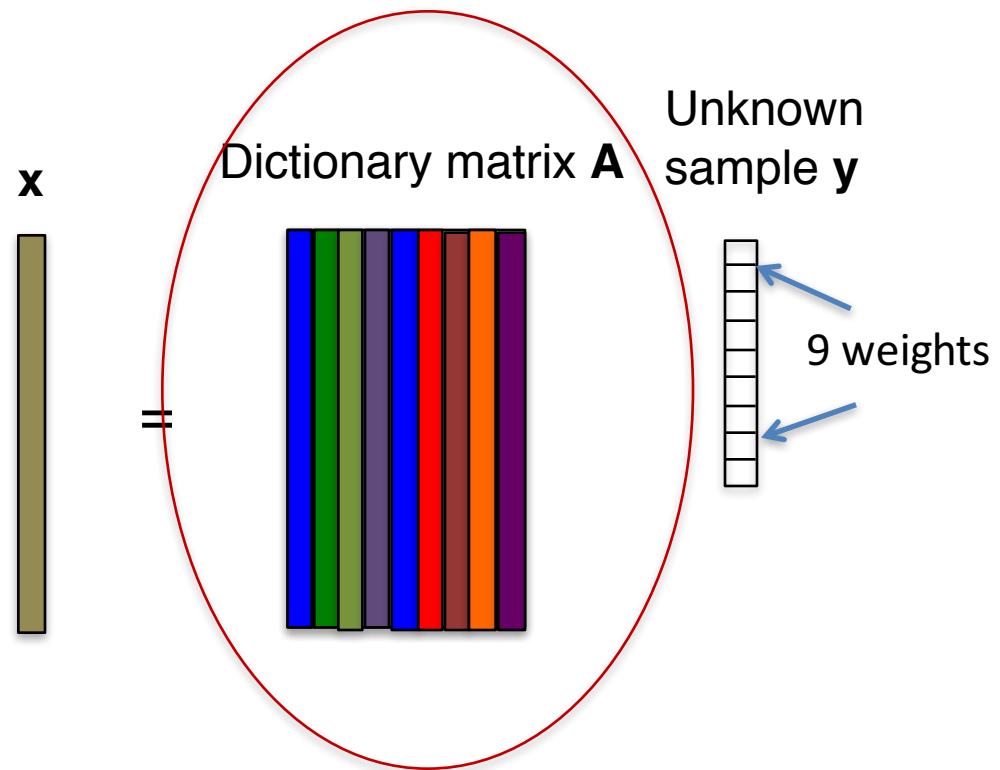
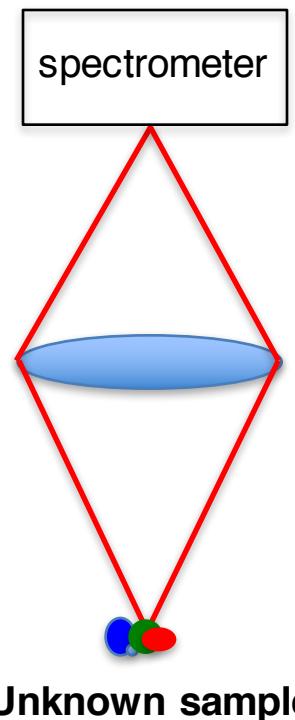
Matrix equation: $x=Ay$

This is your model!

Goal: Given A and x, find y

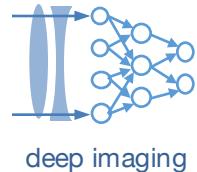


Mathematical model for spectral unmixing



Data in **A** can be thought of , in some sense, as “training data”

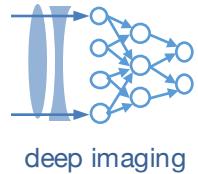
Cost function for spectral unmixing



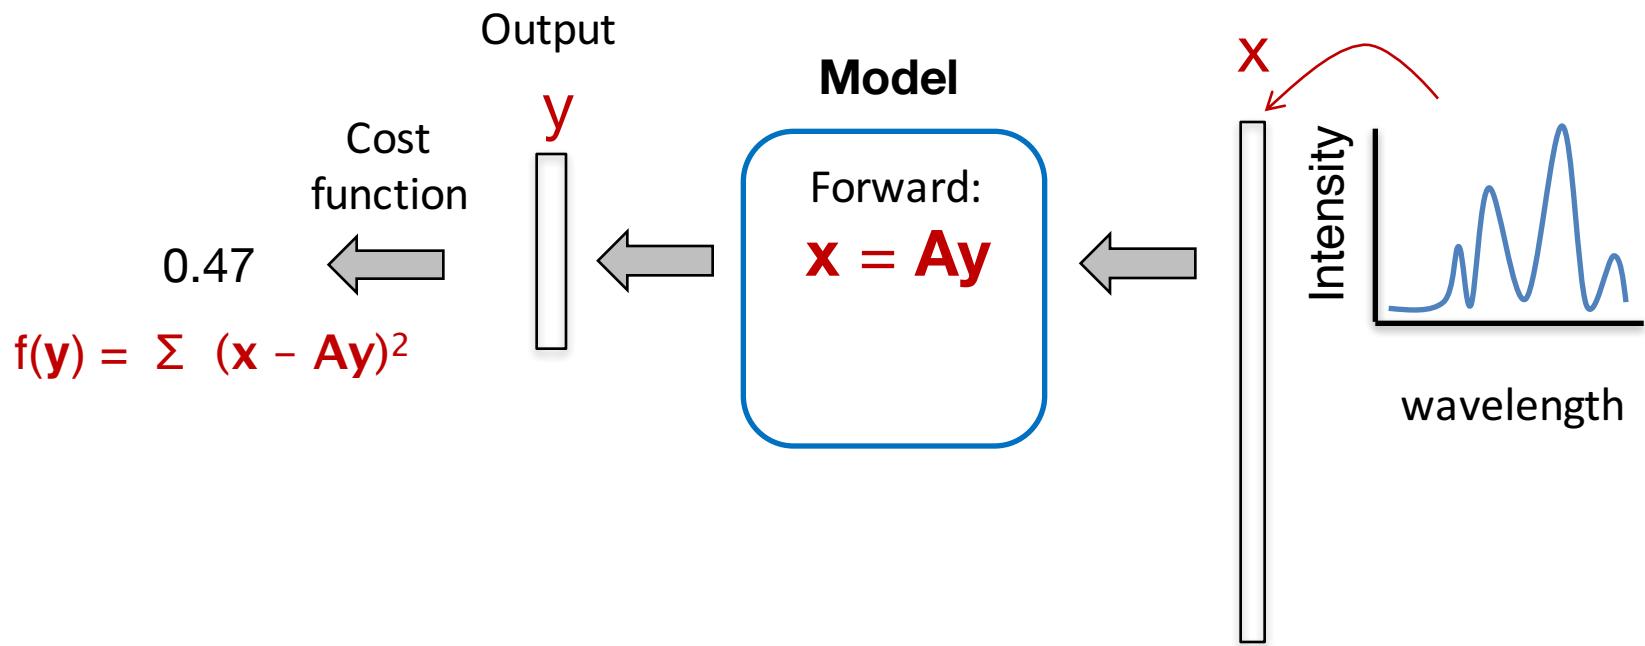
$\mathbf{x} = \mathbf{Ay}$ won't always be true, due to noise (actually, $\mathbf{x} = \mathbf{Ay} + \mathbf{n}$)

Common cost function is minimum mean-squared error:

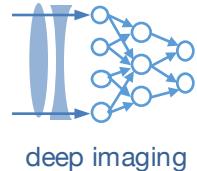
$$\text{Cost function } f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{Ay})^2$$



Optimization pipeline for denoising



Cost function for spectral unmixing



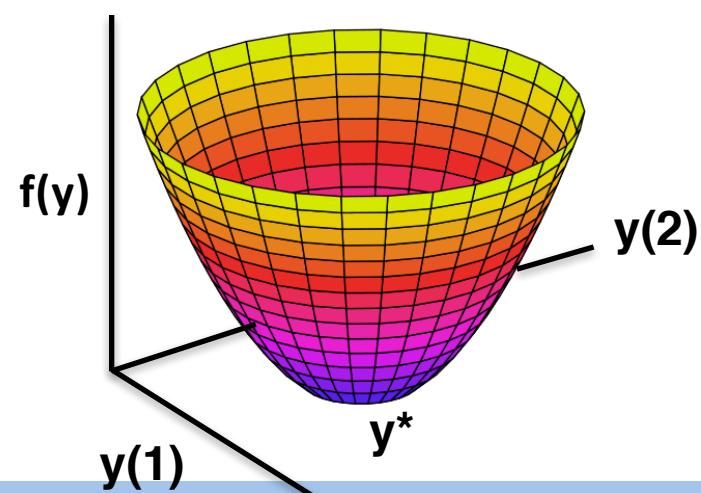
$\mathbf{x} = \mathbf{A}\mathbf{y}$ won't always be true, due to noise (actually, $\mathbf{x} = \mathbf{A}\mathbf{y} + \mathbf{n}$)

Common cost function is minimum mean-squared error:

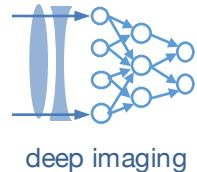
$$\text{Cost function } f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

Find mixture \mathbf{y} of known spectra \mathbf{A} that is as close as possible to measurement \mathbf{x}

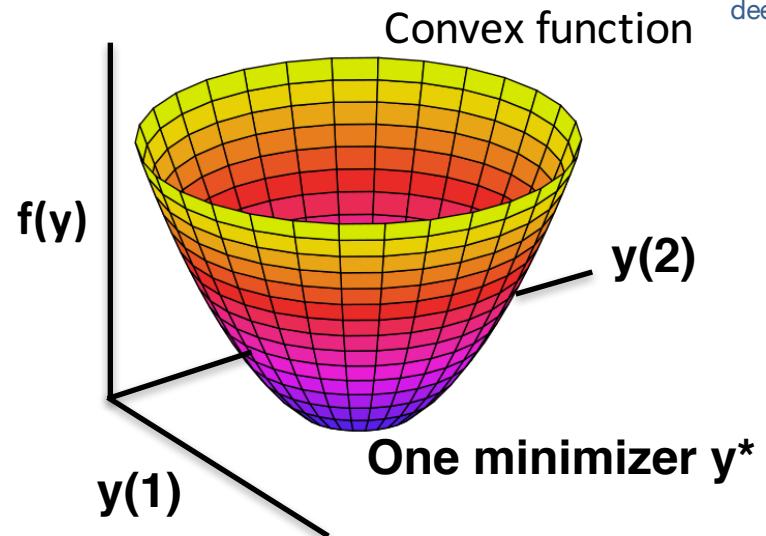
$$\mathbf{y}^* = \text{minimize } f(\mathbf{y})$$



Cost function for spectral unmixing



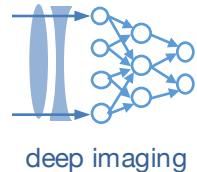
$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$



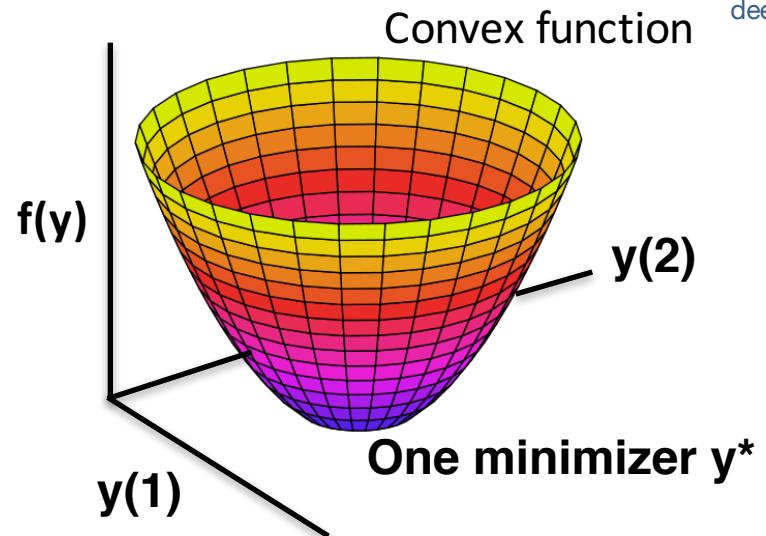
$f(\mathbf{y})$ is convex, so finding \mathbf{y}^* is easy via its gradient:



Cost function for spectral unmixing



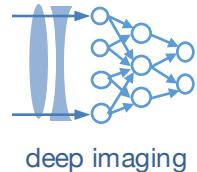
$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$



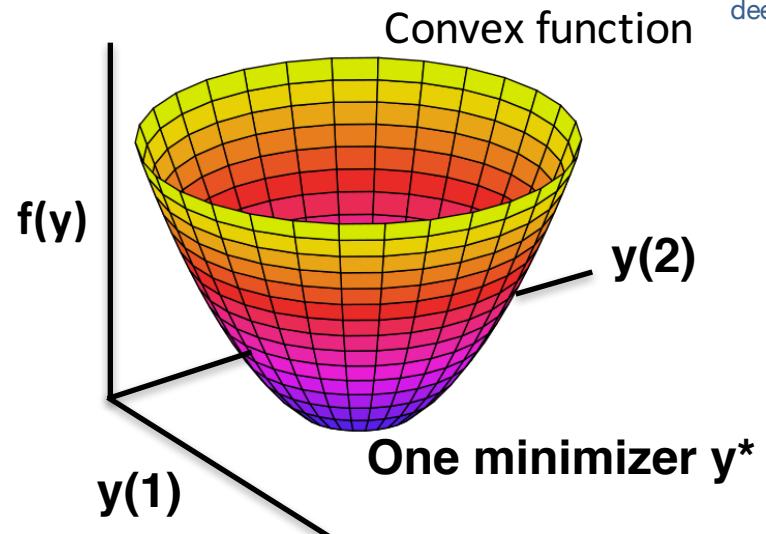
$f(\mathbf{x})$ is convex, so finding \mathbf{x}^* is easy via its gradient:

$$\frac{d}{d\mathbf{y}} f(\mathbf{y}) = \frac{d}{d\mathbf{y}} \sum (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

Cost function for spectral unmixing



$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

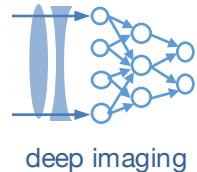


$f(\mathbf{x})$ is convex, so finding \mathbf{x}^* is easy via its gradient:

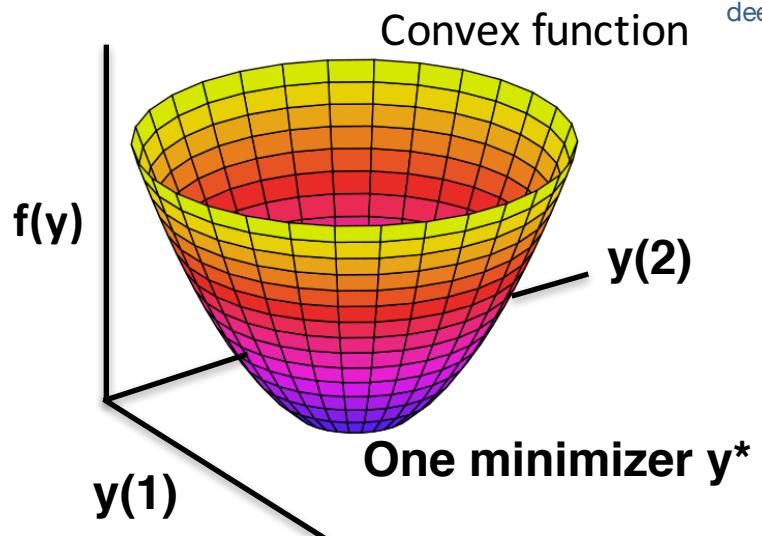
$$\frac{d}{d\mathbf{y}} f(\mathbf{y}) = \frac{d}{d\mathbf{y}} \sum (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

$$df/d\mathbf{y} = \sum d/d\mathbf{y} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

Cost function for spectral unmixing



$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$



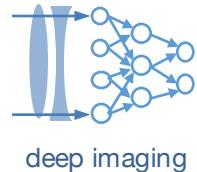
$f(\mathbf{x})$ is convex, so finding \mathbf{x}^* is easy via its gradient:

$$\frac{d}{d\mathbf{y}} f(\mathbf{y}) = \frac{d}{d\mathbf{y}} \sum (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

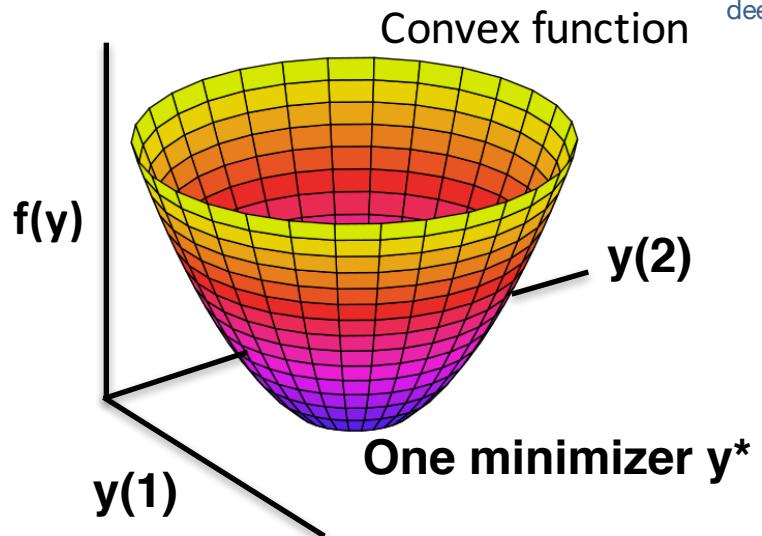
$$df/d\mathbf{y} = \sum d/d\mathbf{y} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

$$df/d\mathbf{y}[j] = \sum 2 (\mathbf{x} - \mathbf{A}\mathbf{y}) \cdot^* \mathbf{a}(:,j)$$

Cost function for spectral unmixing



$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$



$f(\mathbf{x})$ is convex, so finding \mathbf{x}^* is easy via its gradient:

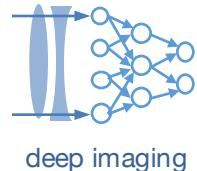
$$\frac{d}{d\mathbf{y}} f(\mathbf{y}) = \frac{d}{d\mathbf{y}} \sum (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

$$df/d\mathbf{y} = \sum d/d\mathbf{y} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$

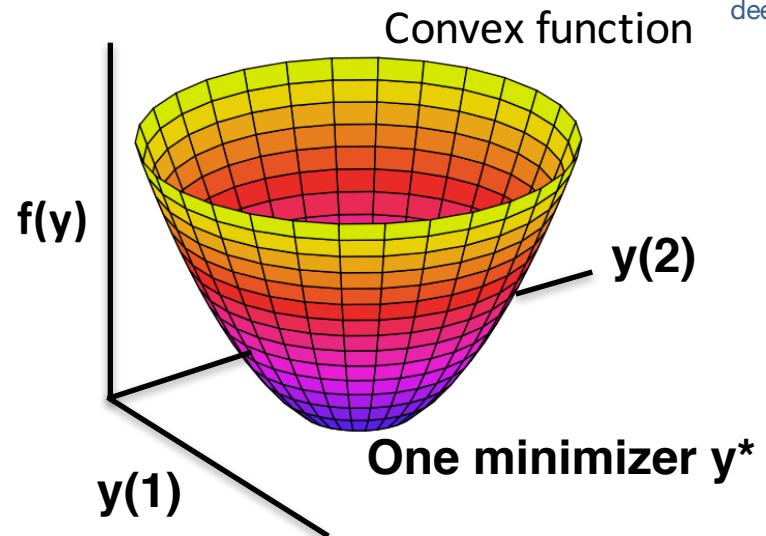
$$df/d\mathbf{y}[j] = \sum 2 (\mathbf{x} - \mathbf{A}\mathbf{y}) \cdot^* \mathbf{a}(:,j)$$

$$df/d\mathbf{y} = \mathbf{A}^T (\mathbf{x} - \mathbf{A}\mathbf{y}^*)$$

Cost function for spectral unmixing



$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$



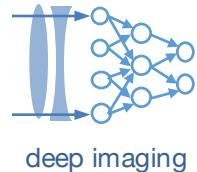
Method 1: *Gradient descent* – follow gradient downhill to solution \mathbf{y}^*

Algorithm 4.1 An algorithm to minimize $f(\mathbf{x}) = \frac{1}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2$ with respect to \mathbf{x} using gradient descent, starting from an arbitrary value of \mathbf{x} .

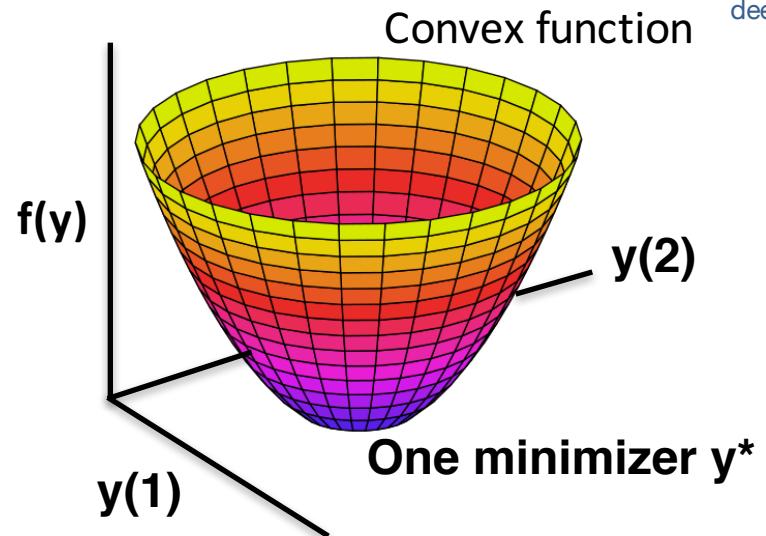
Set the step size (ϵ) and tolerance (δ) to small, positive numbers.

```
while  $\|\mathbf{A}^\top \mathbf{Ax} - \mathbf{A}^\top \mathbf{b}\|_2 > \delta$  do
     $\mathbf{x} \leftarrow \mathbf{x} - \epsilon (\mathbf{A}^\top \mathbf{Ax} - \mathbf{A}^\top \mathbf{b})$ 
end while
```

Cost function for spectral unmixing



$$f(\mathbf{y}) = \sum_{\text{spectral measurements}} (\mathbf{x} - \mathbf{A}\mathbf{y})^2$$



Method 2: *Direct solution* – set derivative to 0 to find \mathbf{y}^* directly

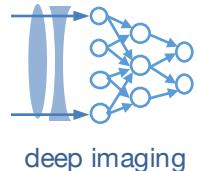
$$\frac{df}{d\mathbf{y}} = \mathbf{A}^T (\mathbf{b} - \mathbf{A}\mathbf{y}^*) = 0 \quad \leftarrow \mathbf{y}^* \text{ is where gradient of } f(\mathbf{y}) \text{ is zero}$$

$$\mathbf{A}^T \mathbf{x} = \mathbf{A}^T \mathbf{A} \mathbf{y}^* \quad \longrightarrow$$

$$(\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{x} = \mathbf{y}^*$$

"Moore-Penrose Pseudo-inverse"

(Note: setting gradient to 0 and solving is usually hard to do...)



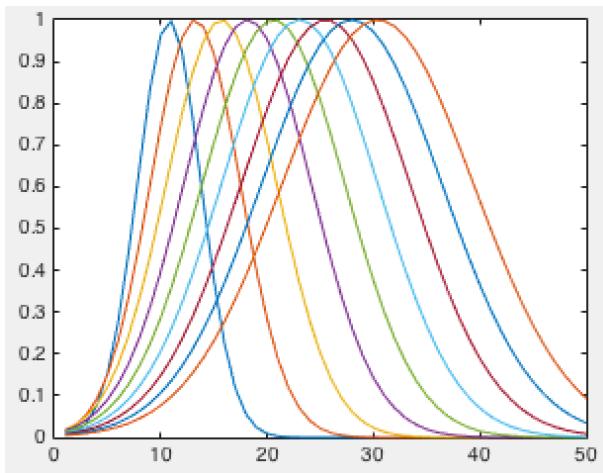
Example unmixing with the pseudo-inverse

Moore-Penrose Pseudo-inverse:

$$\mathbf{y}^* = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{x}$$

Example dictionary A

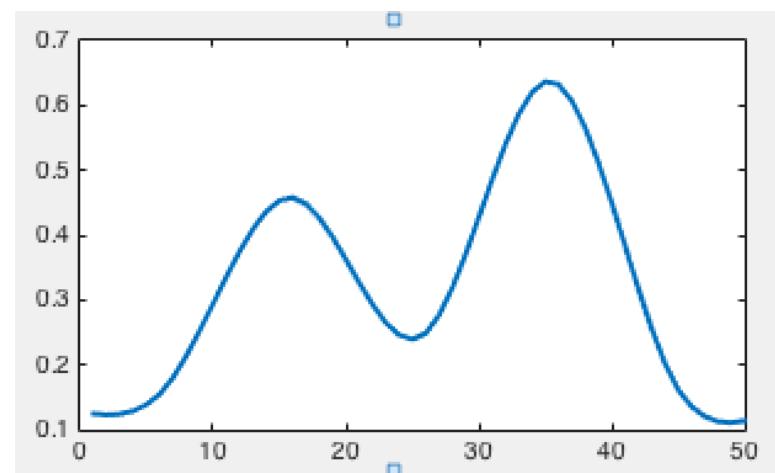
9 spectra



Example y,
compute Ay

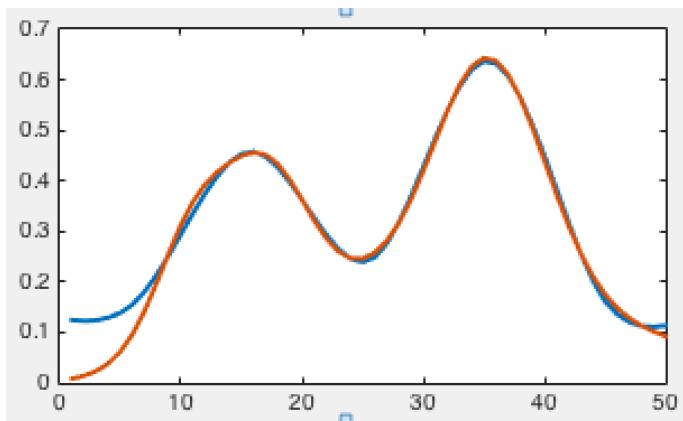


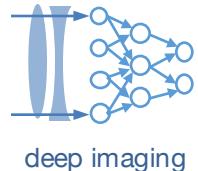
Example detected spectra x



Compute
pseudo-inverse,
 $\mathbf{x}^* = \mathbf{A}\mathbf{y}^*$ is red
curve:

Good fit!





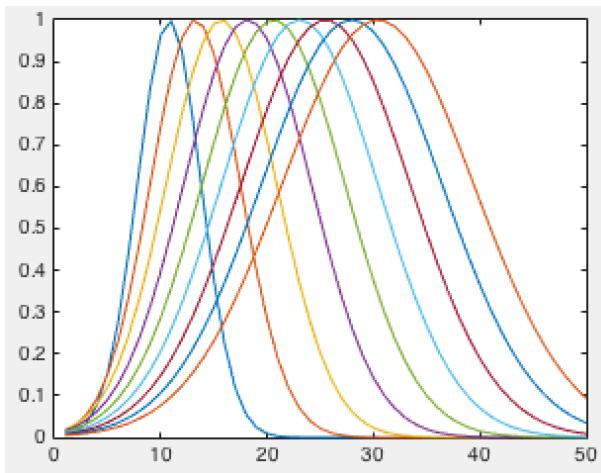
Example unmixing with the pseudo-inverse

Moore-Penrose Pseudo-inverse:

$$\mathbf{y}^* = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{x}$$

Example dictionary A

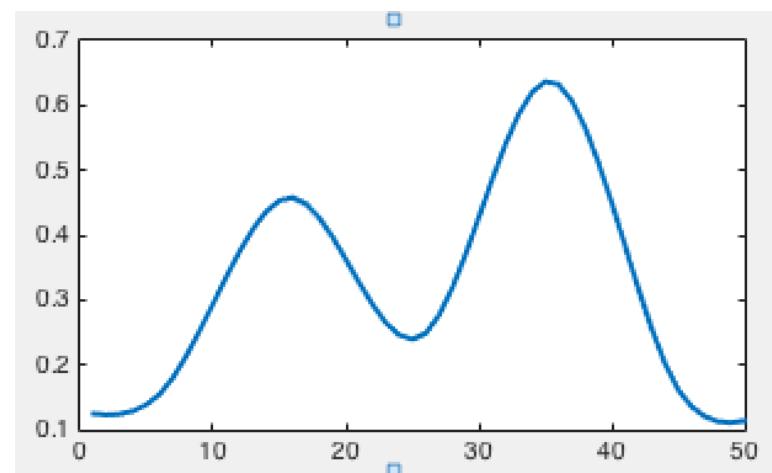
9 spectra



Example y,
compute Ay

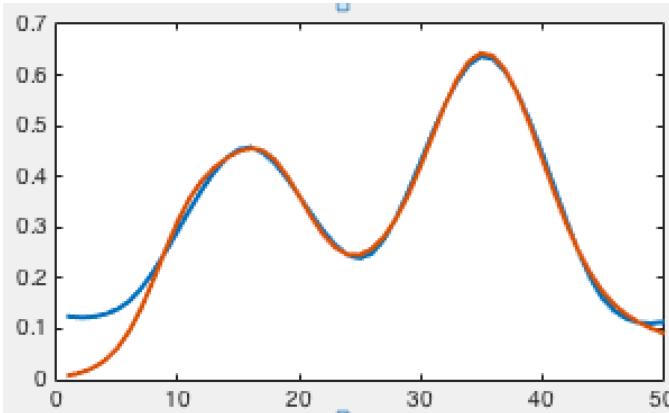


Example detected spectra x



Compute
pseudo-inverse,
 $\mathbf{x}^* = \mathbf{A}\mathbf{y}^*$ is red
curve:

Good fit!



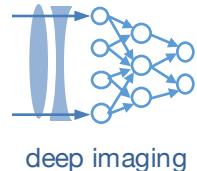
PROBLEM:

$$\mathbf{y}^* = [0.2, -1.1, -1.6, \dots]$$

Solution has negative weights!

Not physically possible...

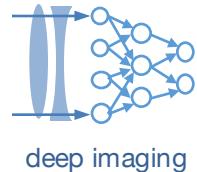
Example unmixing with the pseudo-inverse



Moore-Penrose Pseudo-inverse:
$$\mathbf{y}^* = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{x}$$

```
n = 50; %number of pixels
m = 9; %number of spectral
A=zeros(n,m); %known dictionary of spectra
for j=1:m
    A(:,j) = exp(-(linspace(-1,1,n)+.5-.1*j+.2).^2/(.03*j));
end
%Simulate some spectra
b = imresize(rand([5,1]),[n 1]);
x_opt = A\b;      ← Pseudo-inverse = one line
%Show results
figure;plot(b); hold all; plot(A*x_opt);
```

Spectral un-mixing with a positivity constraint



Option 1: Add a constraint

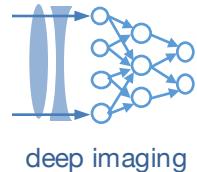
Minimize $f(\mathbf{x}) = \sum (\mathbf{b} - \mathbf{Ax})^2$ Convex cost function

Subject to $\mathbf{x} \geq 0$ Convex constraint

*When you have constraints, can use **CVX**, convex toolbox for Matlab

<http://cvxr.com/cvx/>

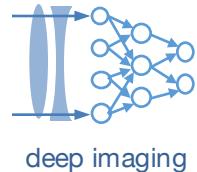
Spectral un-mixing with a positivity constraint



Option 1: Add a constraint

```
%%%%%
addpath '/users/Roarke/Documents/Matlab/cvx'; cvx_setup;
cvx_begin
    variable xc(m);
    minimize( norm(A*xc-b) );
    subject to
        xc  $\geq 0$ ;
cvx_end
%Show results
figure;plot(b); hold all; plot(A*xc);
%%%%%
```

Spectral un-mixing with a positivity constraint



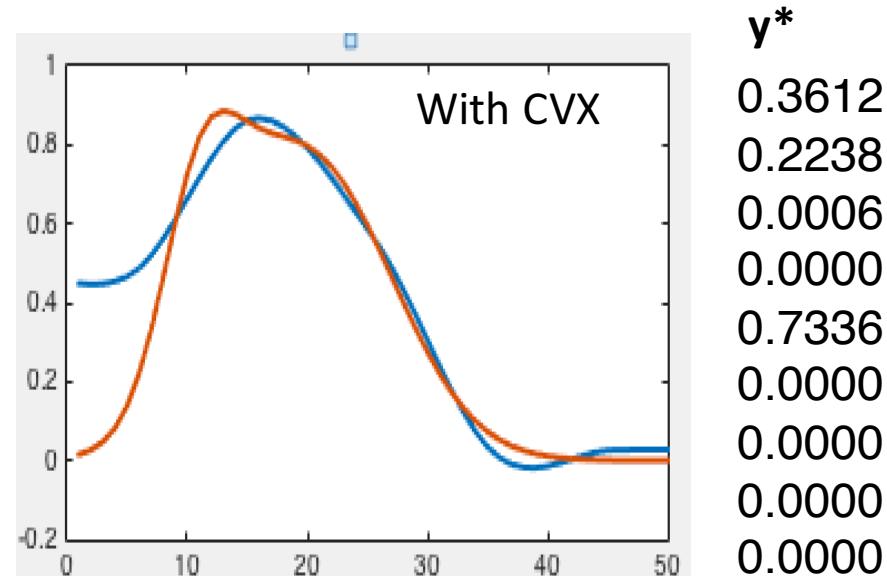
Option 1: Add a constraint

Minimize $f(\mathbf{y}) = \sum (x - \mathbf{A}\mathbf{y})^2$ Convex cost function

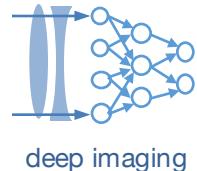
Subject to $\mathbf{y} \geq 0$ Convex constraint

*When you have constraints, can use **CVX**, convex toolbox for Matlab

<http://cvxr.com/cvx/>



Spectral un-mixing with a positivity constraint



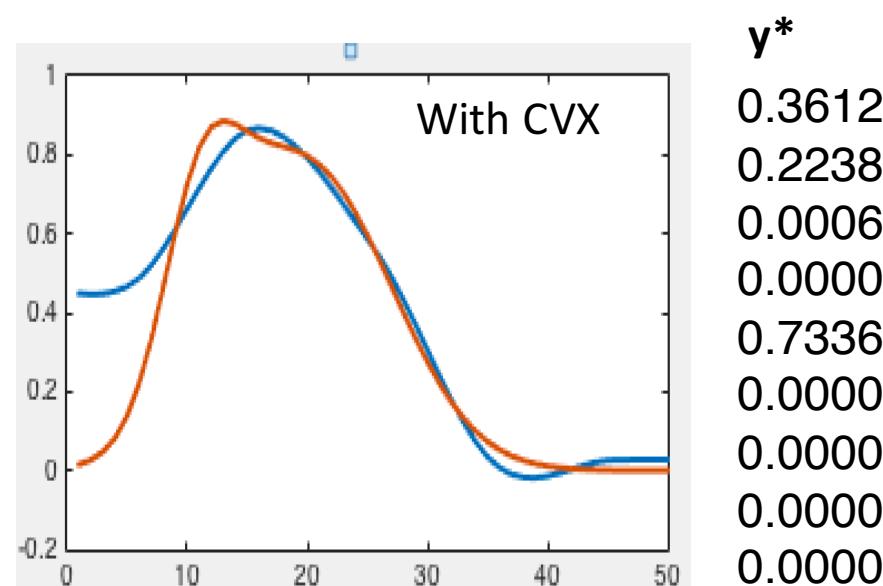
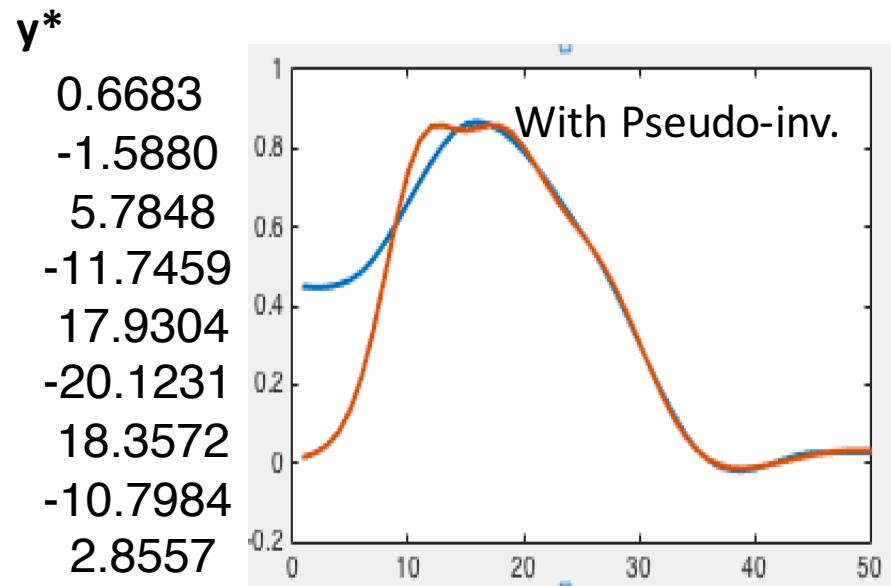
Option 1: Add a constraint

Minimize $f(\mathbf{y}) = \sum (\mathbf{x} - \mathbf{Ay})^2$ Convex cost function

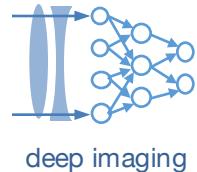
Subject to $\mathbf{y} \geq 0$ Convex constraint

*When you have constraints, can use **CVX**, convex toolbox for Matlab

<http://cvxr.com/cvx/>



Spectral un-mixing with a positivity constraint



Option 2: Modify cost function

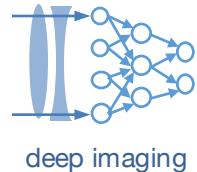
$$\text{Minimize } f(\mathbf{z}) = \sum (\mathbf{x} - \mathbf{Az}^2)^2$$

$\mathbf{z}^2 = \mathbf{y}$ is dummy variable, will change cost function and gradient

*When you don't have constraints but can find the gradient, use **Minfunc**

<https://www.cs.ubc.ca/~schmidtm/Software/minFunc.html>

Spectral un-mixing with a positivity constraint



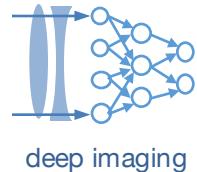
Option 2: Modify cost function

$$\text{Minimize } f(\mathbf{z}) = \sum (\mathbf{x} - \mathbf{Az}^2)^2$$

$\mathbf{z}^2 = \mathbf{y}$ is dummy variable, will change cost function and gradient

```
%%%%%
%3. Minfunc
addpath '/users/Roarke/Documents/Matlab/minFunc_2012';
startVec = ones(m,1);
spectrum_anonymous = @(startVec)spectrum(startVec, b, A);
%Evaluate with minfunc
[xm, msevalue, moreinfo] = minFunc(@(startVec)spectrum_anonymous(startVec), startVec, options);
figure; plot(b); hold all; plot(A*abs(xm).^2);
```

Spectral un-mixing with a positivity constraint



Option 2: Modify cost function

$$\text{Minimize } f(\mathbf{z}) = \sum (\mathbf{x} - \mathbf{Az}^2)^2$$

$\mathbf{z}^2 = \mathbf{y}$ is dummy variable, will change cost function and gradient

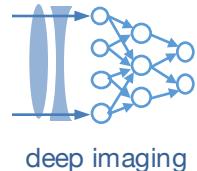
```
%%%%%
%3. Minfunc
addpath '/users/Roarke/Documents/Matlab/minFunc_2012';
startVec = ones(m,1);
spectrum_anonymous = @(startVec)spectrum(startVec, b, A);
%Evaluate with minfunc
[xm, msevalue, moreinfo] = minFunc(@(startVec)spectrum_anonymous(startVec), startVec, options);
figure;plot(b); hold all; plot(A*abs(xm).^2);
```

```
function [err_function, grad_function] = spectrum(input_vec, b, A)

%for direct pseudo-inverse - no constraints or dummy
%err_function = norm(A*input_vec - b);
%grad_function = A'*(A*input_vec - b);

err_function = norm(A*abs(input_vec).^2 - b);
grad_function = A'*((A*abs(input_vec).^2 - b) .* conj(A*input_vec));
```

Spectral un-mixing with a positivity constraint



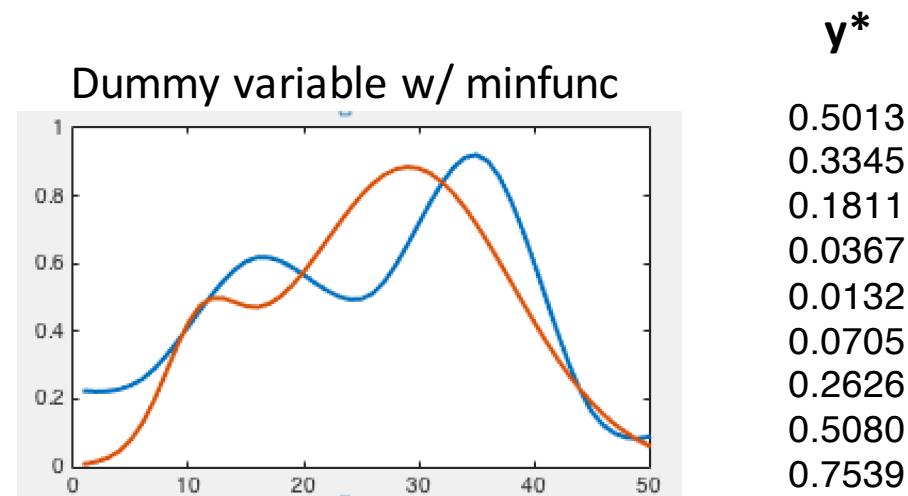
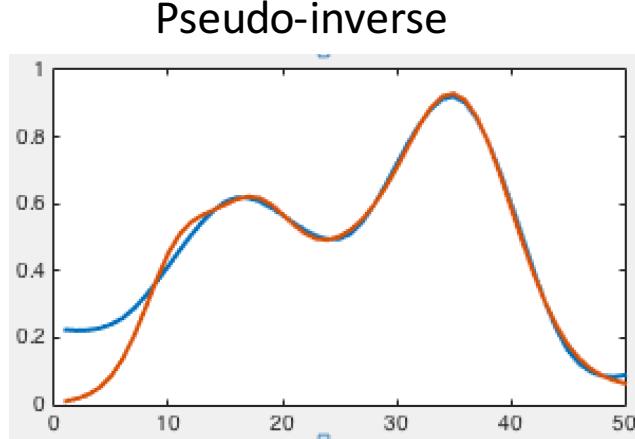
Option 2: Modify cost function

$$\text{Minimize } f(\mathbf{z}) = \sum (\mathbf{x} - \mathbf{A}\mathbf{z})^2$$

$\mathbf{z}^2 = \mathbf{y}$ is dummy variable, will change cost function and gradient

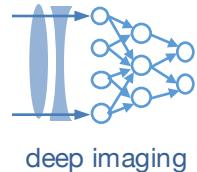
*When you don't have constraints but can find the gradient, use **Minfunc**

<https://www.cs.ubc.ca/~schmidtm/Software/minFunc.html>



Not working too well, gradient could be wrong?

Other bells and whistles



1) Sometimes see solutions where x values get really big

Fix this with a "*regularizer*":

$$\text{Minimize } f(\mathbf{x}) = \sum (\mathbf{b} - \mathbf{Ax})^2 + C^* \sum (\mathbf{x})^2$$

"Don't let x vary too much"

Choose constant C appropriately

2) If you think your signal is "sparse", then it probably has mostly zeros.
Can include this in your model with an "L1" cost function:

$$\text{Minimize } f(\mathbf{x}) = \sum | \mathbf{b} - \mathbf{Ax} |$$

- An extremely simple modification with pretty strong implications