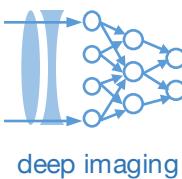


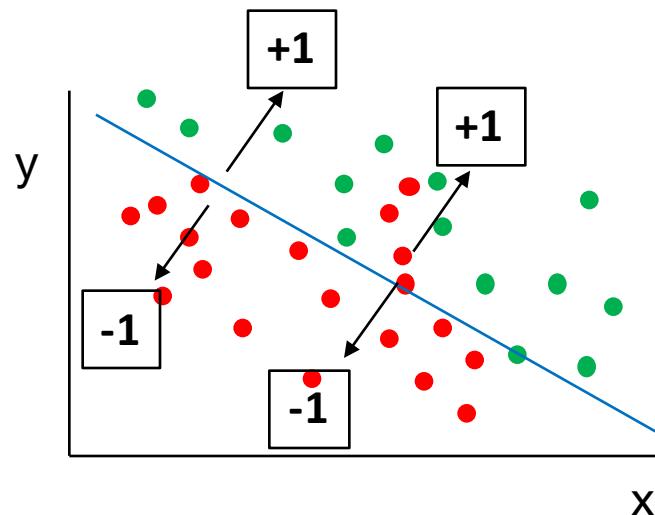
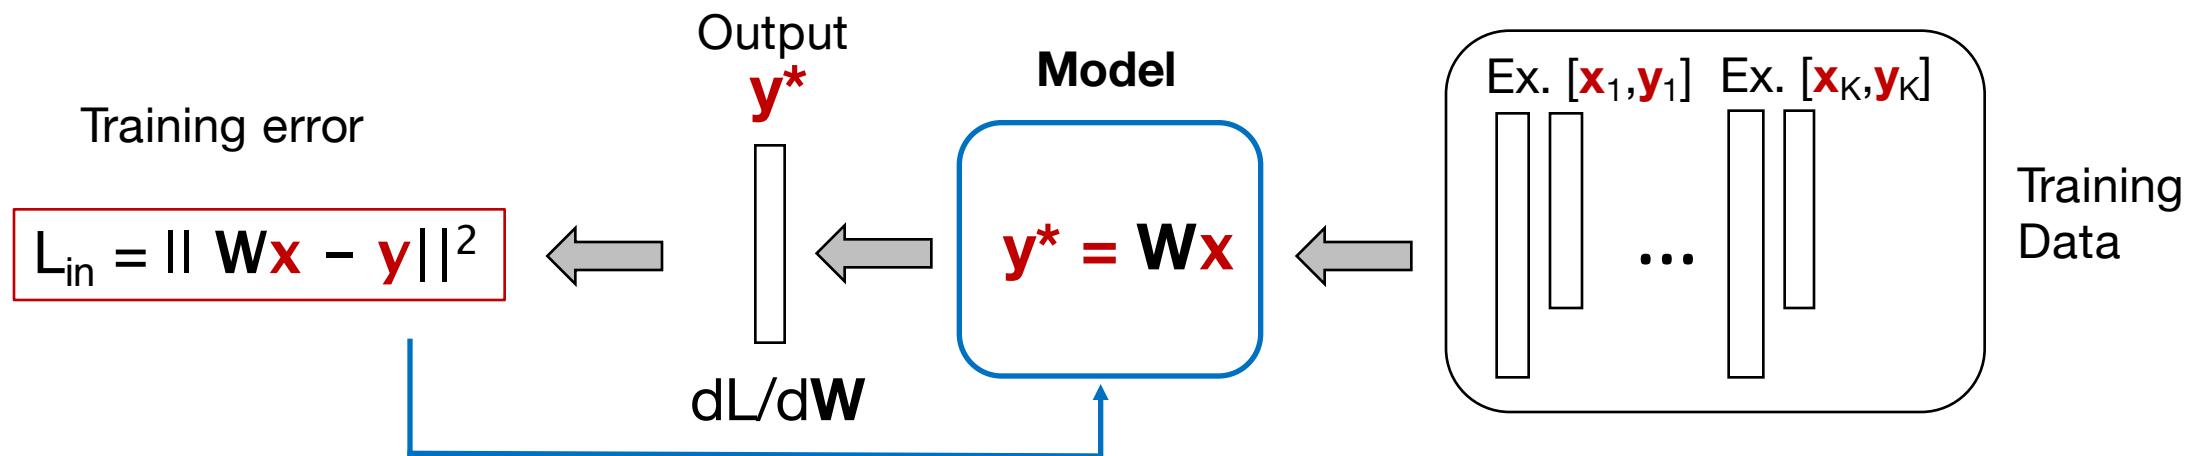
# Lecture 8: Theoretical basics of machine learning

Machine Learning and Imaging

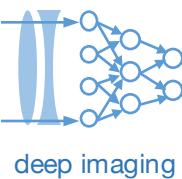
BME 590L  
Roarke Horstmeyer



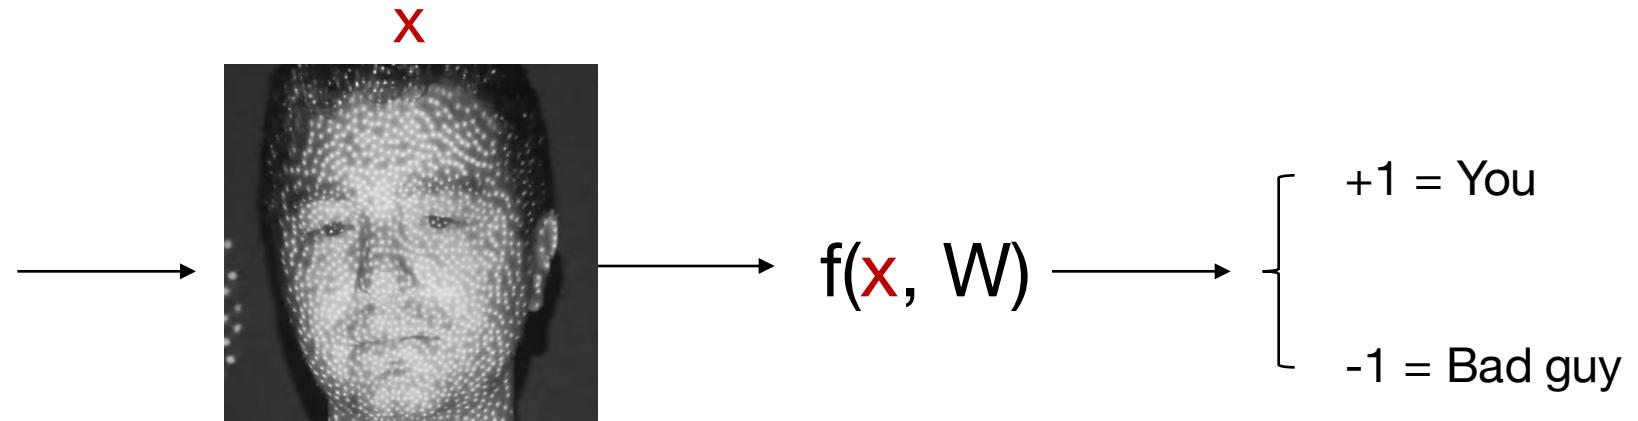
# Last time: the linear classification model – what's not to like?



1. Can only separate data with lines (hyper-planes)...
2. We only allowed for binary labels ( $y = +/- 1$ )
3. Error function  $L_{in}$  inherently makes assumptions about statistical distribution of data



## Cost functions matter: a simple example

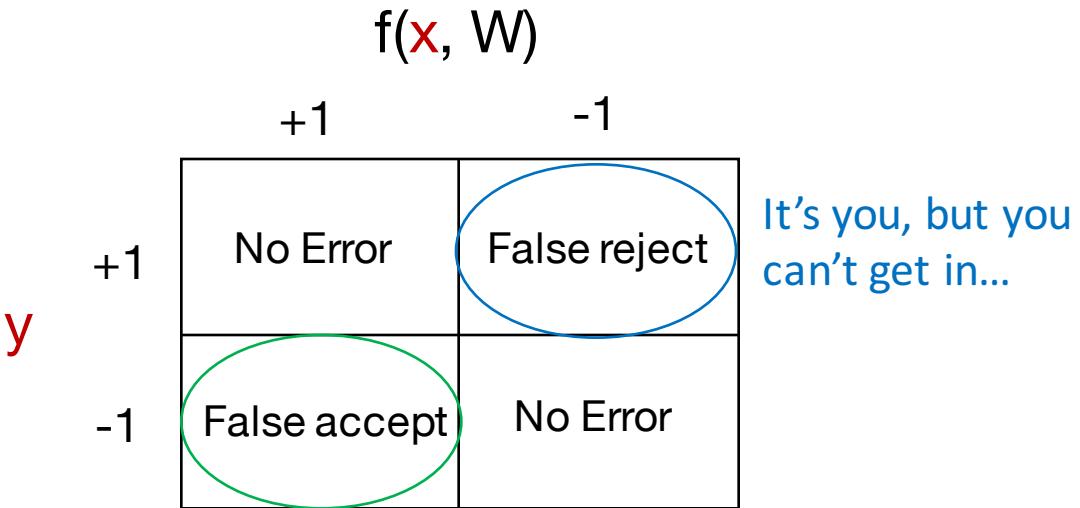


What if you're a CIA agent?

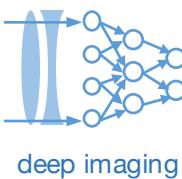
$$L_{in} = 100,000 \operatorname{ReLU}[f(\mathbf{x}, \mathbf{W}) - y] + \operatorname{ReLU}[y - f(\mathbf{x}, \mathbf{W})]$$

BIG penalty  
for intruder

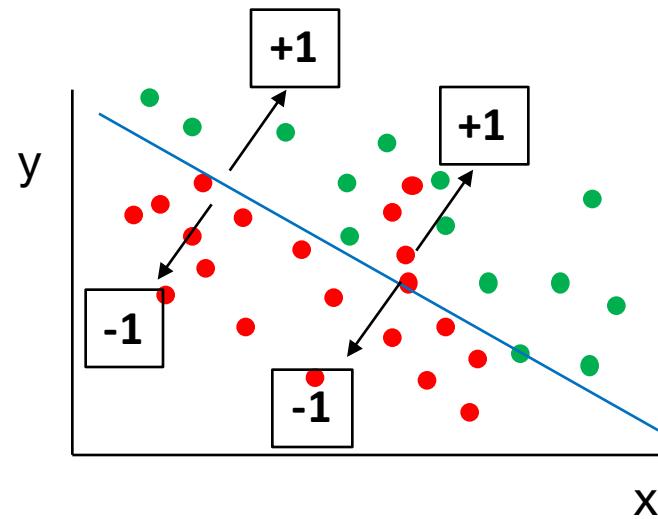
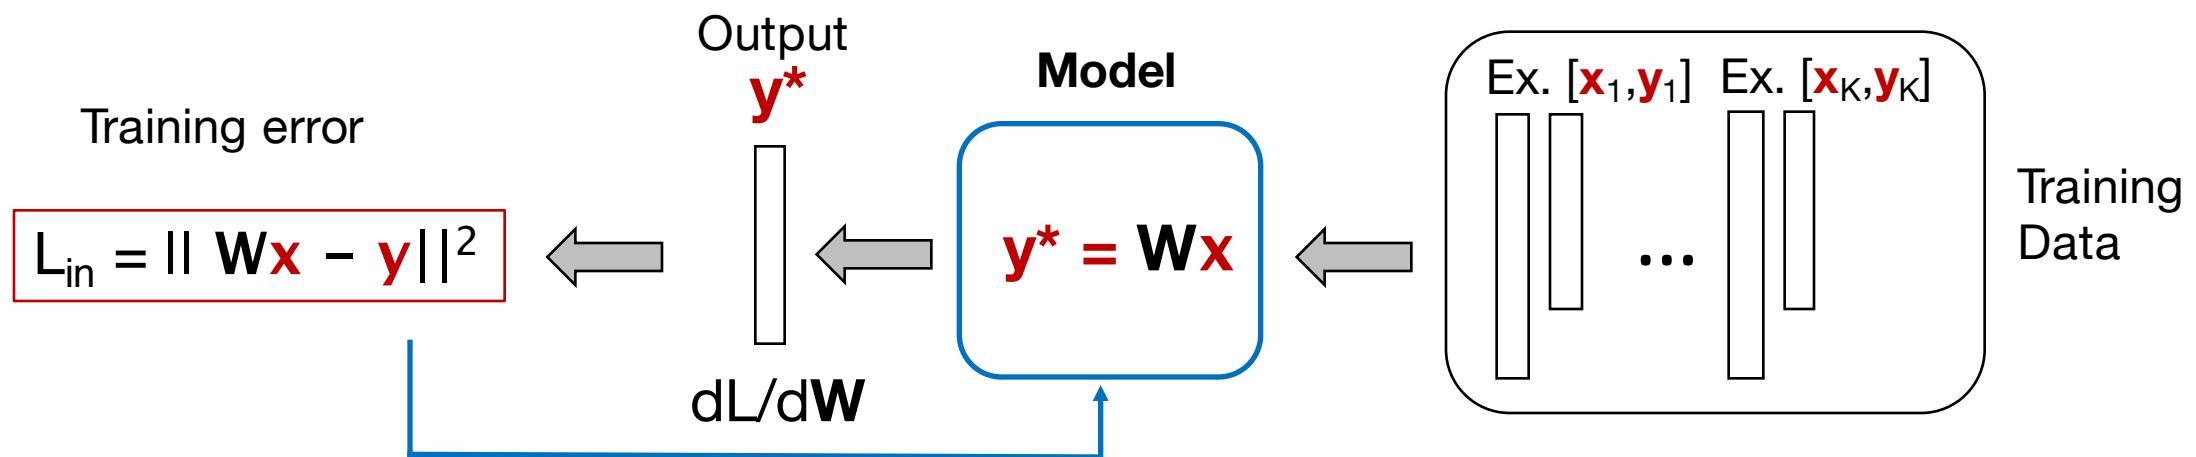
Don't mind about  
annoyance...



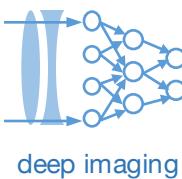
Letting an intruder in



# Last time: the linear classification model – what's not to like?



1. Can only separate data with lines (hyper-planes)...
2. We only allowed for binary labels ( $y = +/- 1$ )
3. Error function  $L_{in}$  inherently makes assumptions about statistical distribution of data



## Deriving cost function for logistic classification for probabilistic outputs

$$\text{Maximize } P(y_1, y_2, \dots, y_N | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N) = \prod_{n=1}^N \theta(y_n \mathbf{w}^T \mathbf{x}_n)$$

$$\text{Minimize } -\frac{1}{N} \ln \left( \prod_{n=1}^N \theta(y_n \mathbf{w}^T \mathbf{x}) \right)$$

$$\text{Minimize } \frac{1}{N} \sum_{n=1}^N \ln \left( \frac{1}{\theta(y_n \mathbf{w}^T \mathbf{x})} \right)$$

Use relationship

$$\theta(a) = \frac{1}{1 + e^{-a}}$$

$$\text{Minimize } L_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \mathbf{w}^T \mathbf{x}})$$

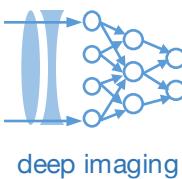
$$L_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (y_n - \mathbf{w}^T \mathbf{x})^2$$

Cross entropy error for logistic classification

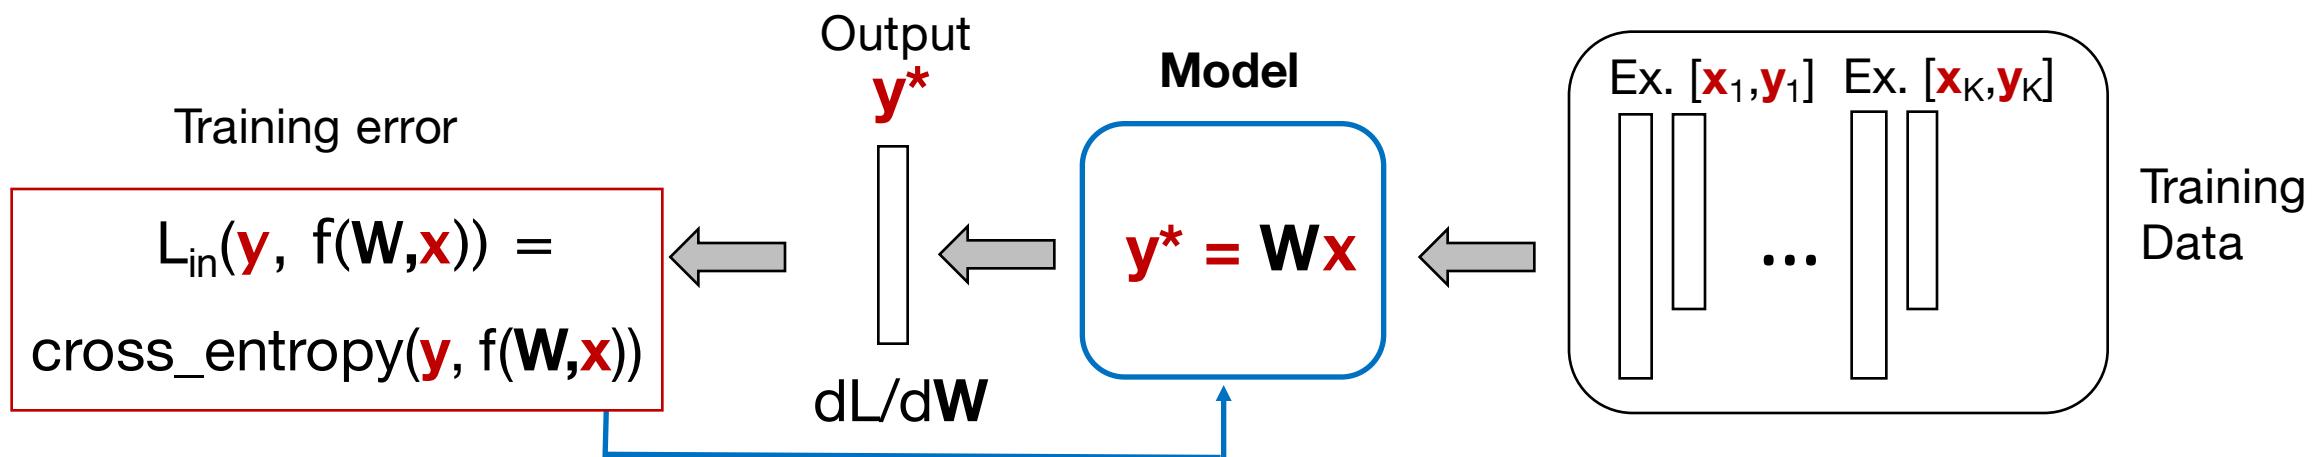
Requires iterative solution to minimize

Mean-square error for linear classification

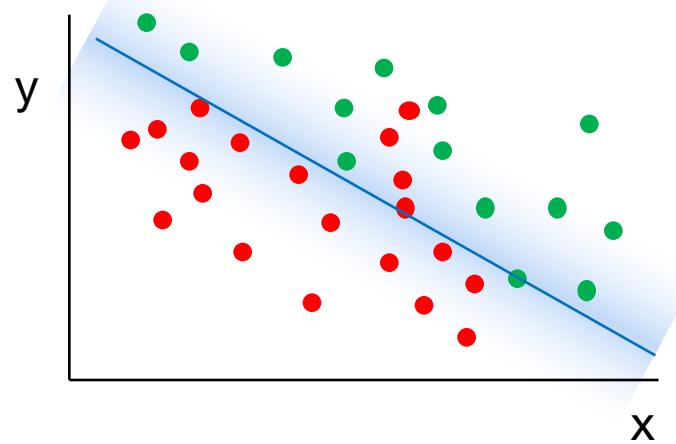
Closed form solution available

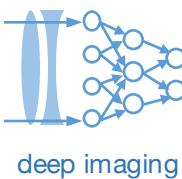


# The linear classification model – what's not to like?

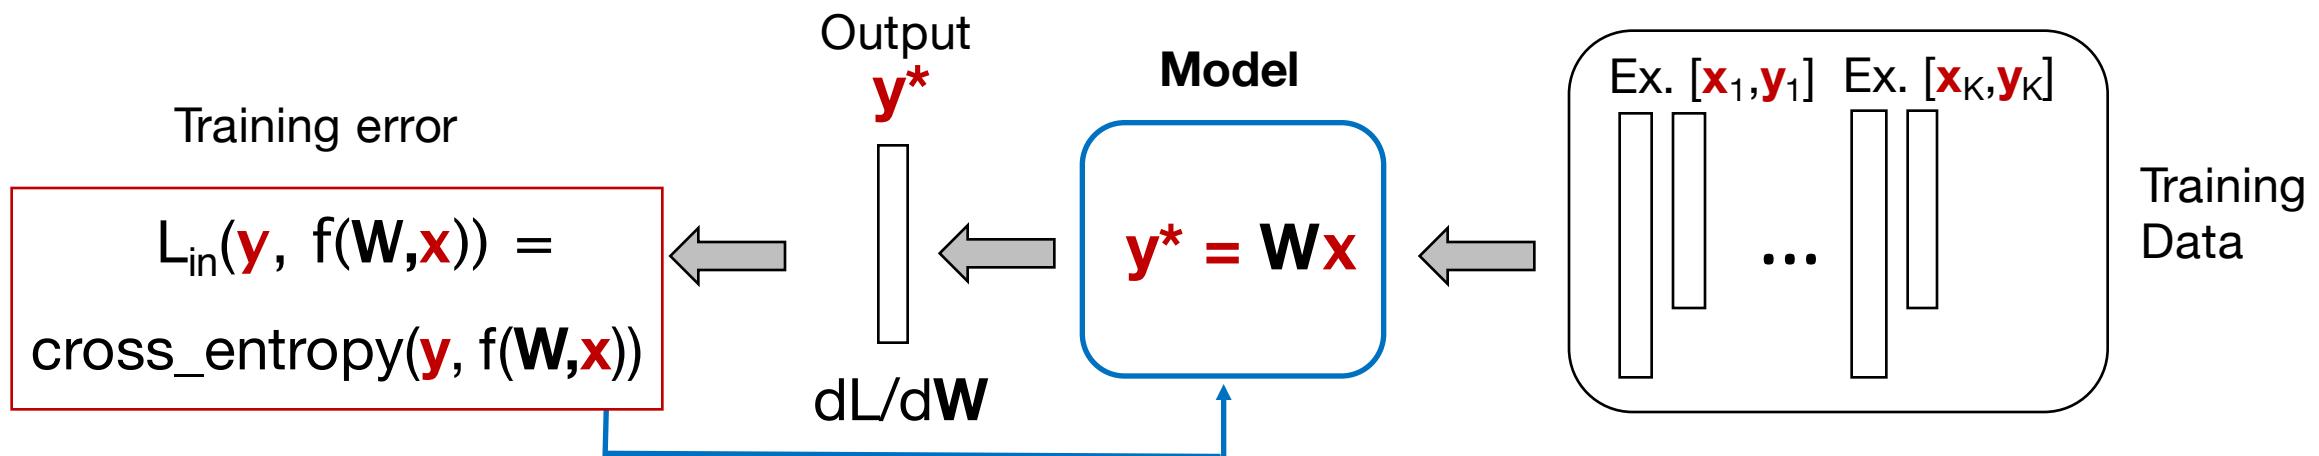


Probabilistic mapping to  $y$

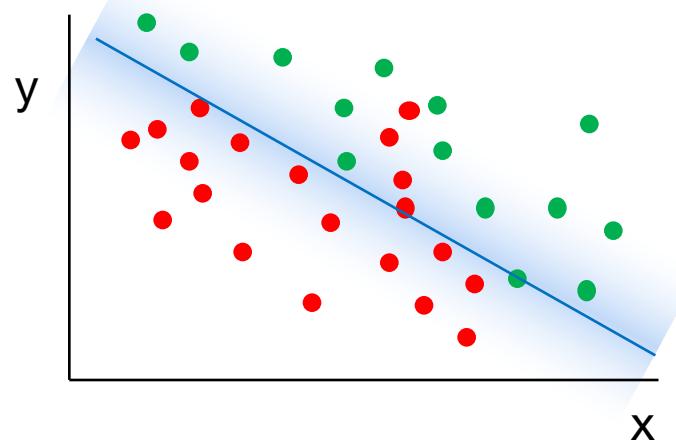




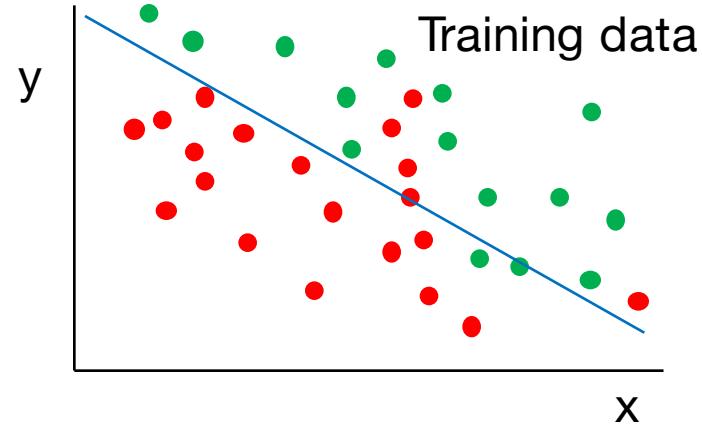
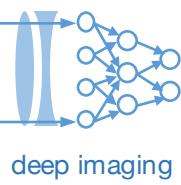
# The linear classification model – what's not to like?



Probabilistic mapping to  $y$



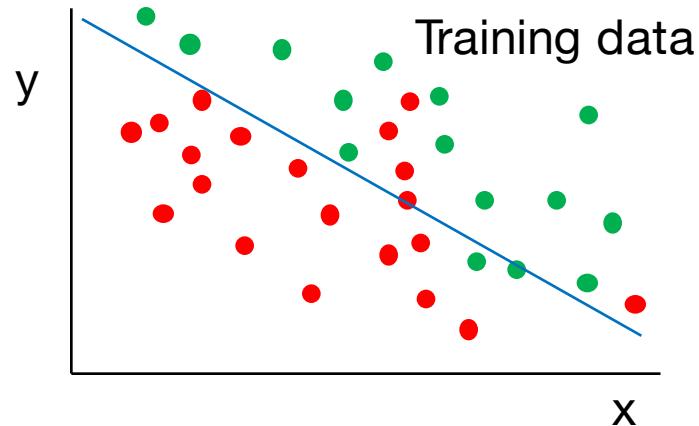
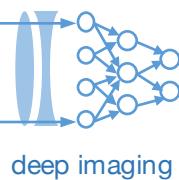
1. Can only separate data with lines (hyper-planes)...
2. We only allowed for binary labels ( $y = +/- 1$ )
3. Error function  $L_{in}$  inherently makes assumptions about statistical distribution of data



$$f = W_1x$$

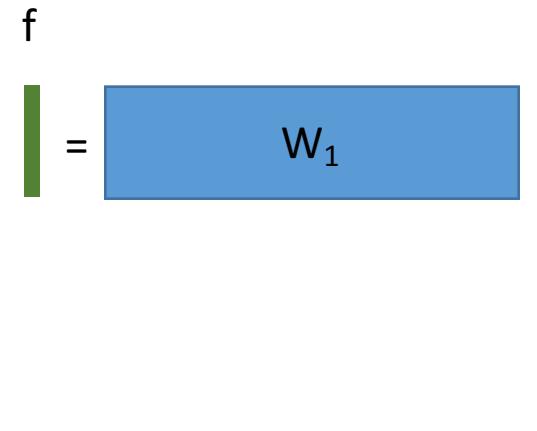
Learned  $f$ : not flexible

$$f = W_1x$$



$$f = W_1 x$$

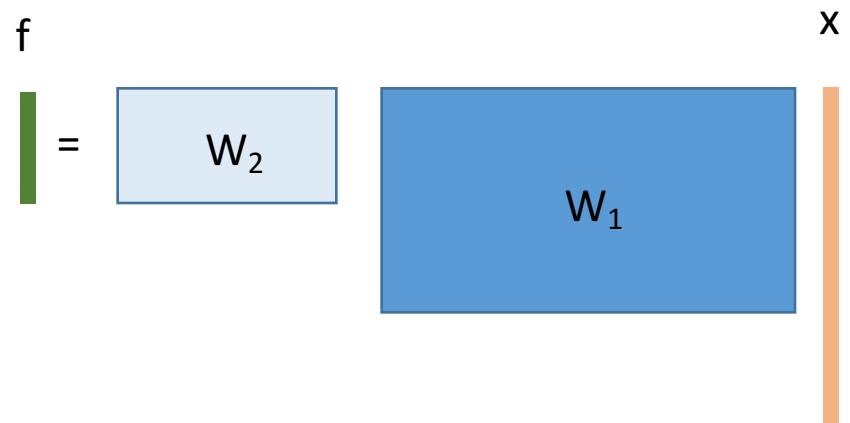
Learned  $f$ : not flexible



Can we add flexibility by multiplying with another weight matrix?

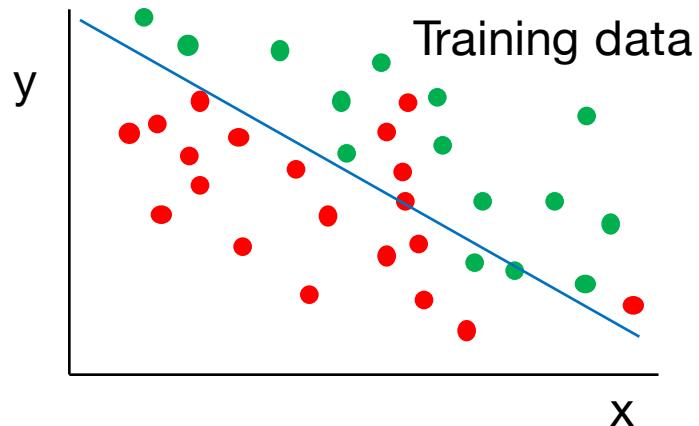
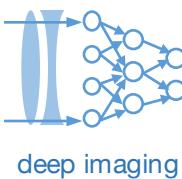
$$\begin{cases} f_1 = W_1 x + b_1 \\ f_2 = W_2 f_1 + b_2 \end{cases}$$

$$f_2 = W_2(W_1 x + b_1) + b_2$$



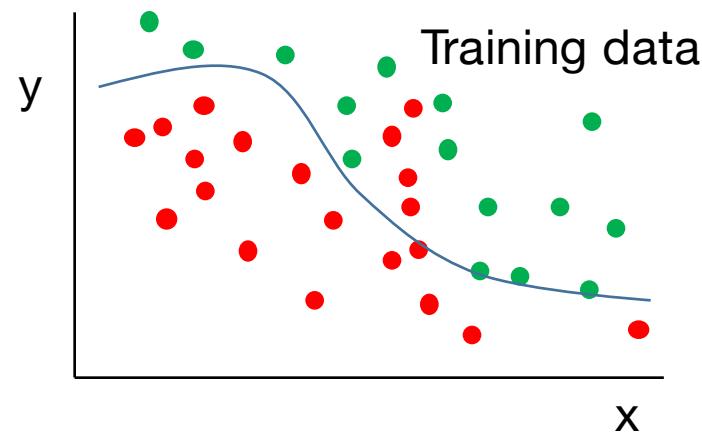
$$f_2 = W' x + b'$$

Unfortunately not...



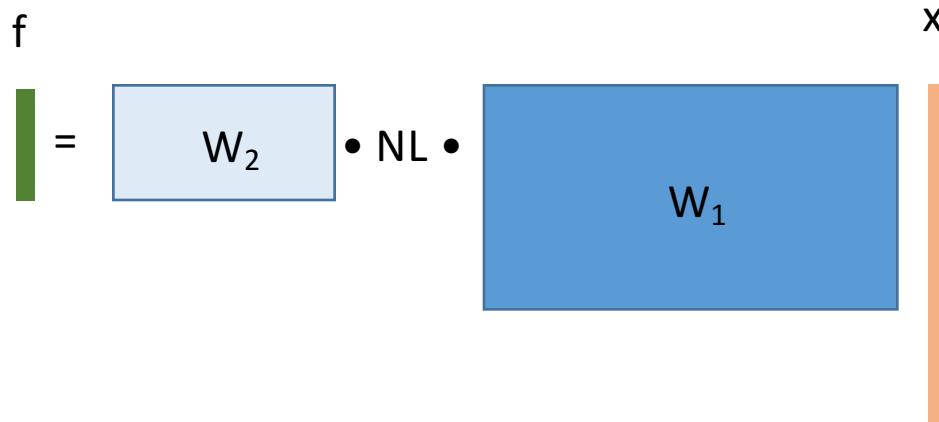
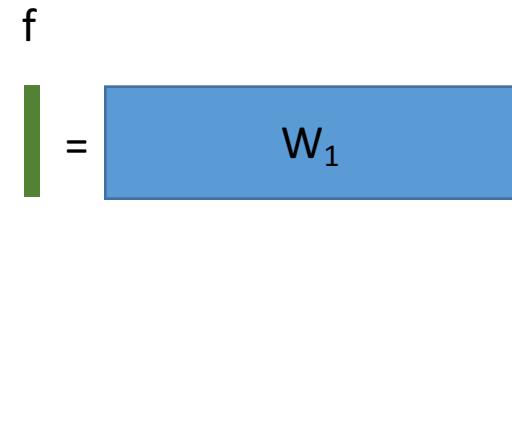
$$f = W_1 x$$

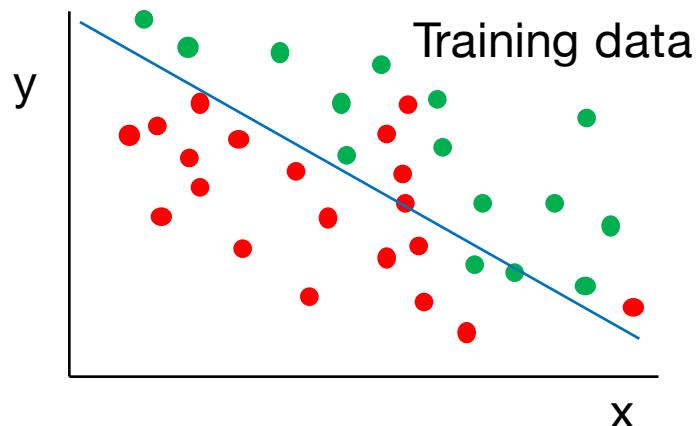
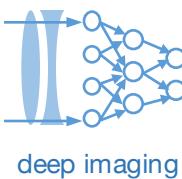
Learned  $f$ : not flexible



$$f = W_2 \max(W_1 x, 0)$$

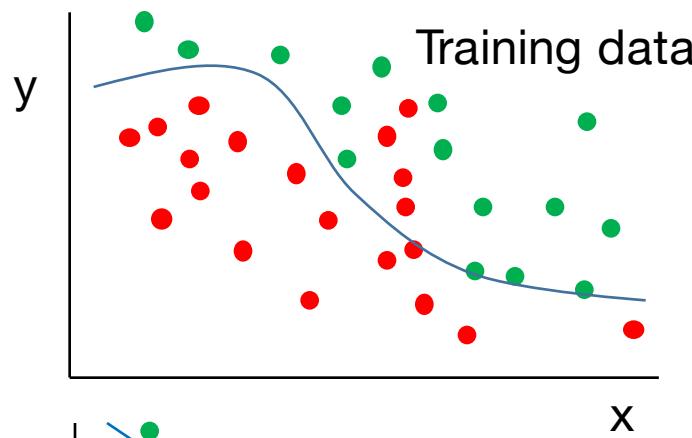
Learned  $f$ : a bit flexible





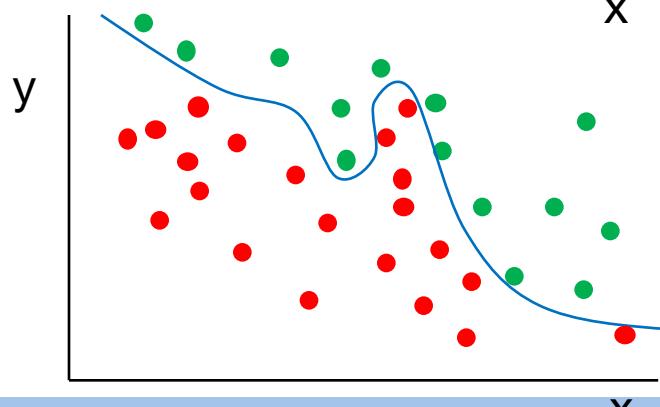
$$f = W_1 x$$

Learned  $f$ : not flexible



$$f = W_2 \max(W_1 x, 0)$$

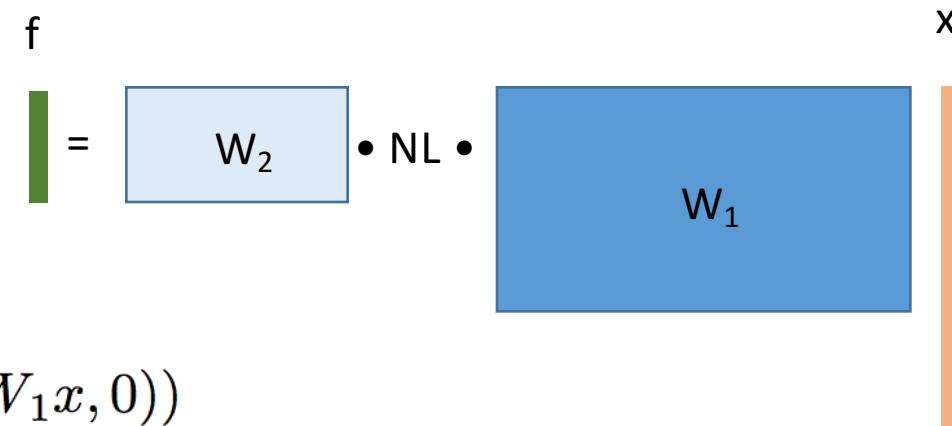
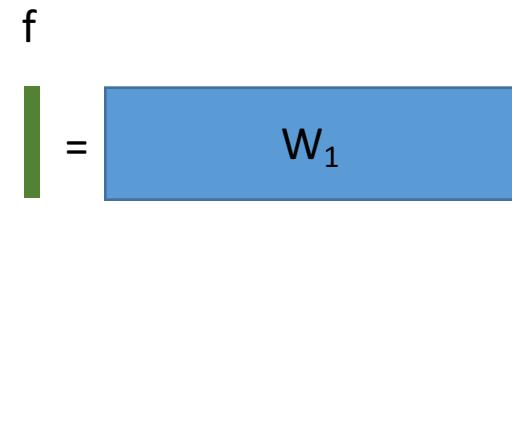
Learned  $f$ : a bit flexible



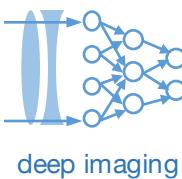
$$f = W_3 \max(0, W_2 \max(W_1 x, 0))$$

Learned  $f$ : more flexible

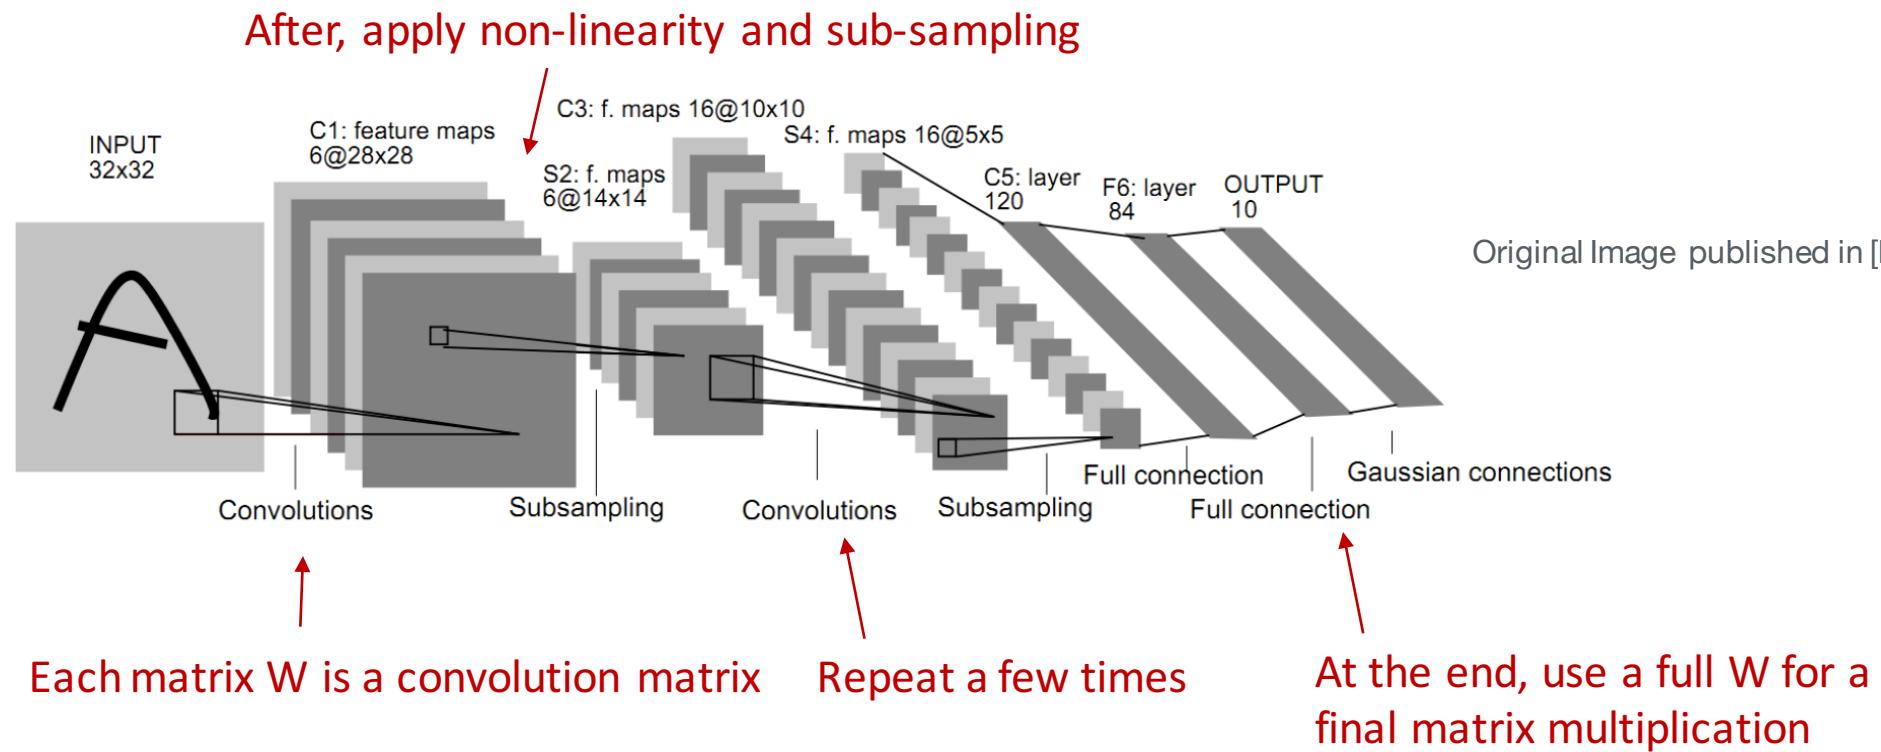
Does it generalize???

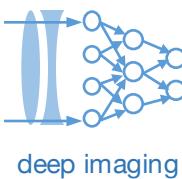


We can keep adding  
these “layers”...

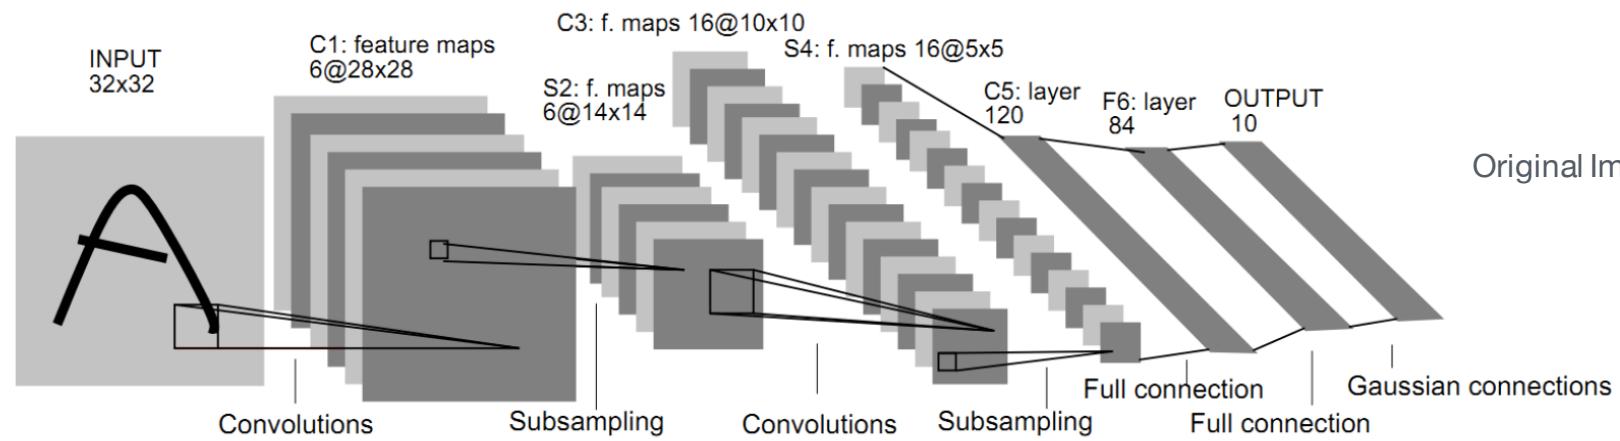


# Getting us to Convolutional Neural Networks



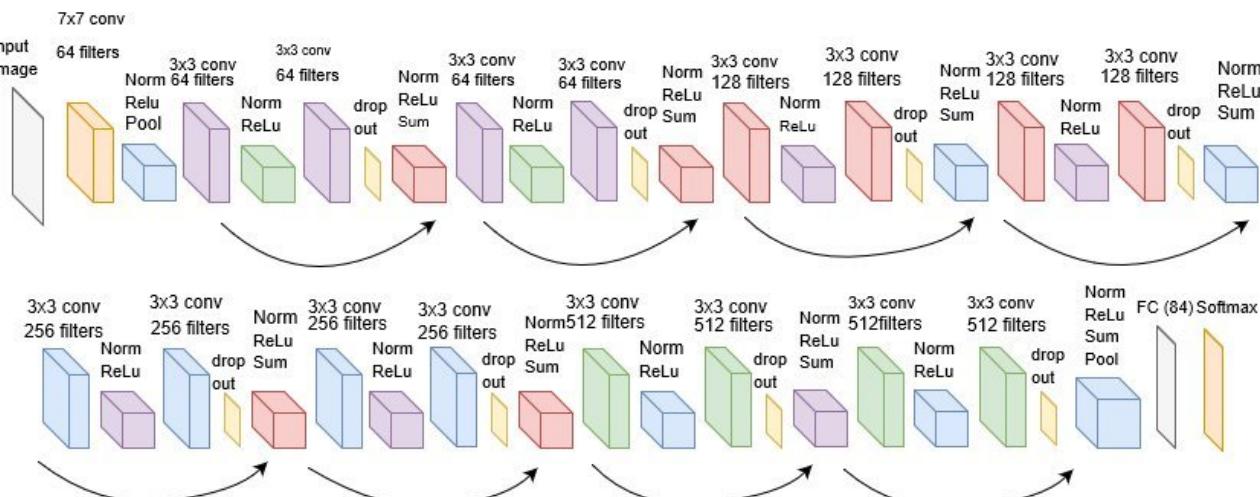


# Getting us to Convolutional Neural Networks



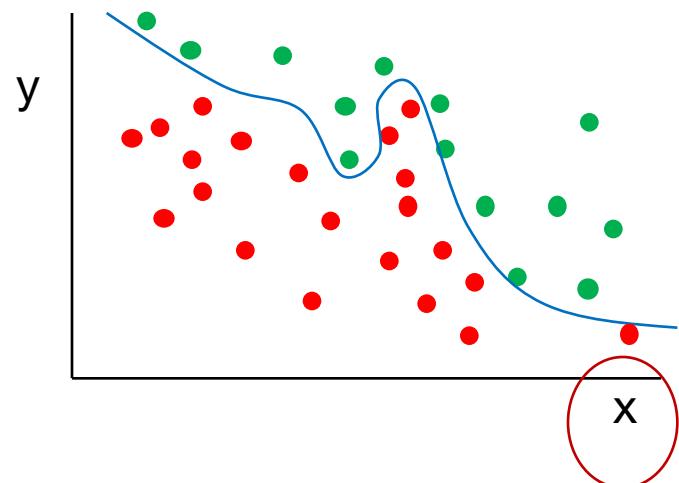
Original Image published in [LeCun et al., 1998]

In practice, this process is repeated many times:



## Aside #1 before convolutional neural network details

Q: Can we try to avoid making these learning models too complicated?

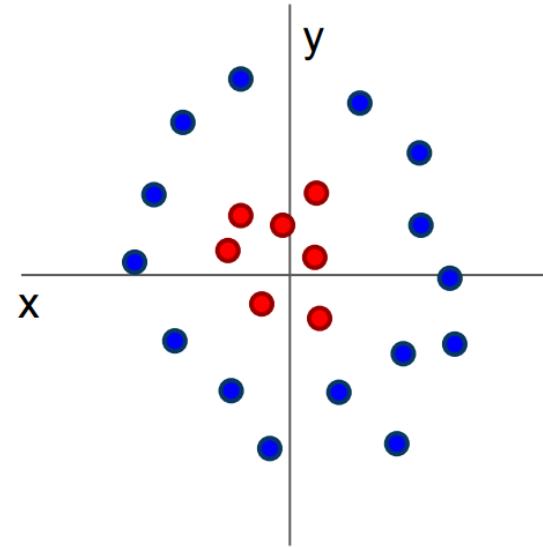


Learned  $f$ : more flexible

Does it generalize???

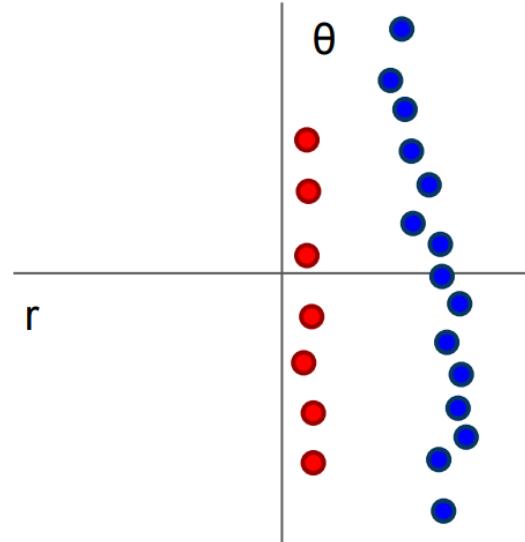
A: Yes, by transforming the data coordinates *before* classification

# Image Features: Motivation



Cannot separate red  
and blue points with  
linear classifier

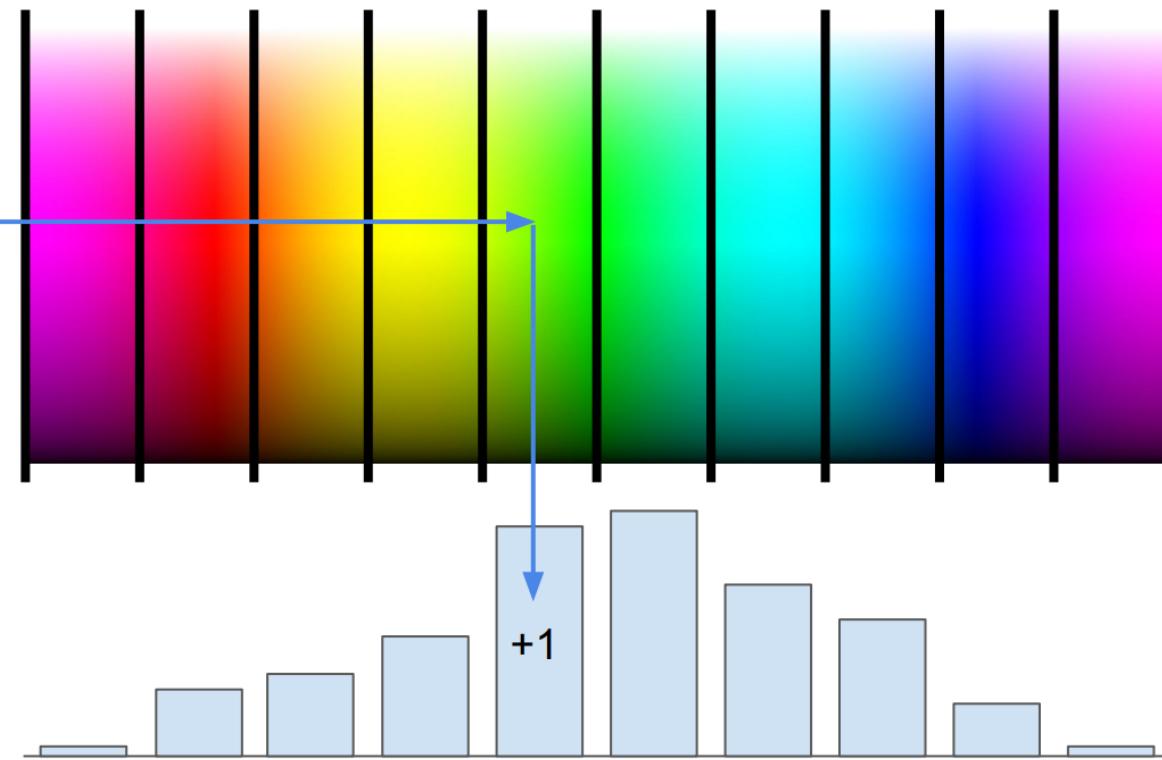
$$f(x, y) = (r(x, y), \theta(x, y))$$



After applying feature  
transform, points can  
be separated by linear  
classifier

From Stanford CS231: <http://cs231n.stanford.edu/>

# Example: Color Histogram



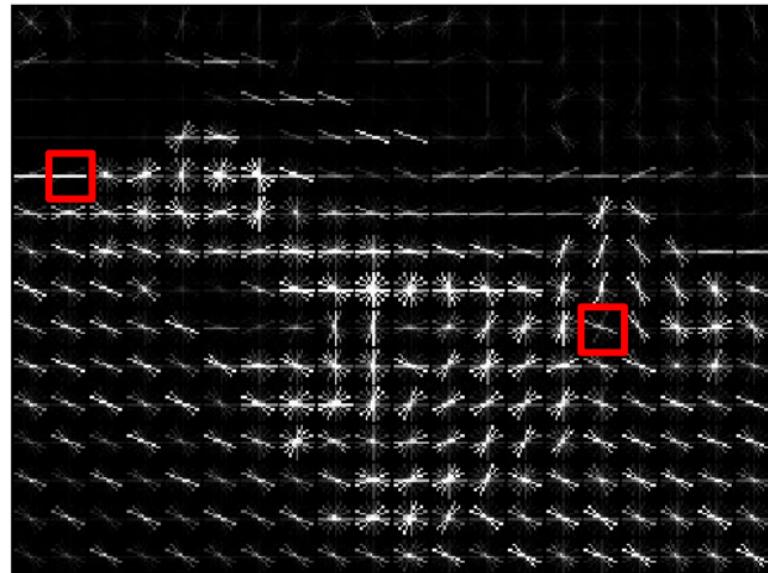
From Stanford CS231: <http://cs231n.stanford.edu/>

# Example: Histogram of Oriented Gradients (HoG)



Divide image into 8x8 pixel regions  
Within each region quantize edge  
direction into 9 bins

Lowe, "Object recognition from local scale-invariant features", ICCV 1999  
Dalal and Triggs, "Histograms of oriented gradients for human detection," CVPR 2005



Example: 320x240 image gets divided  
into 40x30 bins; in each bin there are  
9 numbers so feature vector has  
 $30 \times 40 \times 9 = 10,800$  numbers

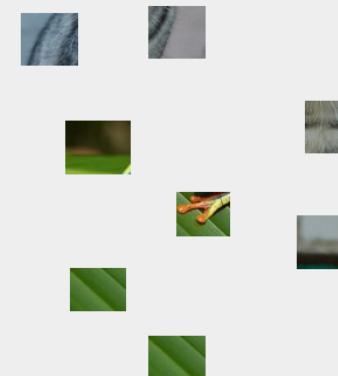
From Stanford CS231: <http://cs231n.stanford.edu/>

# Example: Bag of Words

## Step 1: Build codebook



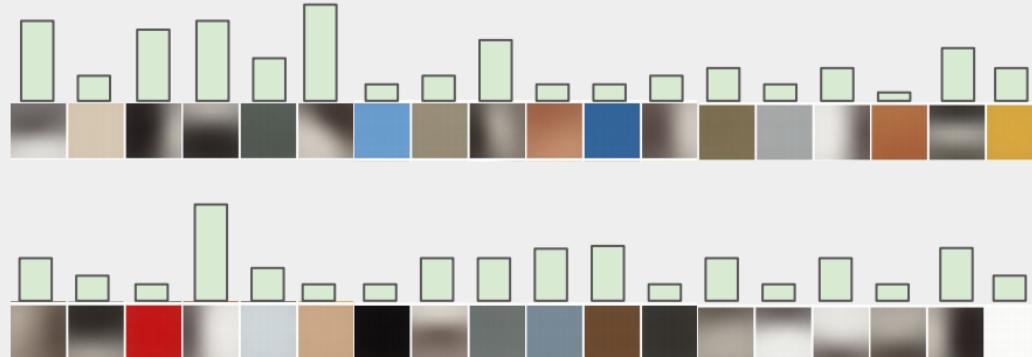
Extract random patches



Cluster patches to  
form “codebook”  
of “visual words”

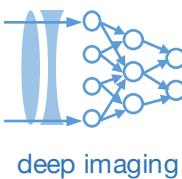


## Step 2: Encode images

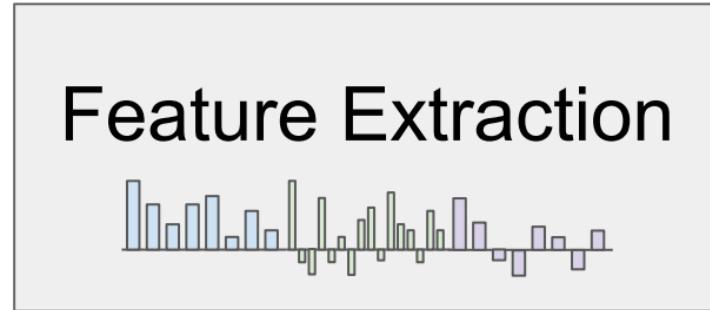


Fei-Fei and Perona, “A bayesian hierarchical model for learning natural scene categories”, CVPR 2005

From Stanford CS231: <http://cs231n.stanford.edu/>



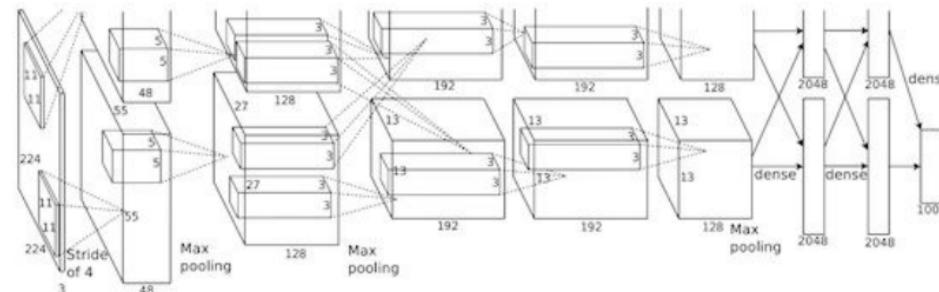
# Image features vs ConvNets



$f$

training

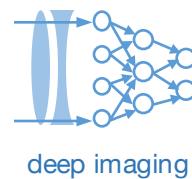
10 numbers giving  
scores for classes



training

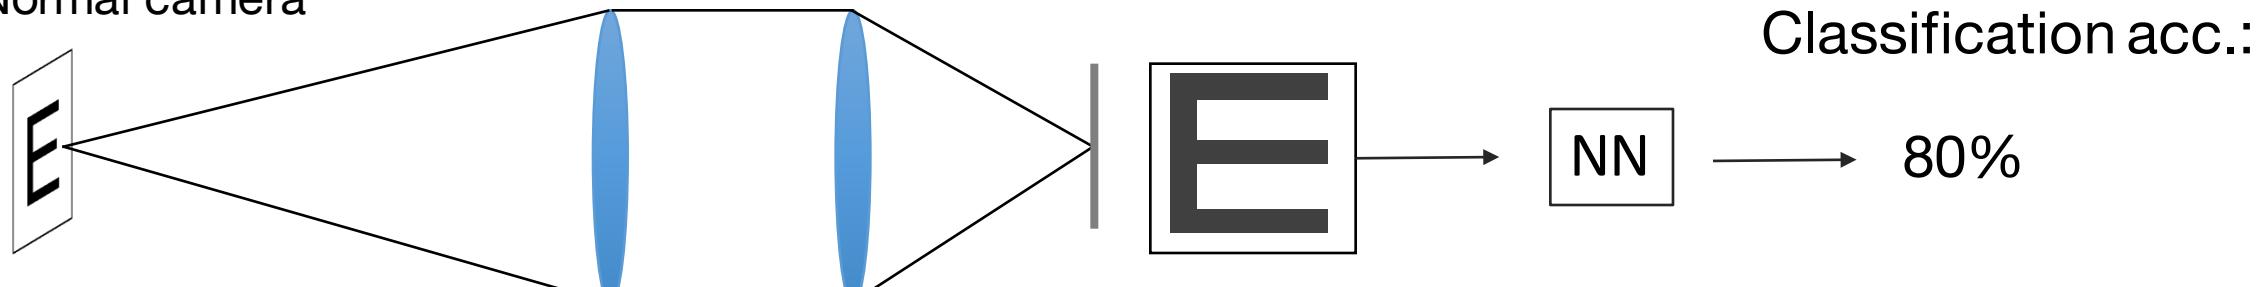
10 numbers giving  
scores for classes

From Stanford CS231: <http://cs231n.stanford.edu/>

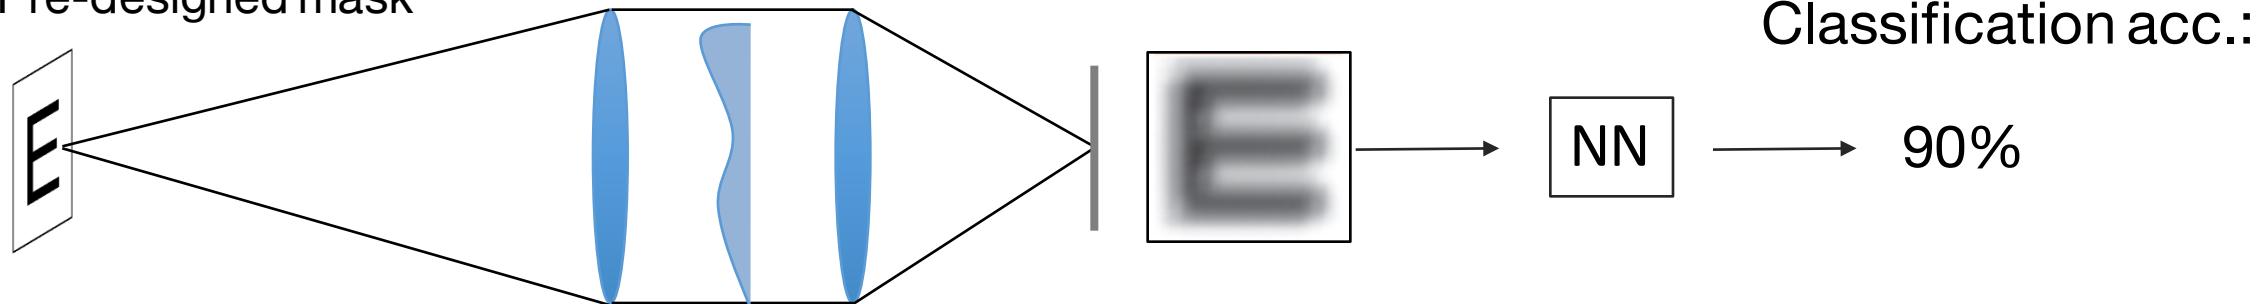


# Hand-crafted versus learned features also applies to imaging

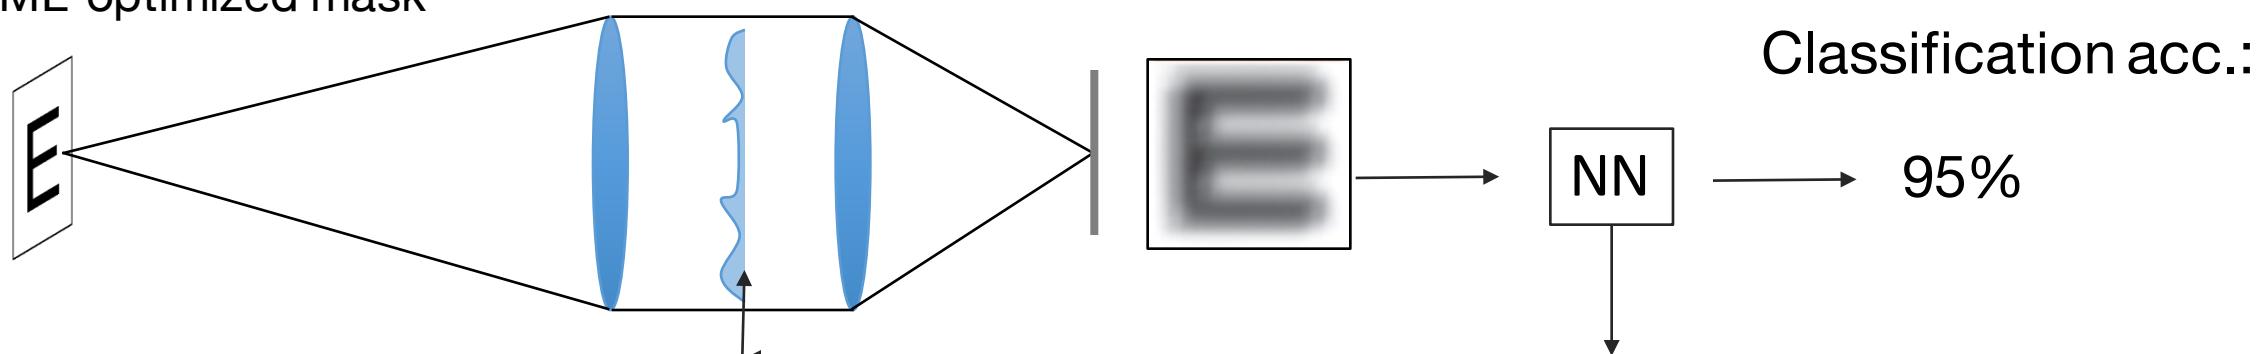
Normal camera

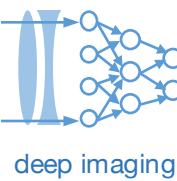


Pre-designed mask



ML-optimized mask

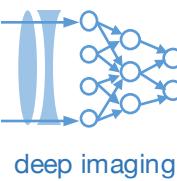




# Statistical Machine Learning in 30 minutes

Two competing goals in machine learning:

1. Can we make sure the in-sample error  $L_{in}(y, f(x, W))$  is small enough?
  - Appropriate cost function
  - “complex enough” model



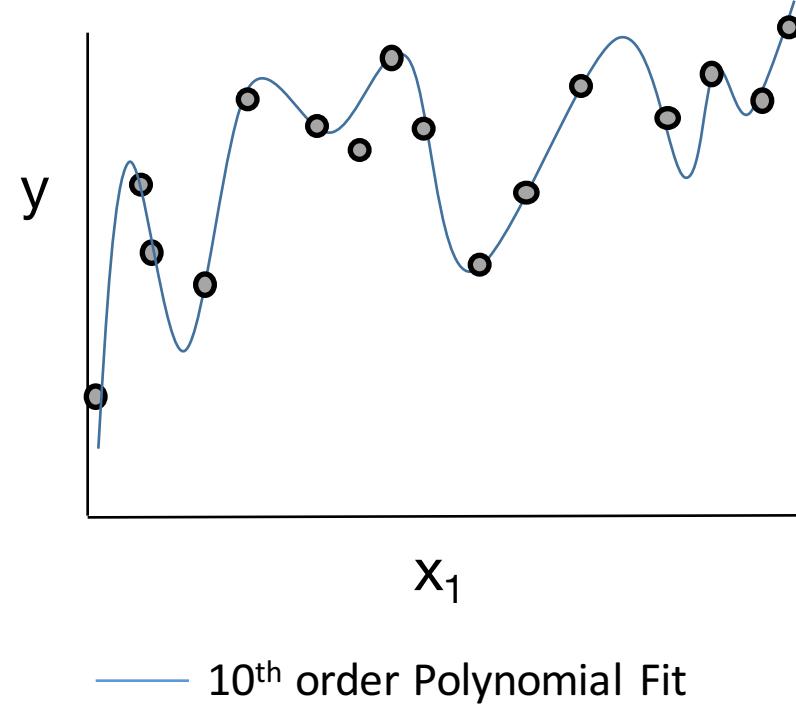
# Statistical Machine Learning in 30 minutes

Two competing goals in machine learning:

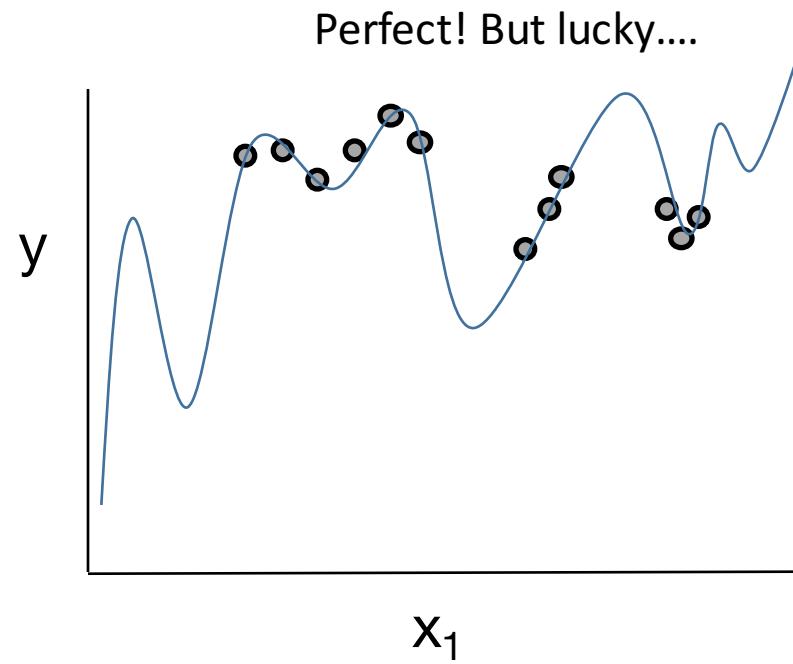
1. Can we make sure the in-sample error  $L_{in}(y, f(x, W))$  is small enough?
  - Appropriate cost function
  - “complex enough” model
2. Can we make sure that  $L_{out}(y, f(x, W))$  is close enough to  $L_{in}(y, f(x, W))$ ?
  - Probabilistic analysis says yes!
  - $|L_{in} - L_{out}|$  bounded from above
  - Bound grows with model capacity (bad)
  - Bound shrinks with # of training examples (good)

## Model overfitting versus underfitting – a thought exercise

Let's fit these “training” data points:

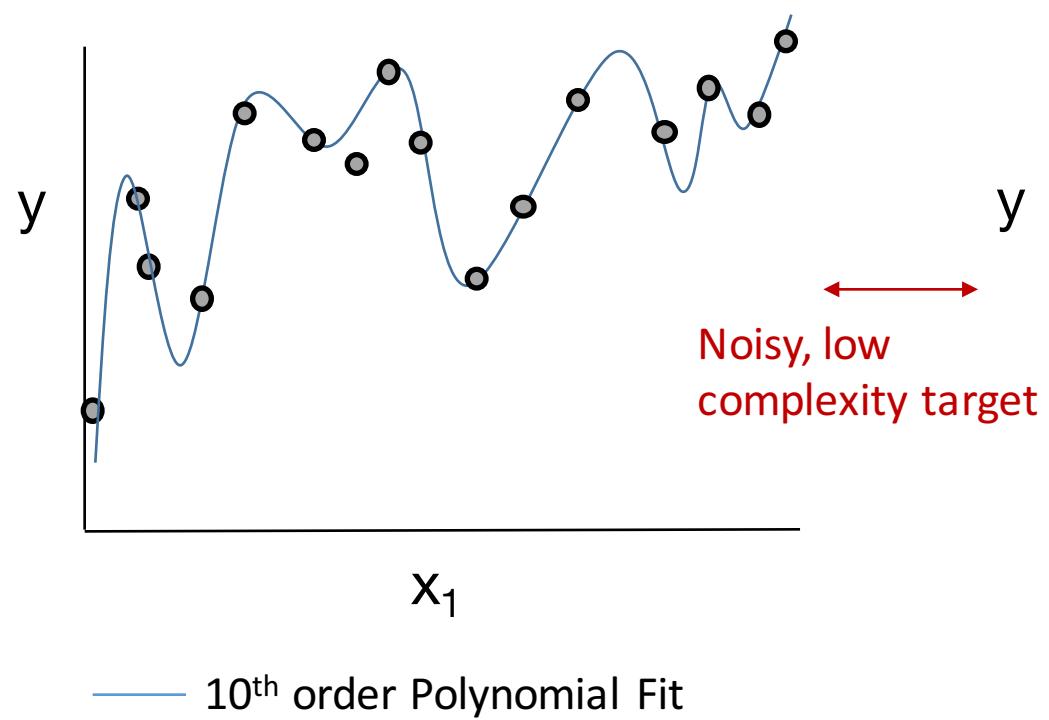


And then here's our testing dataset – good?

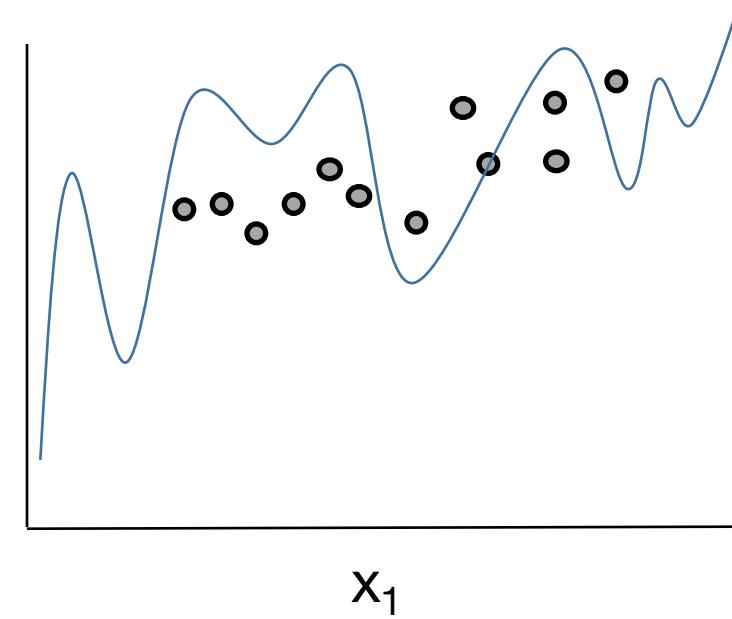


# Model overfitting versus underfitting – a thought exercise

Let's fit these "training" data points:

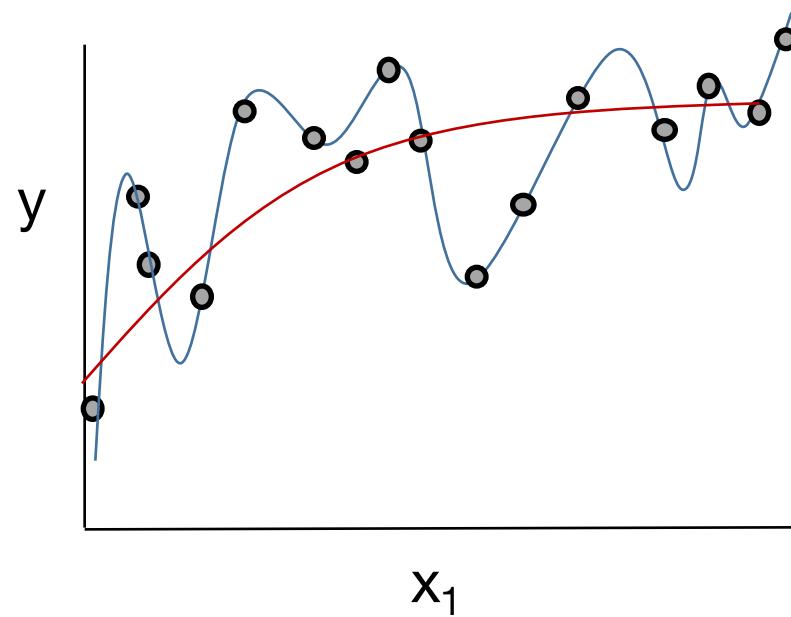


What if our test dataset was this :



# Model overfitting versus underfitting – a thought exercise

Let's fit these "training" data points:

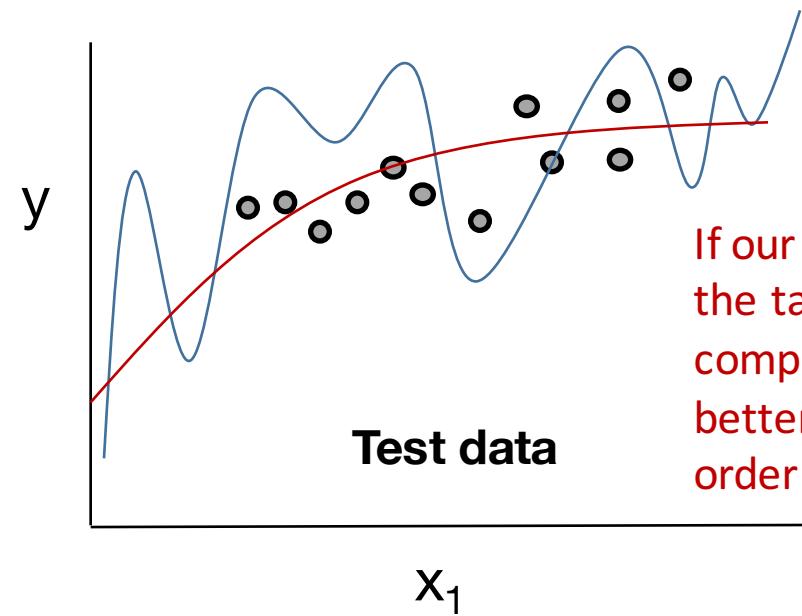
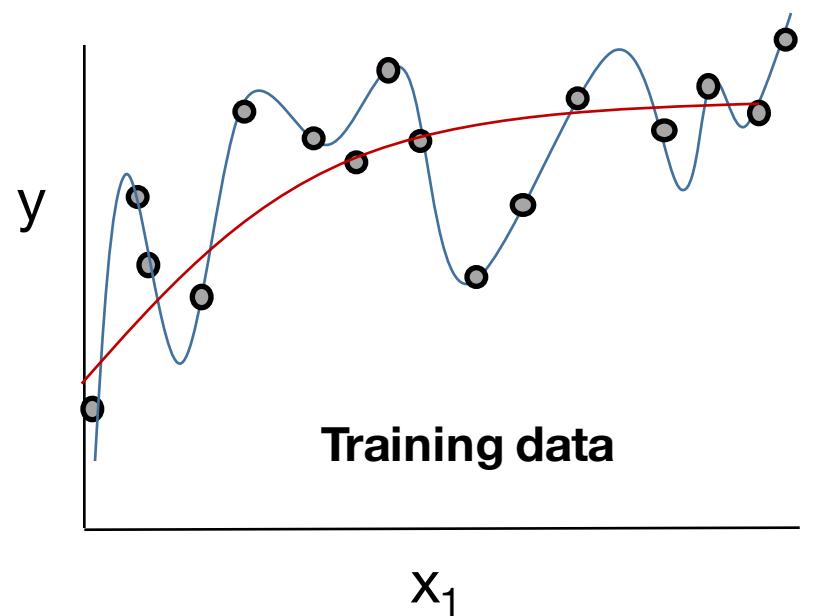


- 10<sup>th</sup> order Polynomial Fit
- 2<sup>nd</sup> order Polynomial Fit

What if our test dataset was this :



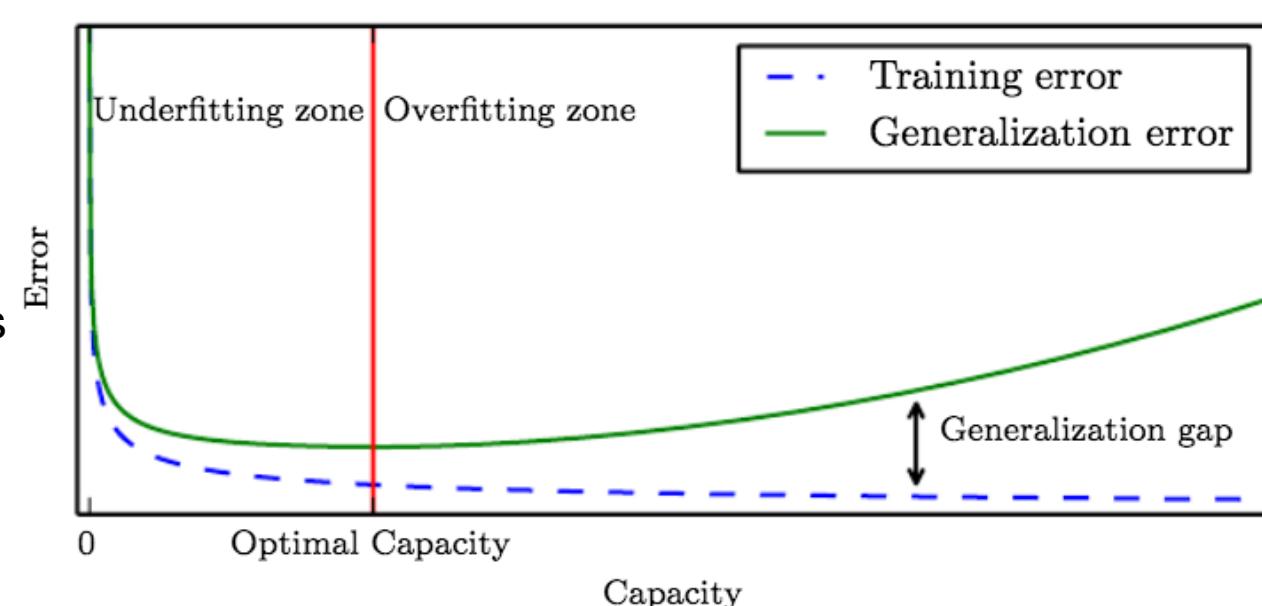
## Model overfitting versus underfitting – a thought exercise



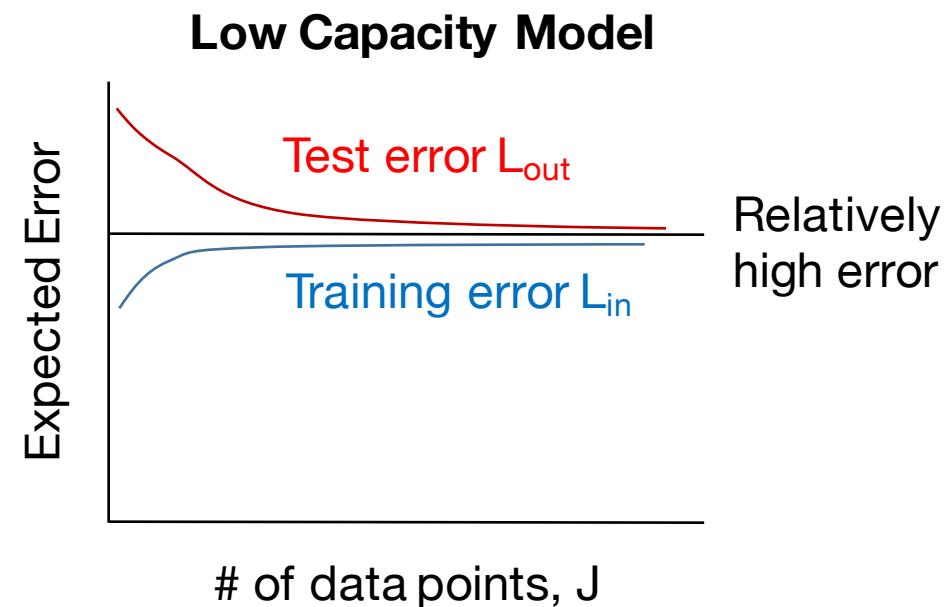
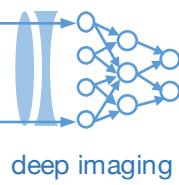
**Model capacity:** ability to fit a wide range of functions

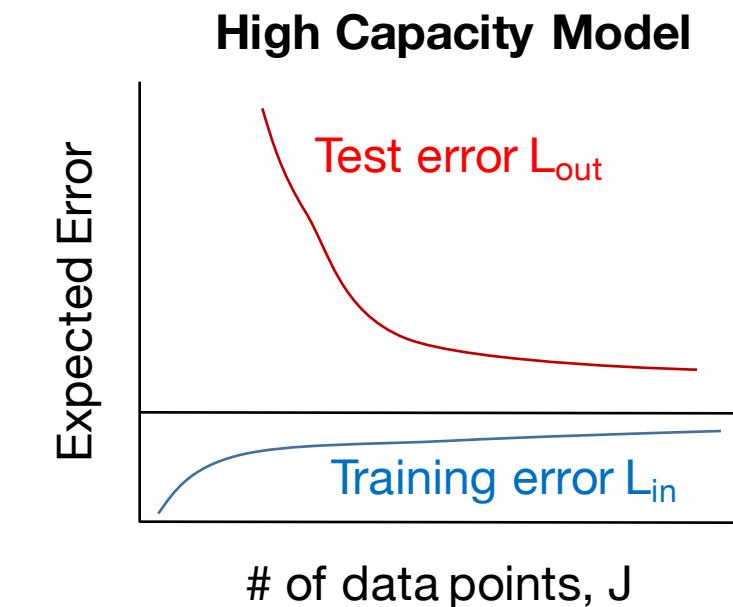
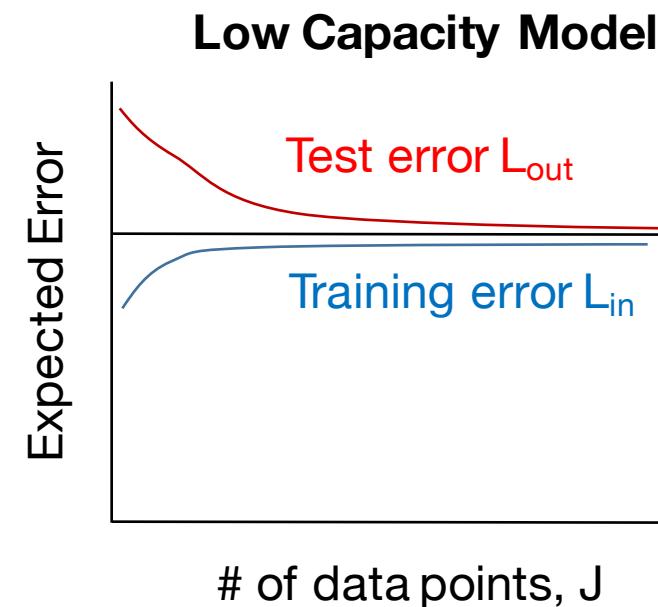
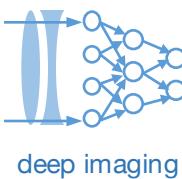
Control capacity through model's hypothesis space (set of functions model can take)

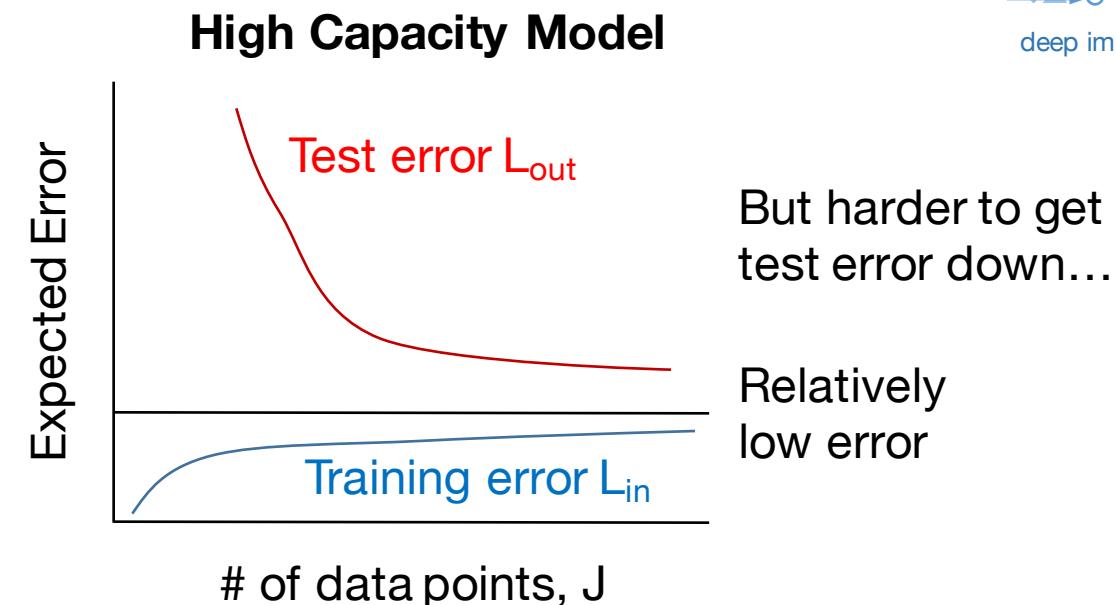
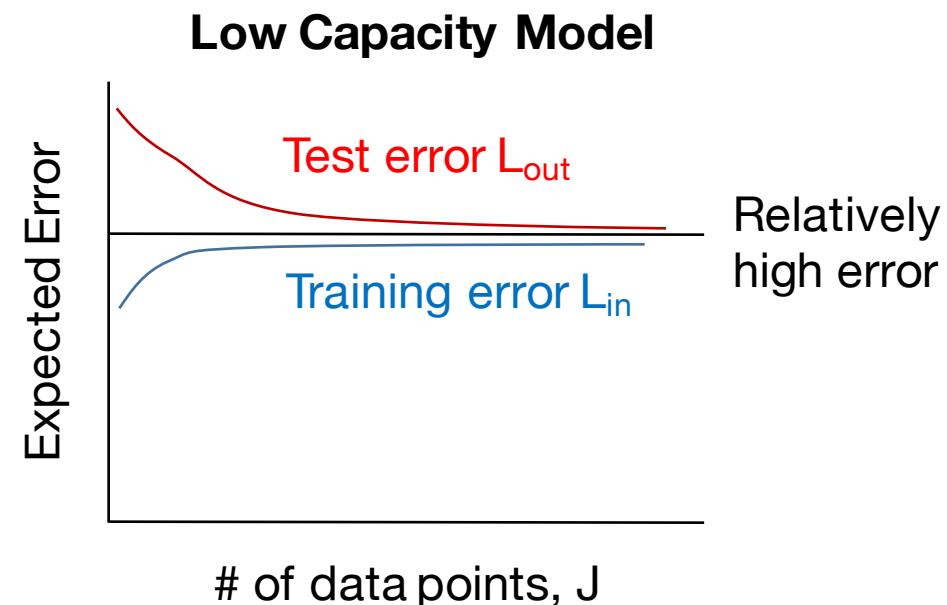
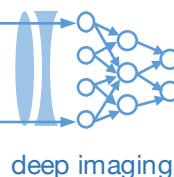
Hard to know ahead of time!



*Deep Learning, I. Goodfellow et al., Fig. 5.3*

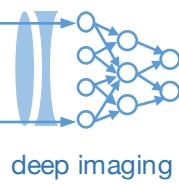




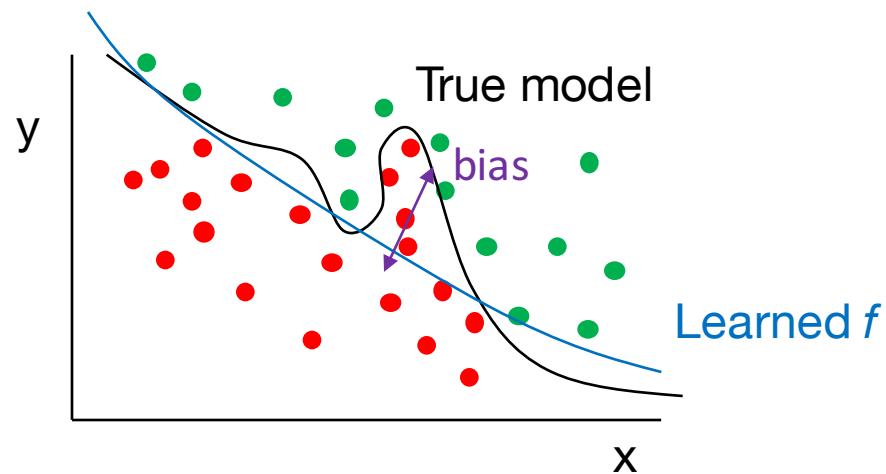
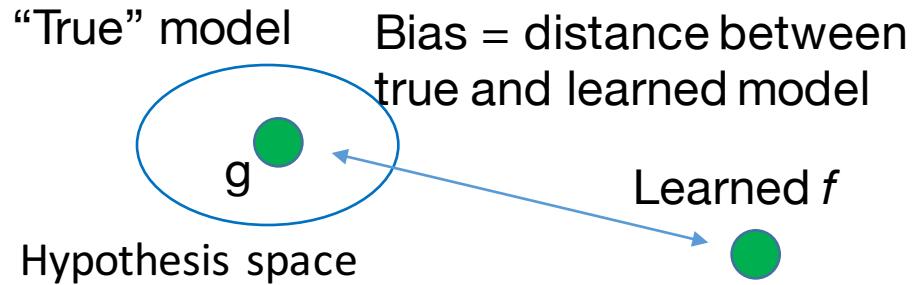


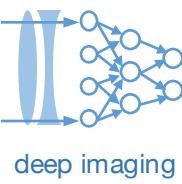
### Take away concepts:

- Can't ever really expect test error to be less than training error
- Complicated models tend to appear to “do better” during training, before trying test data
- When the model gets complicated and you don’t have enough data, challenging to get test error down

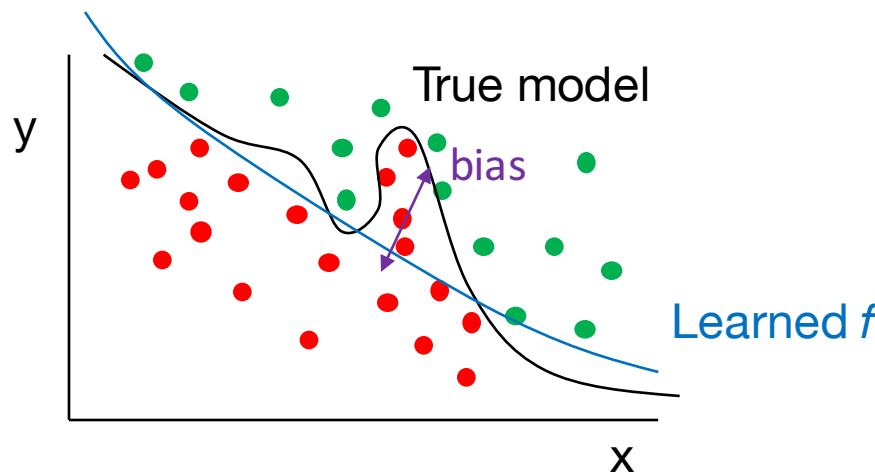
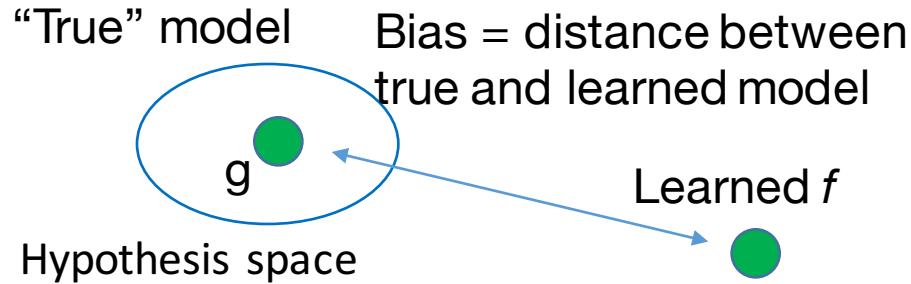


# Model bias versus variance





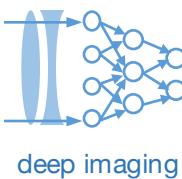
# Model bias versus variance



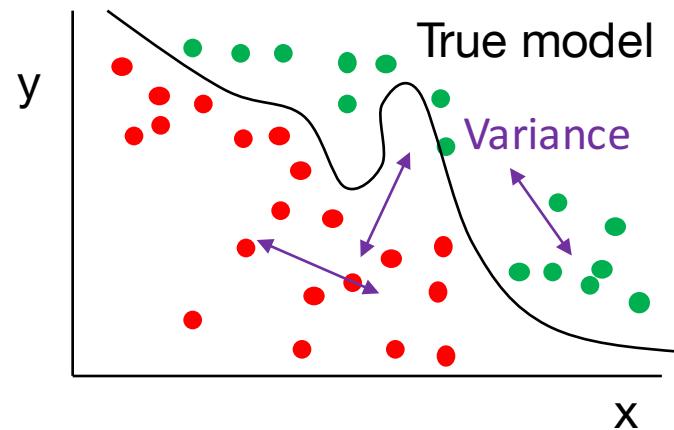
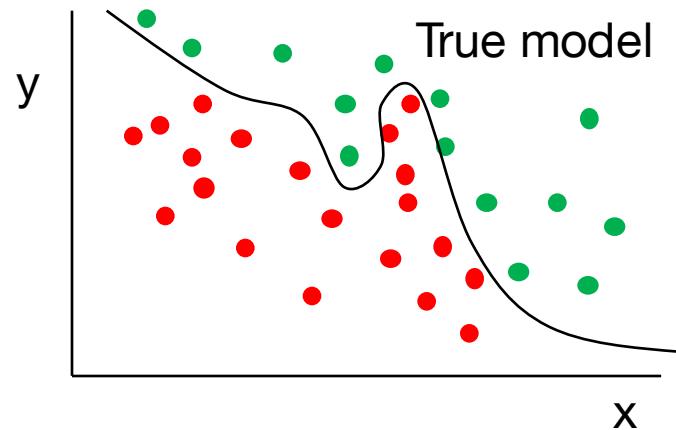
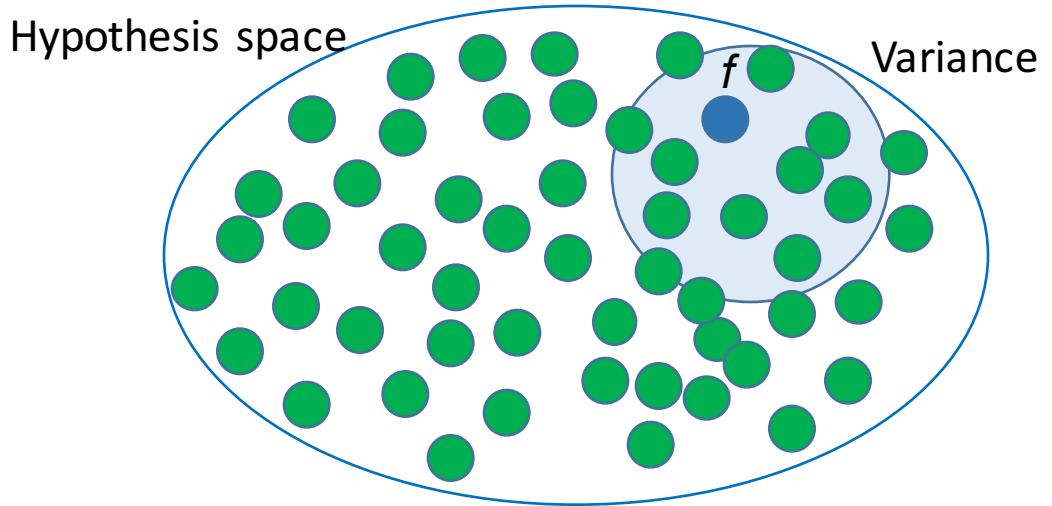
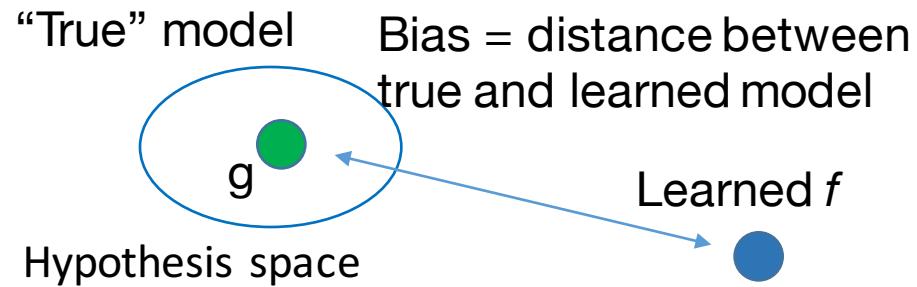
$$\text{Bias} = (g(\mathbf{x}) - f(\mathbf{x}))^2$$

Measures how far our learning model  $f$  is biased away from target function  $g$  (for perfect training data classification)

Models that tend to be “a bit too simple” are biased away from “true” model

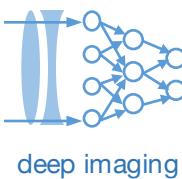


# Model bias versus variance

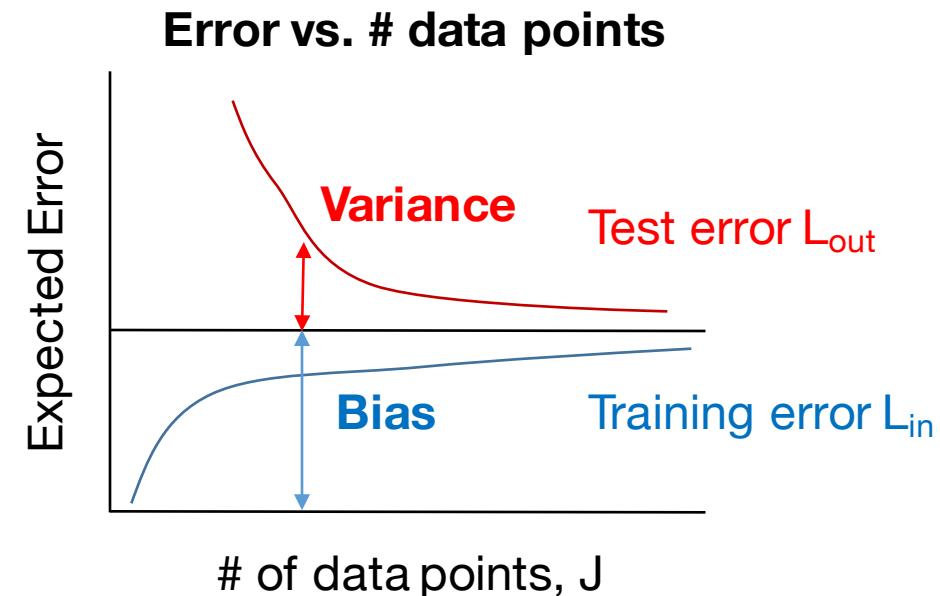
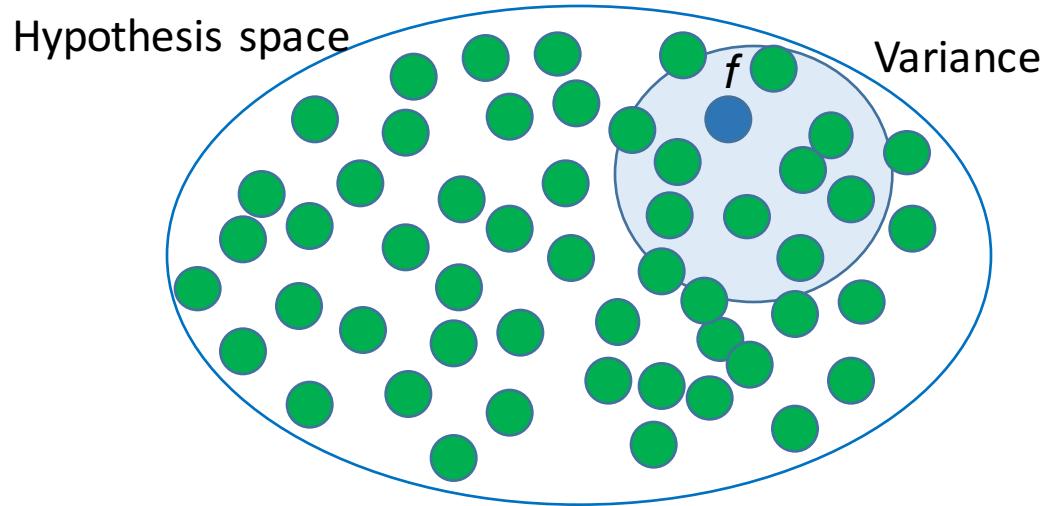
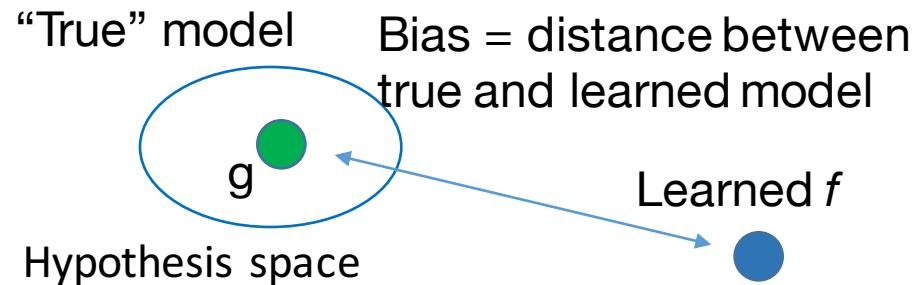


$$\text{Variance} = \text{Var}[g(x)]$$

More complicated datasets exhibit lots of variance between training and test set

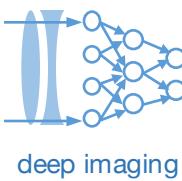


# Model bias versus variance



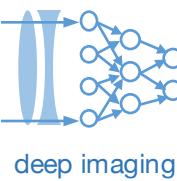
Test Error is sum of model bias and variance!

Goal is to find a model  $f$  that balances between these two quantities for a given dataset



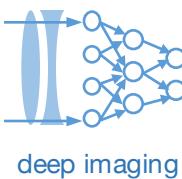
## How to formally define capacity and complexity?

- Short answer: it's complicated...
- Related to something called the *VC Dimension*
  - Can provide theoretical bounds on performance
  - Dimensional bounds rather than scalar bounds...
- I decided not to go into it, but please let me know if you'd like me to!



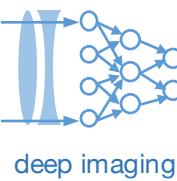
# Conclusions from statistical machine learning

- Conclusion: you want a model that is complex enough to capture variations within high-dimensional space, but not too complex such that it overfits the data
- Want a model with a high capacity, but can still *generalize* to data outside training set
  - More data -> less overfitting, complex target -> more overfitting
- For simple models, we can measure complexity via degrees of freedom, the VC bound and so-on to help us nail down ideal models that can generalize well



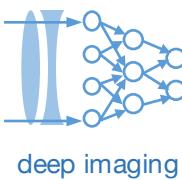
# Conclusions from statistical machine learning

- Conclusion: you want a model that is complex enough to capture variations within high-dimensional space, but not too complex such that it overfits the data
- Want a model with a high capacity, but can still *generalize* to data outside training set
  - More data -> less overfitting, complex target -> more overfitting
- For simple models, we can measure complexity via degrees of freedom, the VC bound and so-on to help us nail down ideal models that can generalize well
- **For DL models:** this will get too hard...here's a few counter-intuitive properties:
  1. A fixed DL *architecture* exhibits data-dependent complexities
    - e.g., “good” DL networks achieve 0 training error on images with random labels, so cannot generalize at all in this case, and are too complex
  2. DL networks with more hidden units leads to *better* generalization (the main finding of the last few years). So deeper models tend to be less complex, actually...
  3. Complexity depends upon loss function and optimization method...



## Important to remember: “No Free Lunch Theorem”

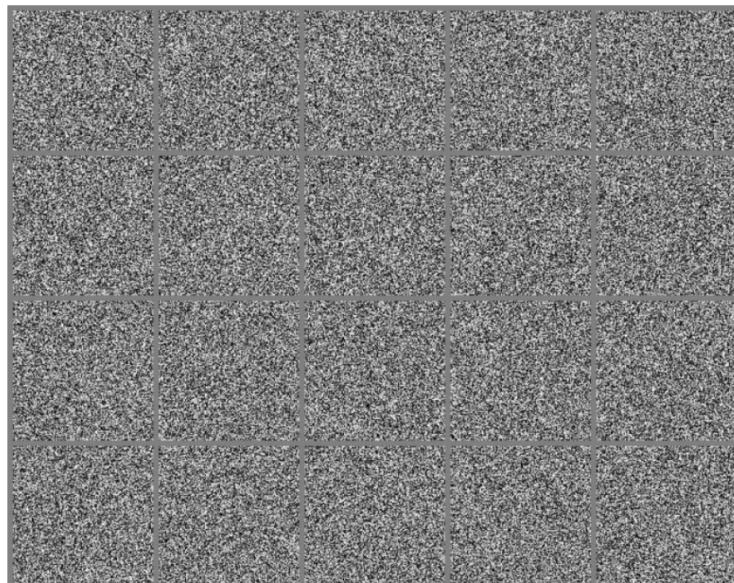
- “Averaged over all possible data-generating distributions, every classification algorithm has the same error rate when classifying previously unobserved points.”
- The most sophisticated DL algorithm has save average performance (averaged over all possible tasks) as the simplest.



## Important to remember: “No Free Lunch Theorem”

- “Averaged over all possible data-generating distributions, every classification algorithm has the same error rate when classifying previously unobserved points.”
- The most sophisticated DL algorithm has save average performance (averaged over all possible tasks) as the simplest.
- Must make assumptions about probability distributions of inputs we’ll encounter in real-world

Set of 20 “images”, random Gaussian distribution



Face at different orientations =  
manifold n-D space

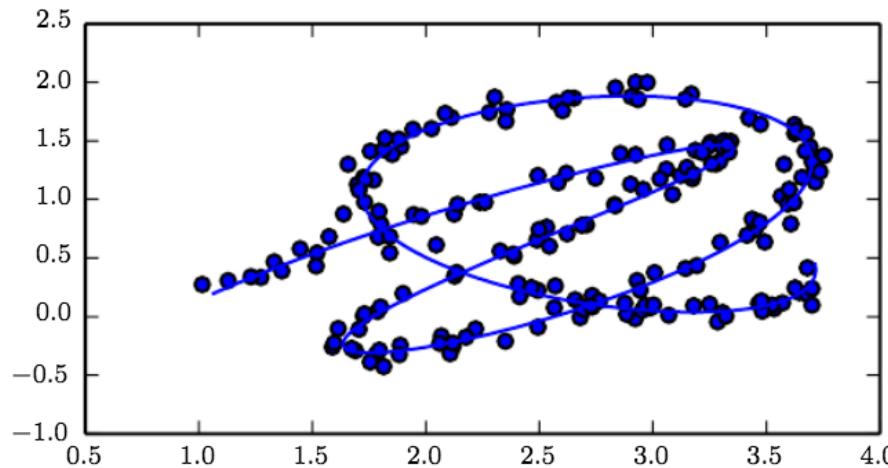


*Deep Learning*, I. Goodfellow et al., Fig. 5.12-13

## Important to remember: “No Free Lunch Theorem”

- “Averaged over all possible data-generating distributions, every classification algorithm has the same error rate when classifying previously unobserved points.”
- The most sophisticated DL algorithm has save average performance (averaged over all possible tasks) as the simplest.
- Must make assumptions about probability distributions of inputs we’ll encounter in real-world

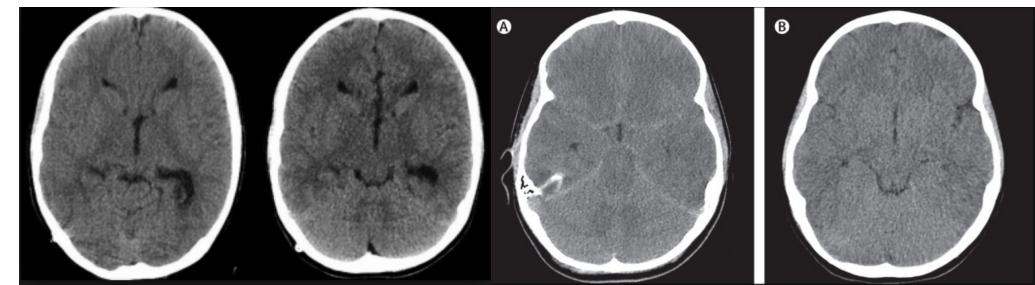
1D Manifold in 2D space



Manifold  
Hypothesis



CT reconstructions of every brain in the world = kD manifold in nD space?



Deep Learning, I. Goodfellow et al., Fig. 5.11