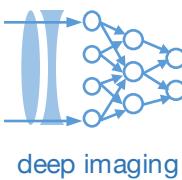


# Lecture 24: Review of Machine Learning + Imaging

Machine Learning and Imaging

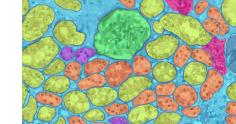
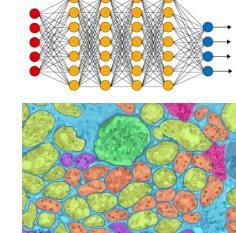
BME 590L  
Roarke Horstmeyer



# What was this class about?

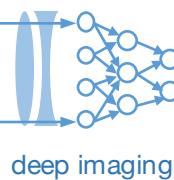
Where your final  
project is aimed

Computer-  
centered hardware  
+ software

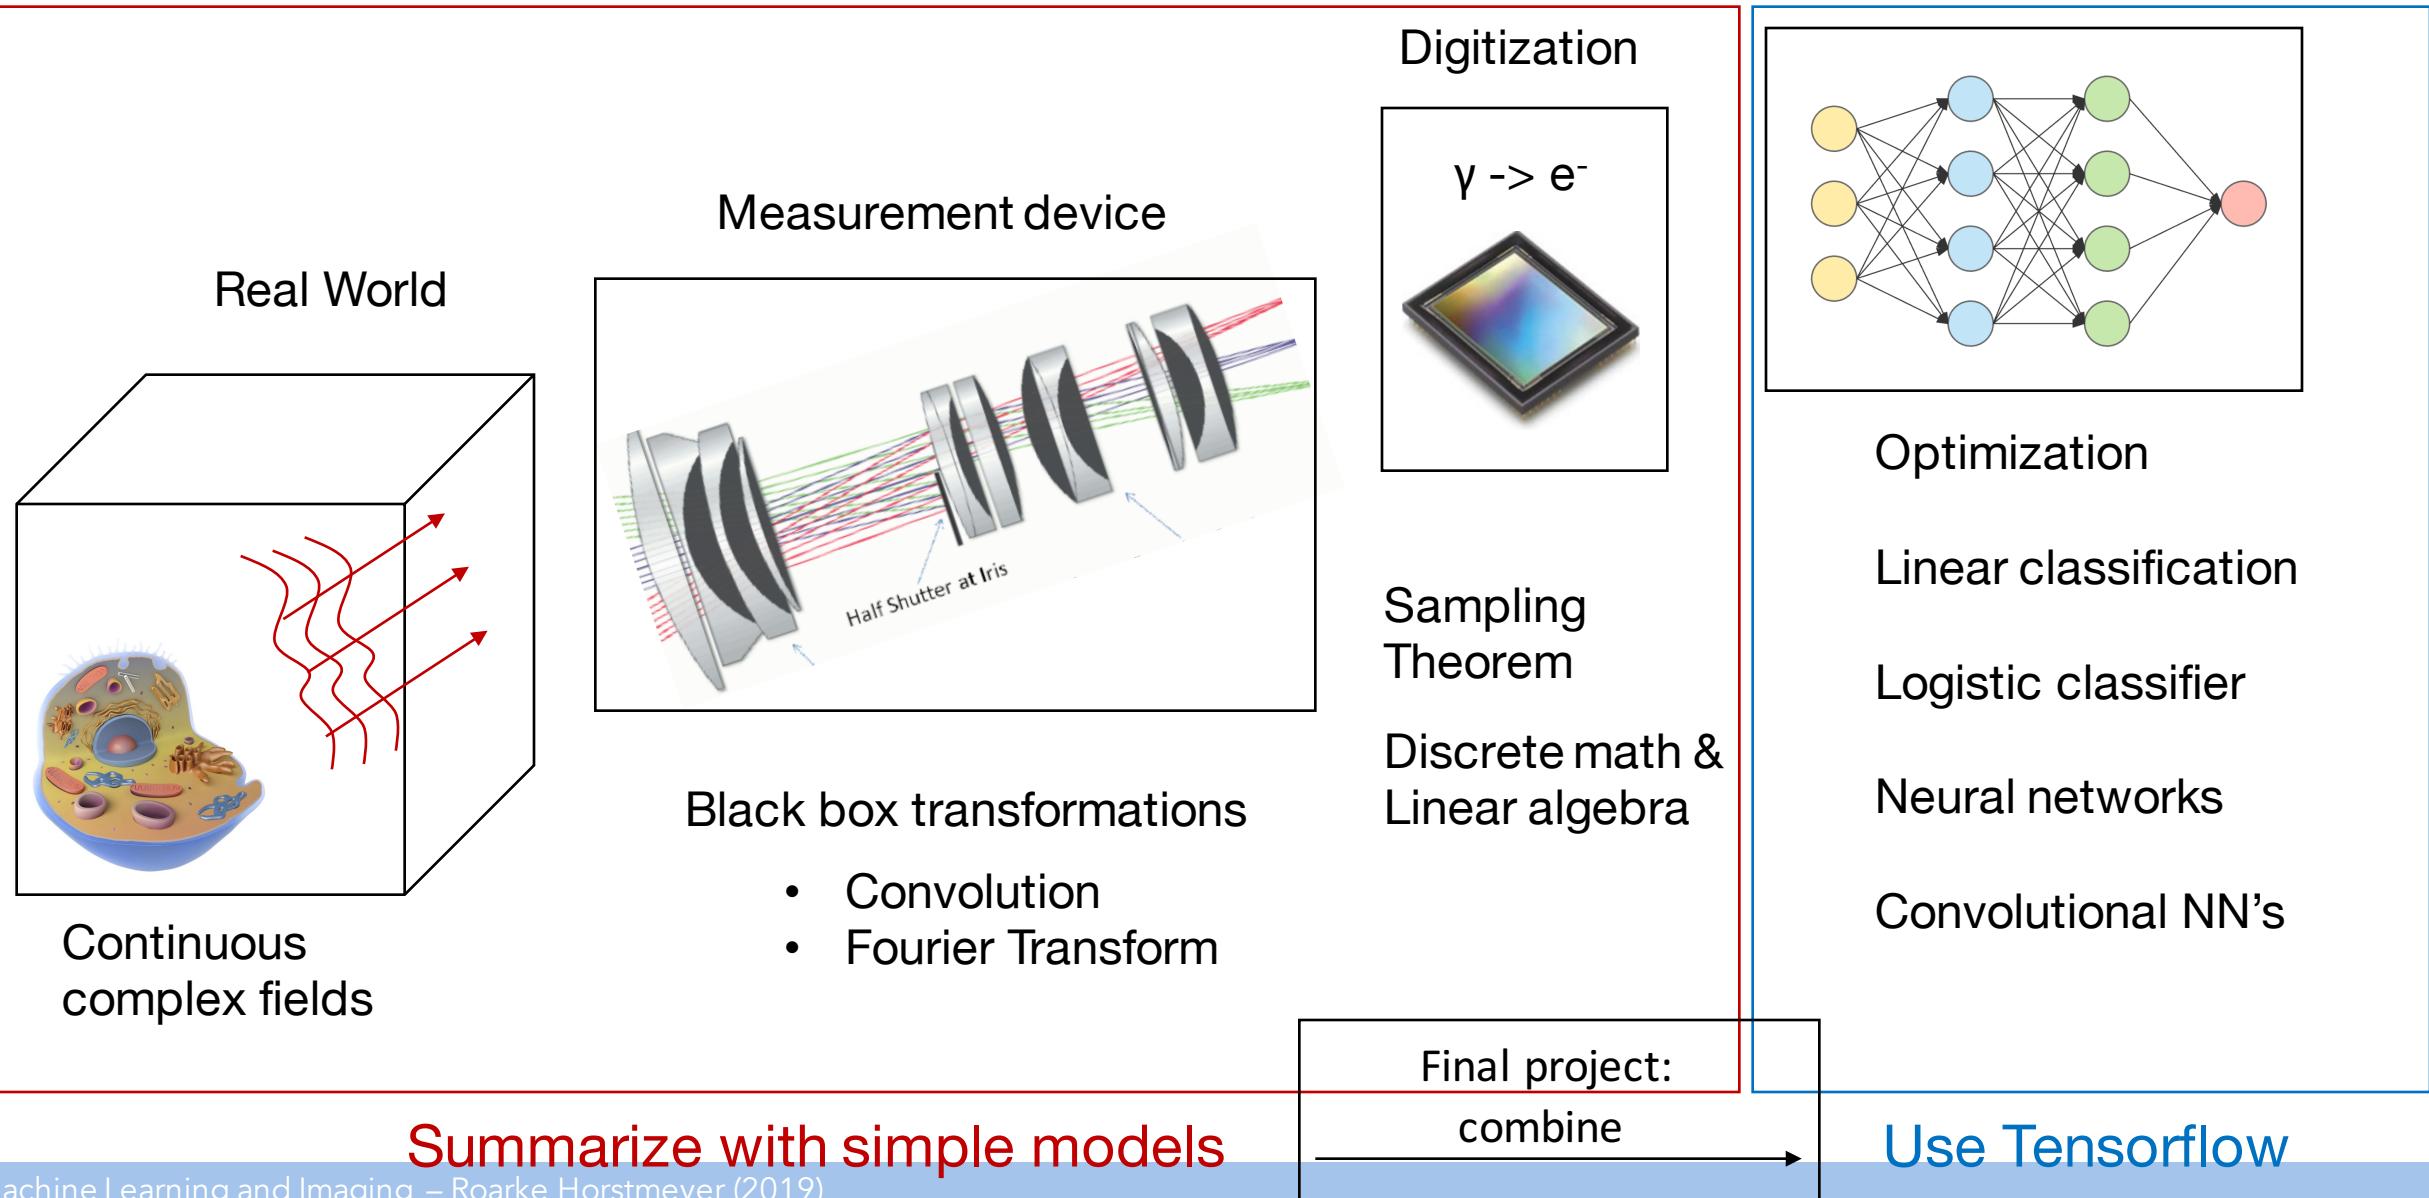


Human-centered hardware design

Computer-centered software design

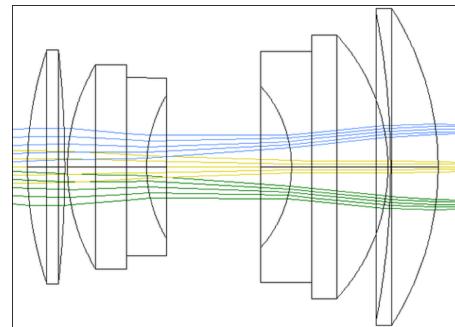
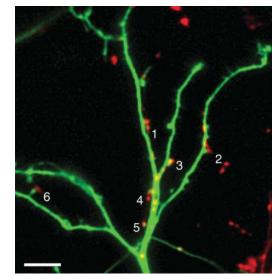


# ML+Imaging pipeline introduction



# Physical models for light propagation to sensor

- Interpretation #1: Radiation (*Incoherent*)
- Model: Rays

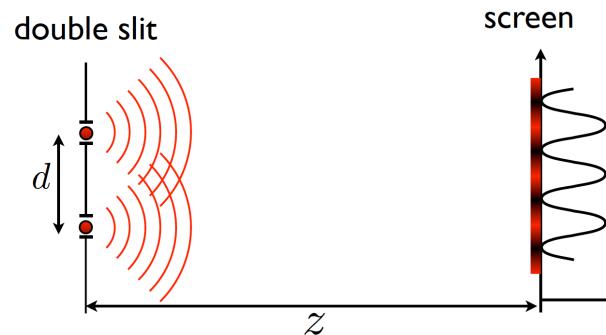
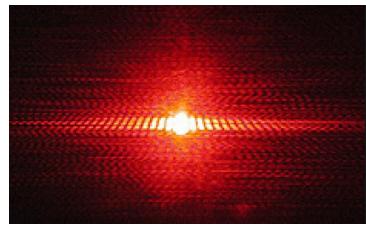


- Real, non-negative

$$I_s = H B S_0$$

- Sample absorption **S**
- Illumination brightness **B**
- Blur in **H**

- Interpretation #2: Electromagnetic wave (*Coherent*)
- Model: Waves



- Complex field

$$I_c = |H C S_c|^2$$

- Sample abs./phase **S**
- Illumination wave **B**
- Blur in **H**

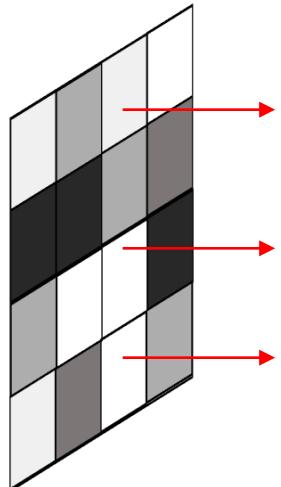
# Mathematical model of for incoherent image formation

- All quantities are real, and non-negative

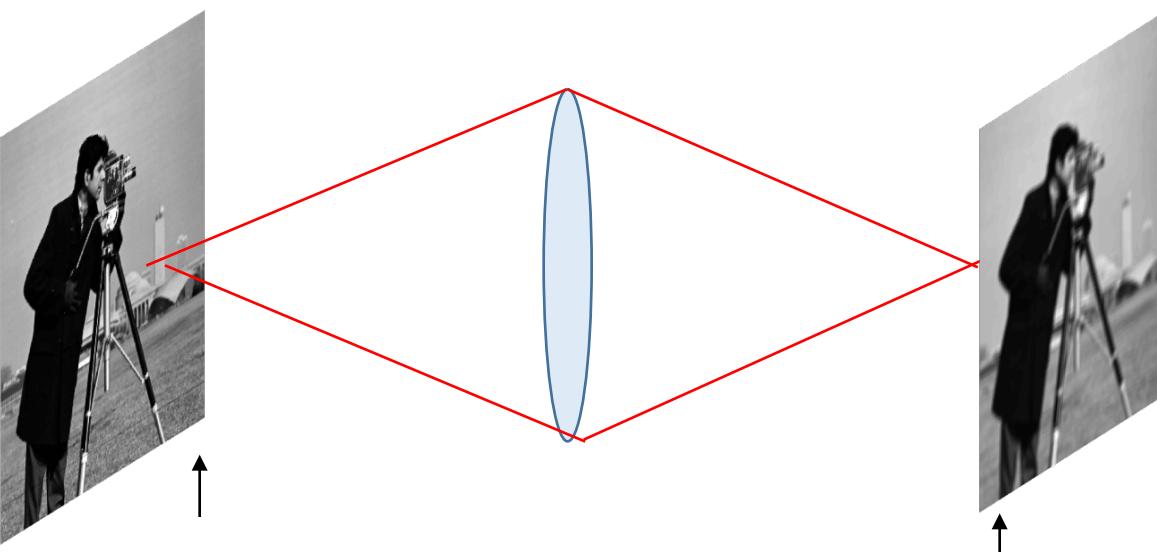
Object absorption:

Illumination brightness:

$$\mathbf{B}(x,y)$$



$$\mathbf{S}_0(x,y)$$



$\mathbf{B} \mathbf{S}_0$   
multiplication

$(\mathbf{B} \mathbf{S}_0) * \mathbf{h}$   
convolution

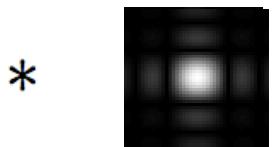
## We can also add in some lens blur

Lenses blur and rescale images:

Input  
intensity



Convolution filter  $h$



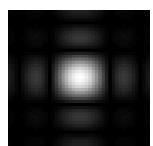
“Incoherent point-  
spread function”



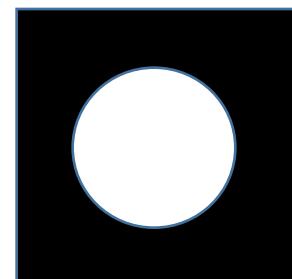
Output intensity

$$h(x,y) = | F[ A(f_x, f_y) ] |^2$$

$h(x,y)$



$$= | F [ A(f_x, f_y) ] |^2$$



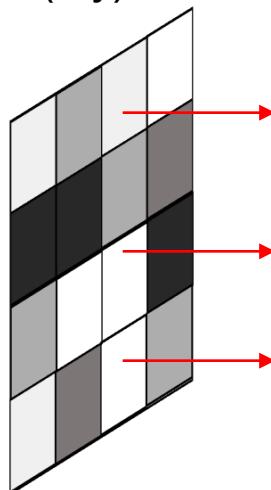
# Mathematical model of for incoherent image formation

- All quantities are real, and non-negative

Object absorption:

Illumination brightness:

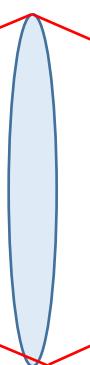
$$\mathbf{B}(x,y)$$



$$\mathbf{S}_0(x,y)$$

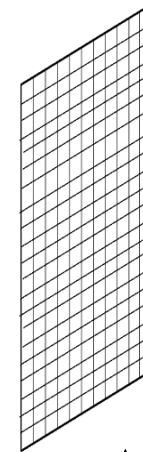


$$\mathbf{B} \mathbf{S}_0$$
  
multiplication



$$(\mathbf{B} \mathbf{S}_0) * \mathbf{h}$$
  
convolution

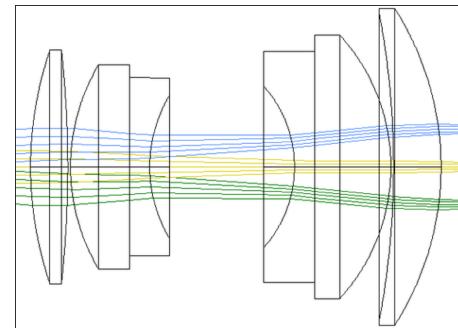
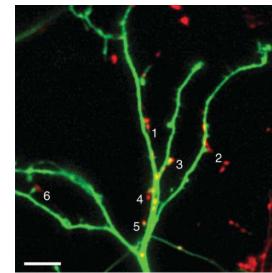
Photons (intensity) hits detector



$$\mathbf{I}_s = \mathbf{H} \mathbf{B} \mathbf{S}_0$$

# Physical models for light propagation to sensor

- Interpretation #1: Radiation (*Incoherent*)
- Model: Rays

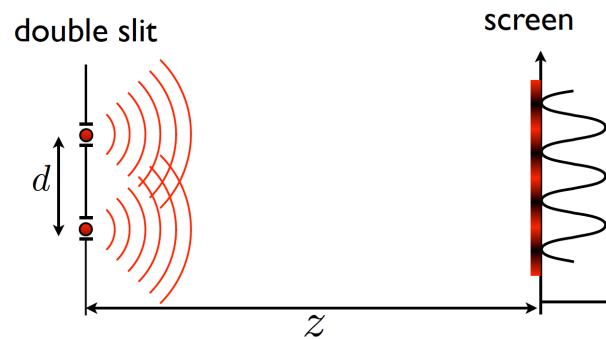
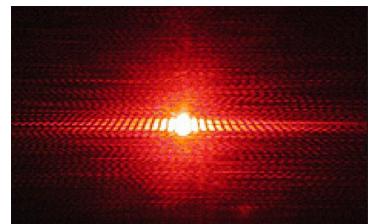


- Real, non-negative

$$I_s = H B S_0$$

- Sample absorption **S**
- Illumination brightness **B**
- Blur in **H**

- Interpretation #2: Electromagnetic wave (*Coherent*)
- Model: Waves



- Complex field

$$I_c = |H C S_c|^2$$

- Sample abs./phase **S**
- Illumination wave **B**
- Blur in **H**

## Let's take a step back: how does light propagate?

Maxwell's equations  
without any charge

$$\nabla \times \vec{\mathcal{E}} = -\mu \frac{\partial \vec{\mathcal{H}}}{\partial t}$$

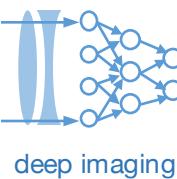
$$\nabla \times \vec{\mathcal{H}} = \epsilon \frac{\partial \vec{\mathcal{E}}}{\partial t}$$

$$\nabla \cdot \epsilon \vec{\mathcal{E}} = 0$$

$$\nabla \cdot \mu \vec{\mathcal{H}} = 0.$$

1. Take the curl of both sides of first equation
2. Substitute 2<sup>nd</sup> and 3<sup>rd</sup> equation
3. Arrive at the wave equation:

$$\nabla^2 \vec{\mathcal{E}} - \frac{n^2}{c^2} \frac{\partial^2 \vec{\mathcal{E}}}{\partial t^2} = 0 \quad n = \left( \frac{\epsilon}{\epsilon_0} \right)^{1/2} \quad c = \frac{1}{\sqrt{\mu_0 \epsilon_0}}.$$



## Let's take a step back: how does light propagate?

Considering light that isn't pulsed over time, we can use the following solution:

$$u(P, t) = A(P) \cos[2\pi\nu t + \phi(P)]$$

$$u(P, t) = \text{Re}\{U(P) \exp(-j2\pi\nu t)\},$$

With this particular solution, we get the following important time-independent equation:

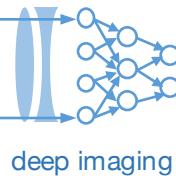
Helmholtz  
Equation

$$(\nabla^2 + k^2)U = 0.$$

$$k = 2\pi n \frac{\nu}{c} = \frac{2\pi}{\lambda},$$

This is an important equation in physics. We won't go into the details, but it leads to the Huygen-Fresnel principle:

$$U(P_2) = \frac{1}{j\lambda} \iint_{\Sigma} U(P_1) \frac{\exp(jkr_{21})}{r_{21}} \cos\theta ds$$



## From the Fresnel approximation to the Fraunhofer approximation

**Fresnel Approximation:**

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

Lets assume that the second plane is “pretty far away” from the first plane. Then,

$$z > \frac{2D^2}{\lambda}$$

1. Expand the squaring

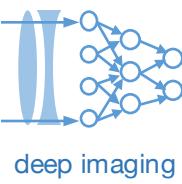
$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint E(x', y', 0) e^{\frac{ik}{2z} (x^2 + y^2)} e^{\frac{ik}{2z} (x'^2 + y'^2)} e^{\frac{ik}{2z} (xx' + yy')} dx' dy'$$

2. Front term comes out, assume second term goes away, then,

$$E(x, y, z) = C \iint E(x', y', 0) e^{\frac{ik}{2z} (xx' + yy')} dx' dy'$$

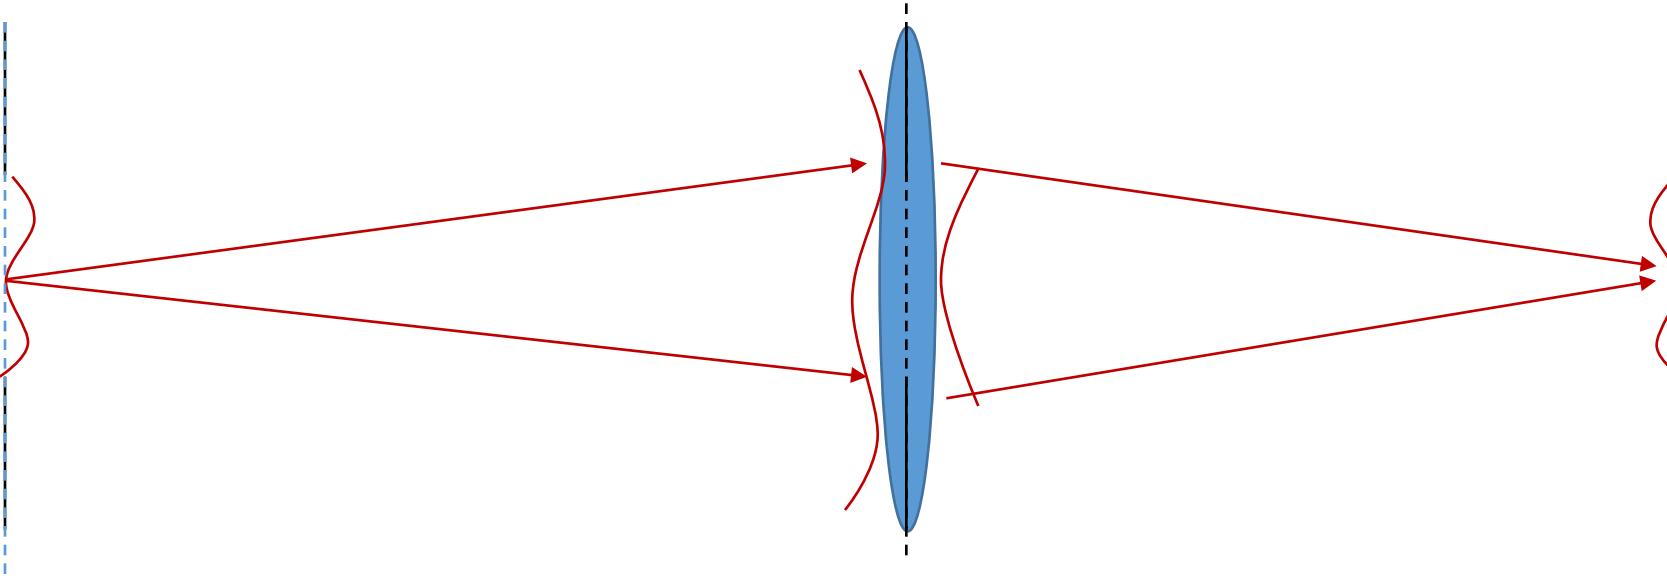
$$C = \frac{e^{ikz}}{i\lambda z} e^{\frac{ik}{2z} (x^2 + y^2)}$$

**Fraunhofer diffraction is a Fourier transform!!!!!!**



## Model of a microscope (or camera) using Fourier transforms:

from lens to sensor, light undergoes an  
***inverse Fourier transform!***

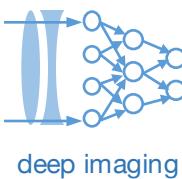


$$E_s(x_s, y_s, 0)$$

$$E'_a(x_d, y_d) = E_a(x_d, y_d)A(x_d, y_d)$$

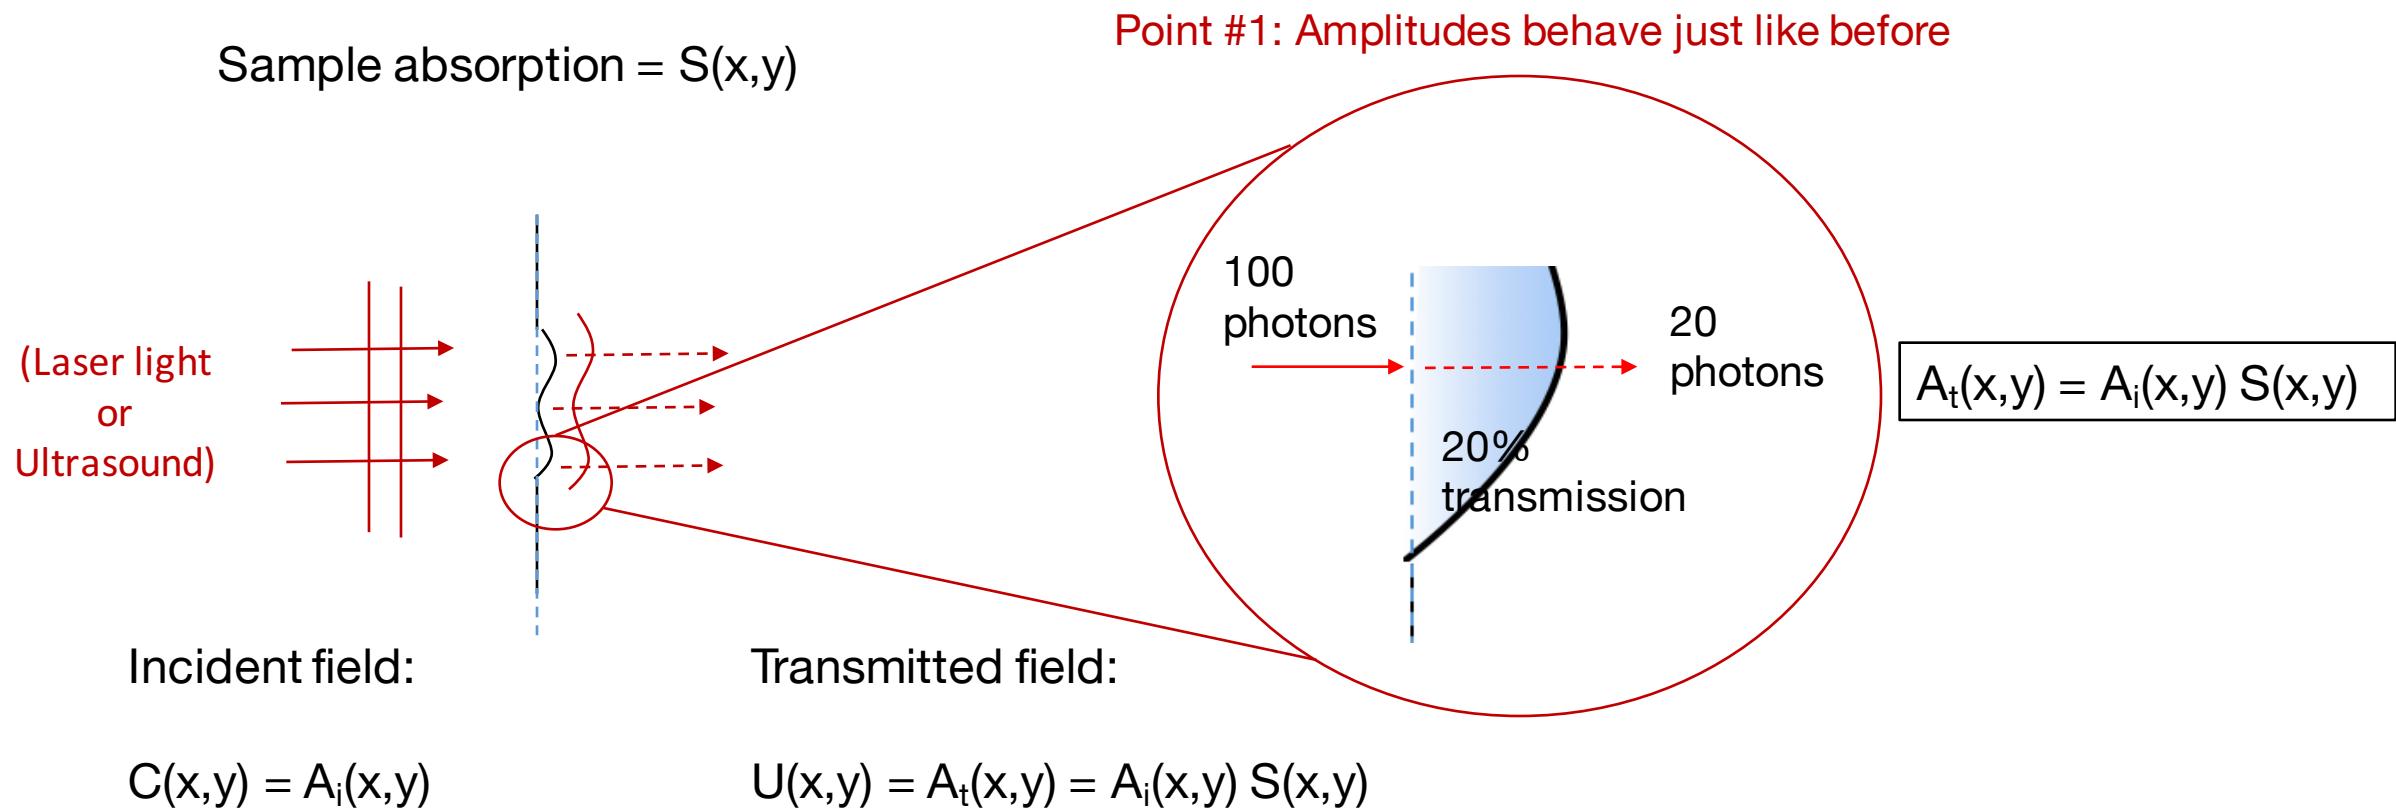
2D Fourier Transform

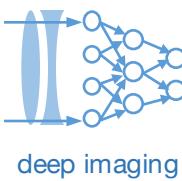
2D inverse Fourier Transform



# Mathematical model of for coherent image formation

- Pretty much the same thing, but now we have an amplitude and a complex phase





# Mathematical model of for coherent image formation

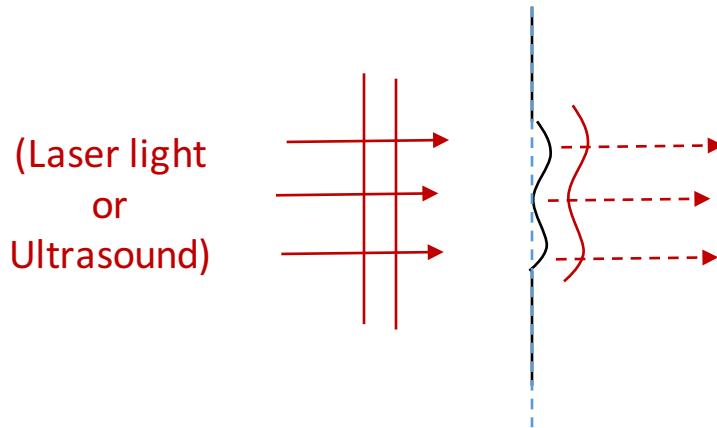
- Pretty much the same thing, but now we have an amplitude and a complex phase

Sample absorption =  $S(x,y)$

**Sample phase delay** =  $\exp[ik\varphi(x,y)]$

**New: complex phase delay**

- Needed to represent wave
- Represents wave delay across space



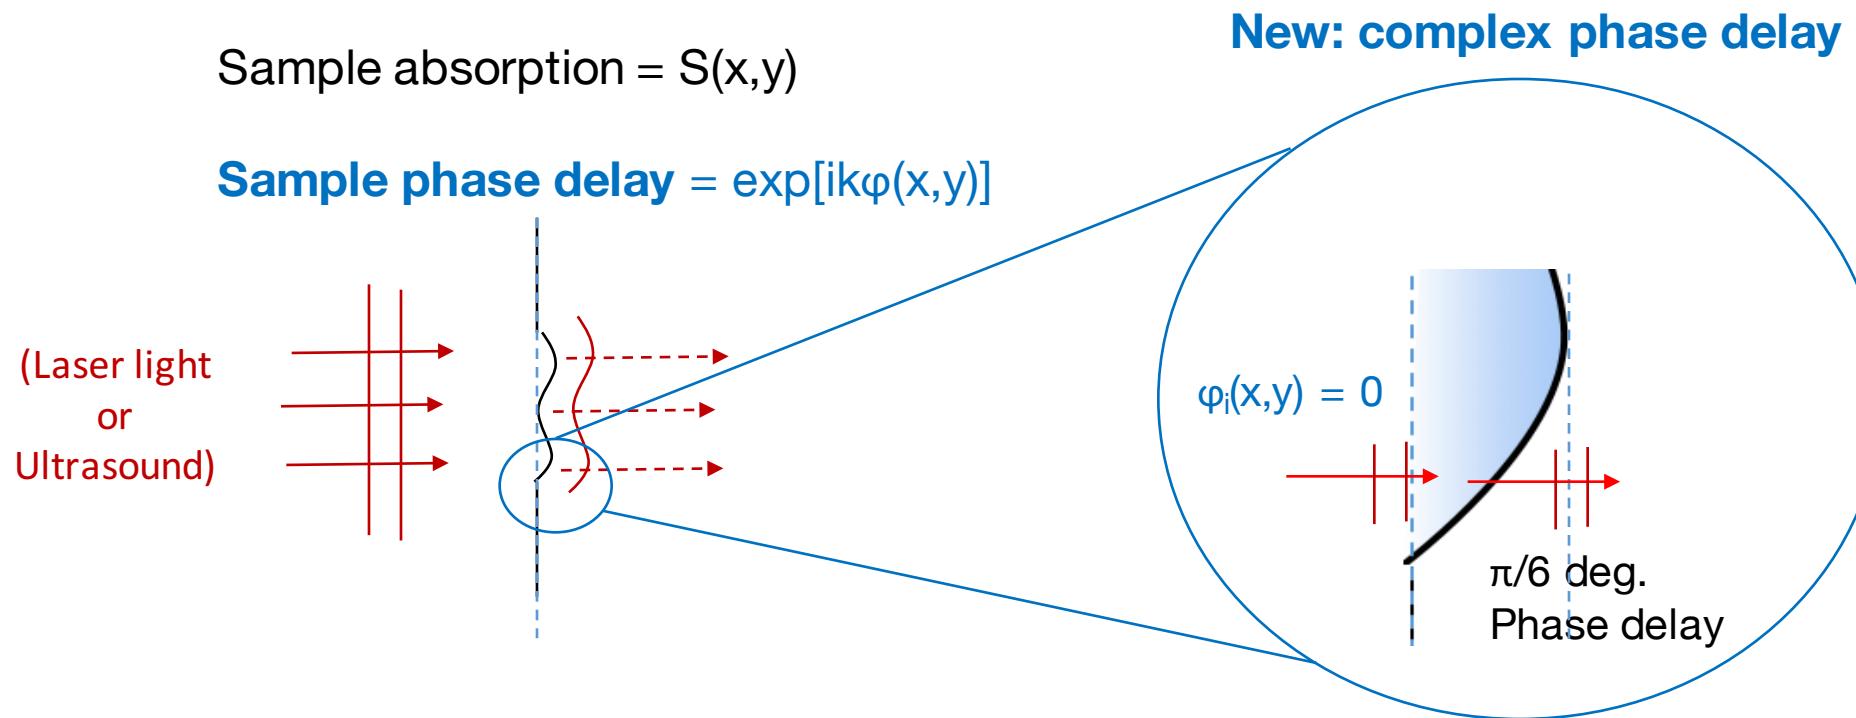
Incident field:

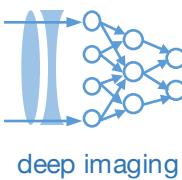
Transmitted field:

$$C(x,y) = A_i(x,y) \exp[ik\varphi_i(x,y)] \quad U(x,y) = A_i(x,y) S(x,y) \exp[ik\varphi_t(x,y)]$$

# Mathematical model of for coherent image formation

- Pretty much the same thing, but now we have an amplitude and a complex phase





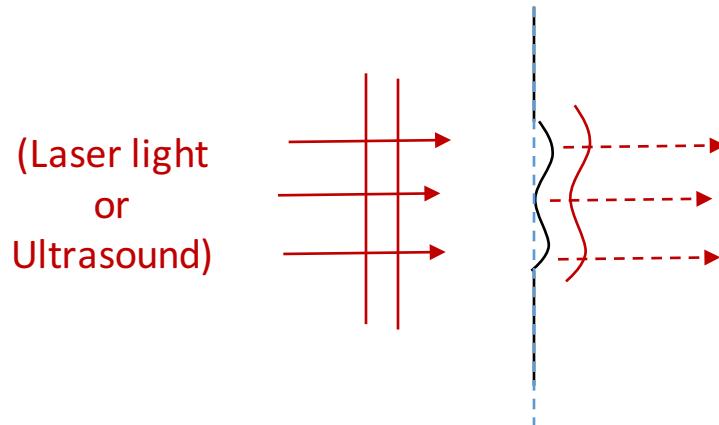
## Mathematical model of for coherent image formation

- Pretty much the same thing, but now we have an amplitude and a complex phase

Sample absorption =  $S(x,y)$

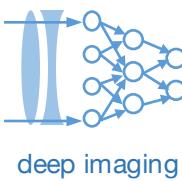
Sample phase delay =  $\exp[ik\varphi(x,y)]$

Output phase is sum of phase delays, product of phasors



$$\varphi_t(x,y) = \varphi(x,y) + \varphi_i(x,y)$$

$$\exp[ik\varphi_t(x,y)] = \exp[ik\varphi_i(x,y)] \exp[ik\varphi(x,y)]$$

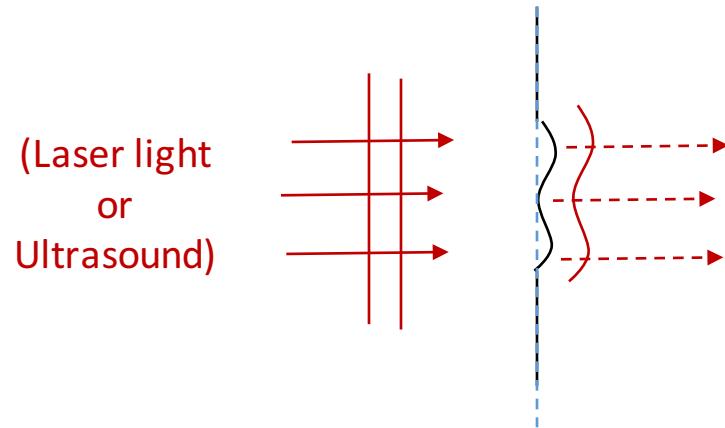


# Mathematical model of for coherent image formation

- Pretty much the same thing, but now we have an amplitude and a complex phase

Sample absorption =  $S(x,y)$

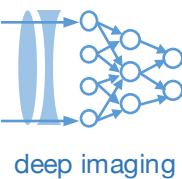
Sample phase delay =  $\exp[ik\varphi(x,y)]$



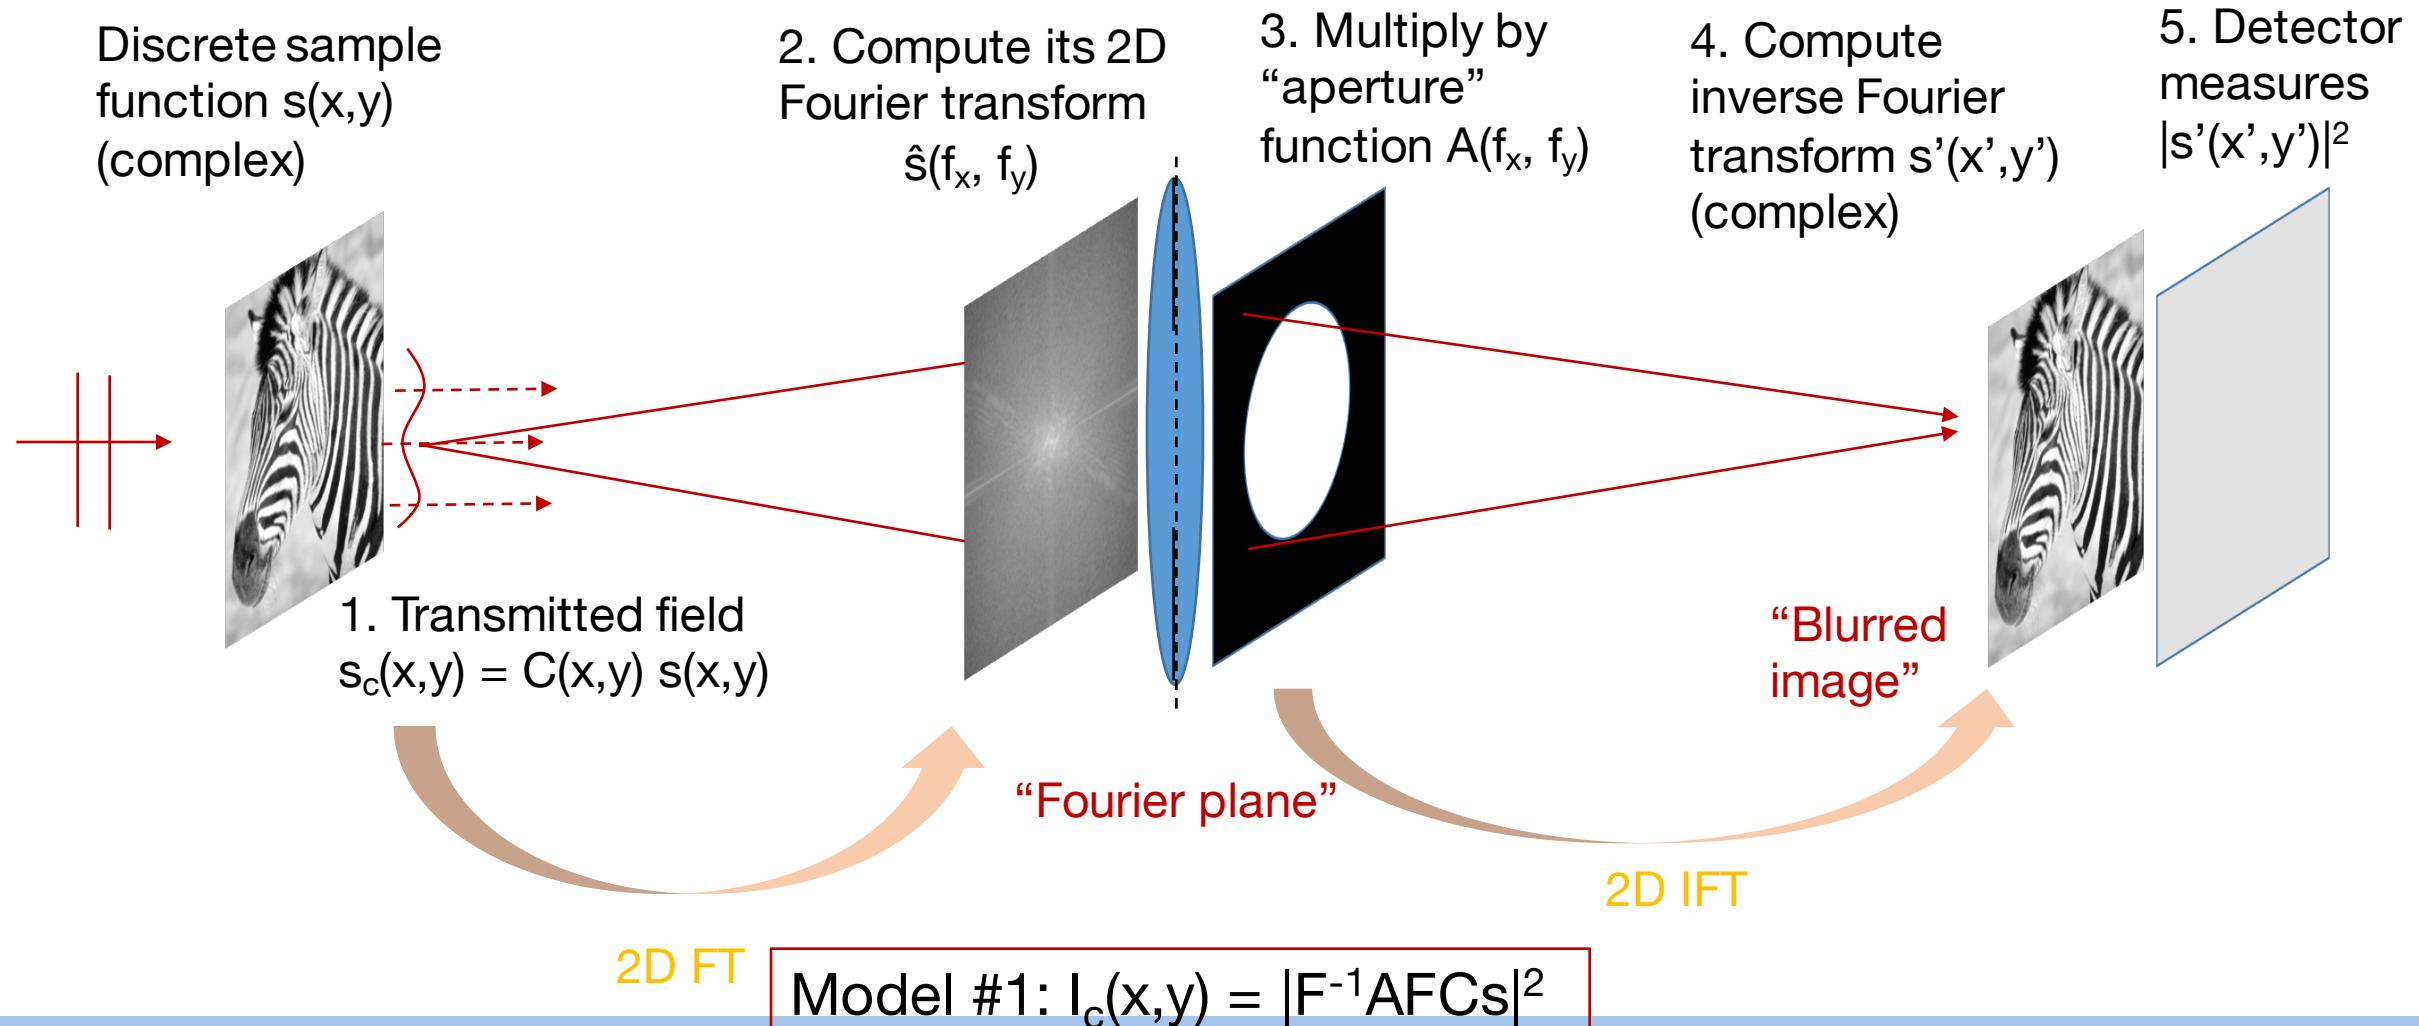
Conclusion:

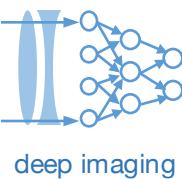
Transmitted field = incident field  $\times$  complex sample :

$$U(x,y) = C(x,y) S(x,y) \exp[ik\varphi(x,y)]$$

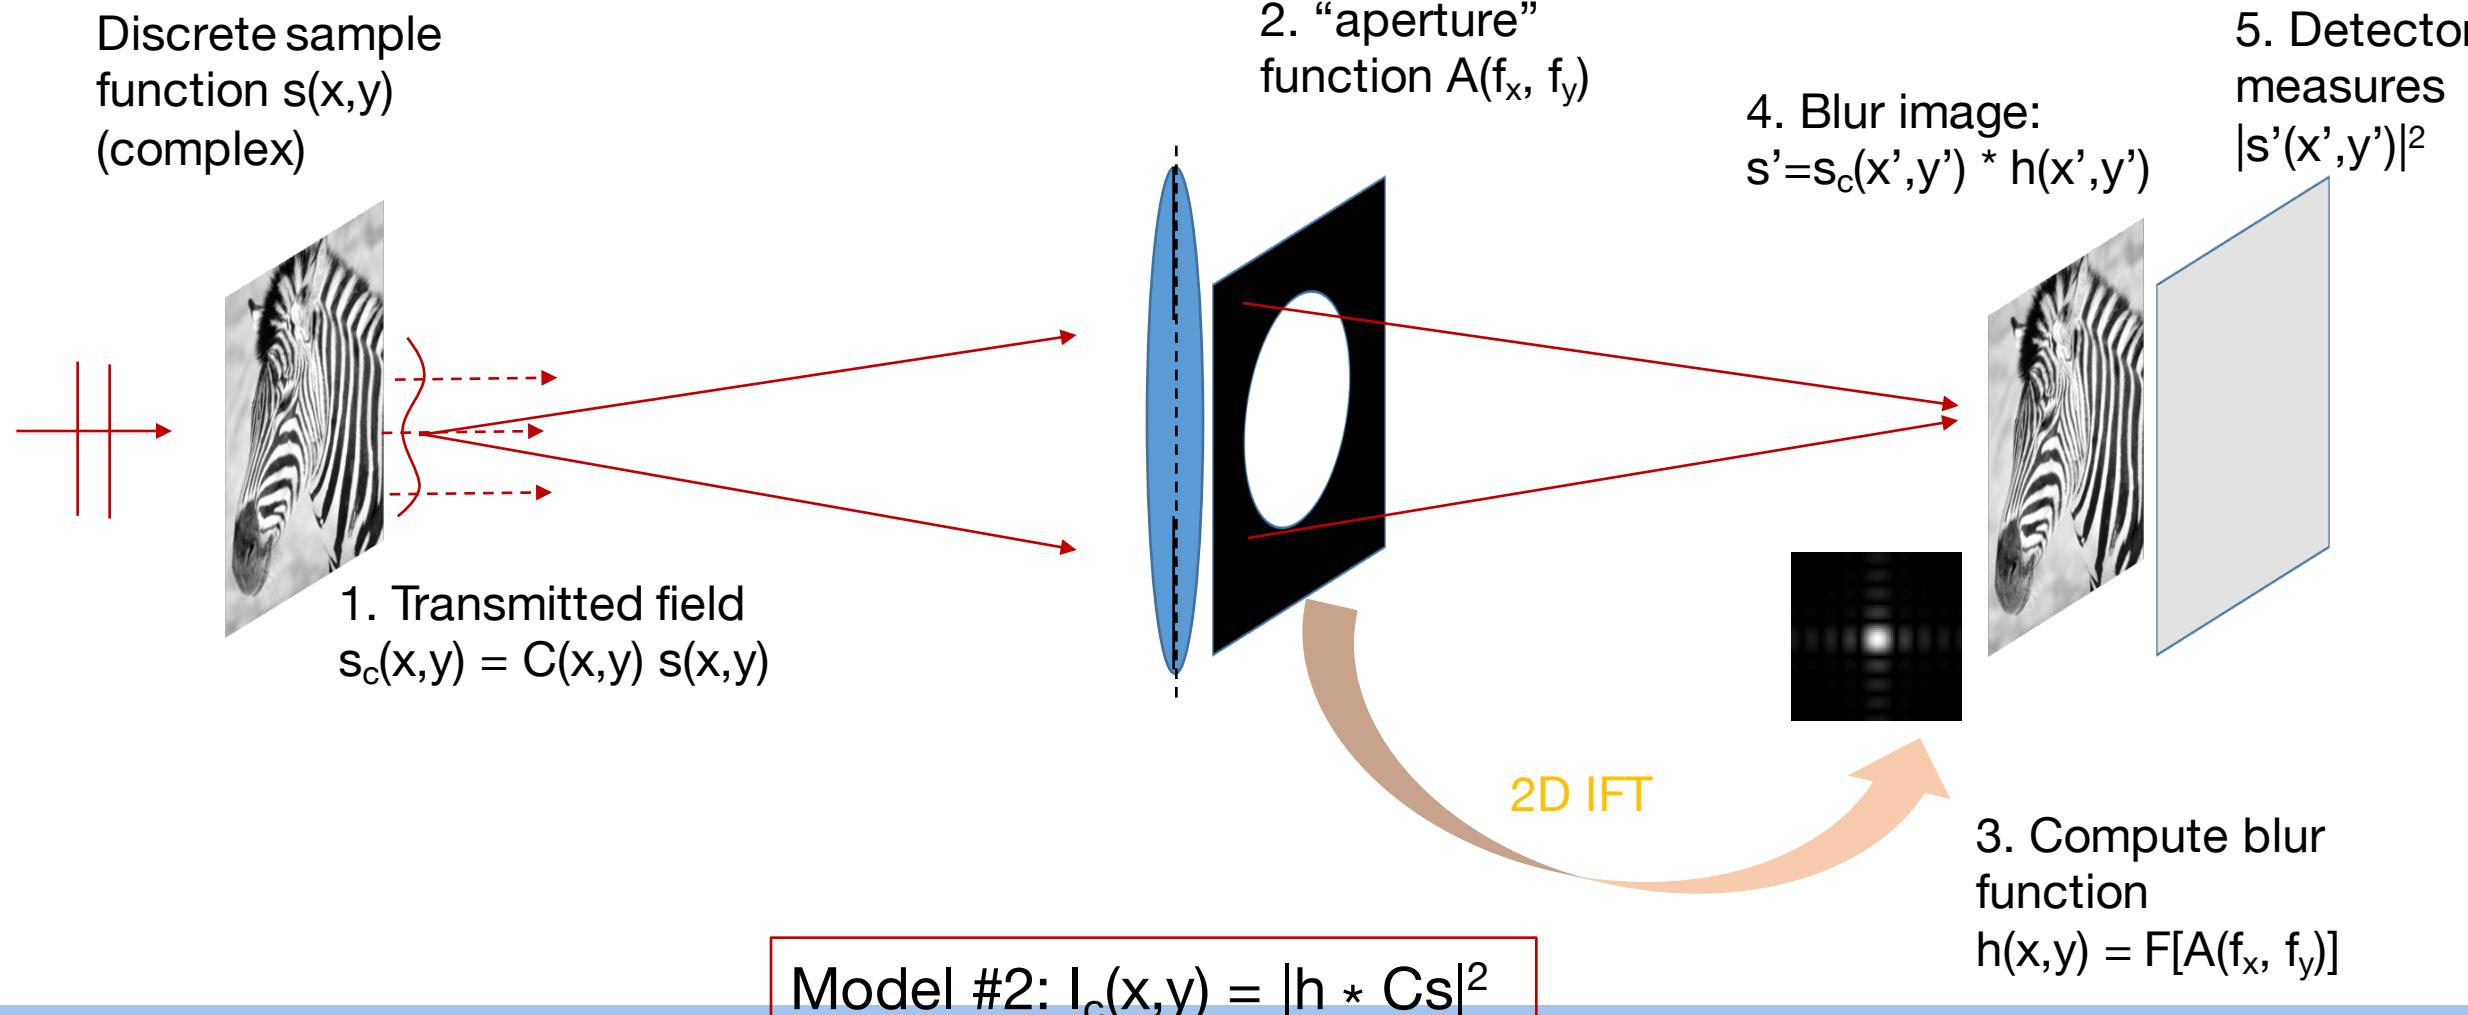


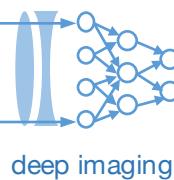
## Model of image formation for wave optics (coherent light):



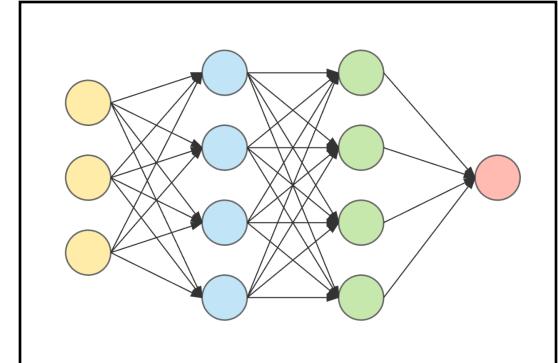


## Model of image formation for wave optics (coherent light):

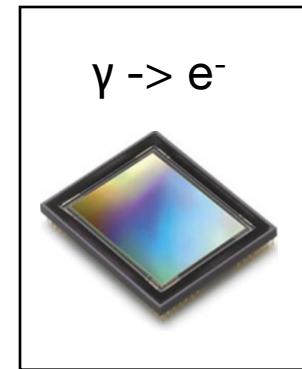




## Machine Learning

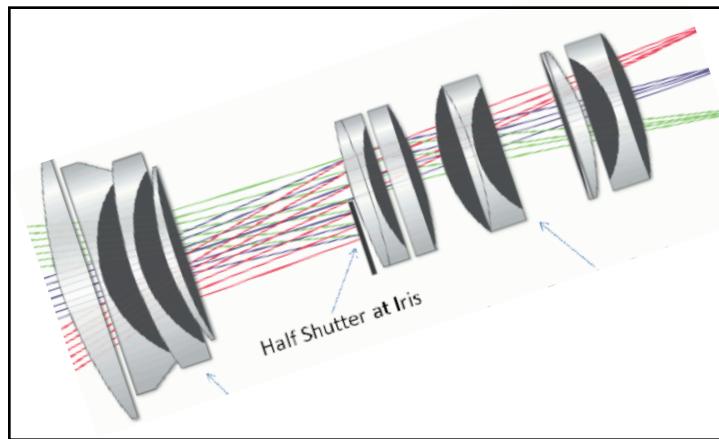


### Digitization



Sampling  
Theorem  
Discrete math &  
Linear algebra

### Measurement device



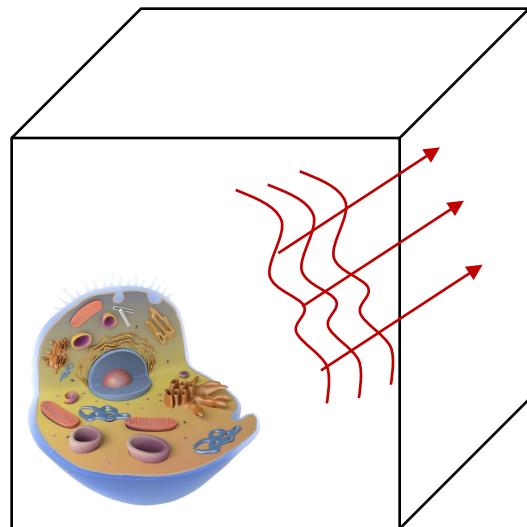
### Black box transformations

- Convolution
- Fourier Transform

$$I_{Inc} = H B S_0$$

$$I_{Coh} = |H C S_C|^2$$

### Real World



Continuous  
complex fields

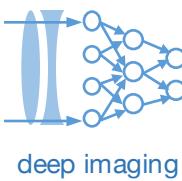
### Optimization

Linear classification

Logistic classifier

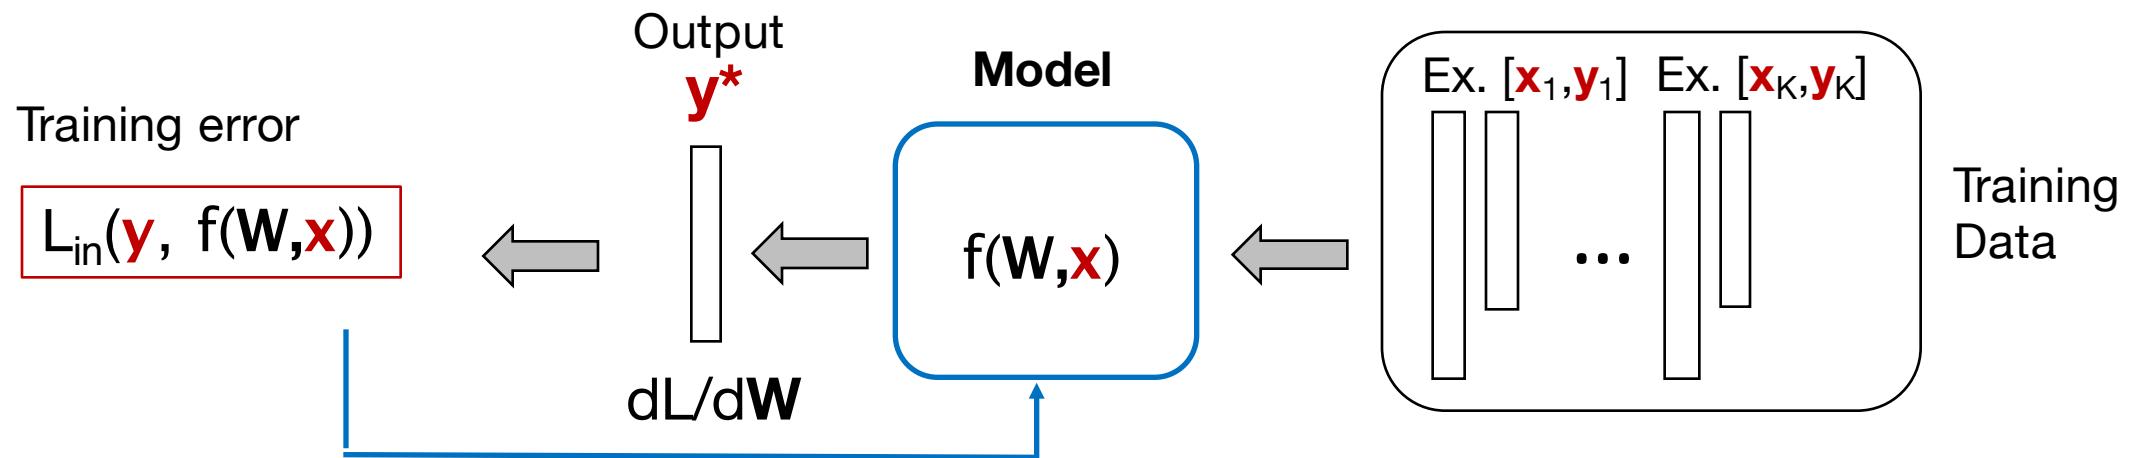
Neural networks

Convolutional NN's



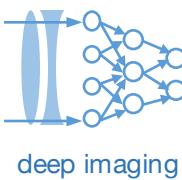
# Summary of machine learning pipeline:

## 1. Network Training



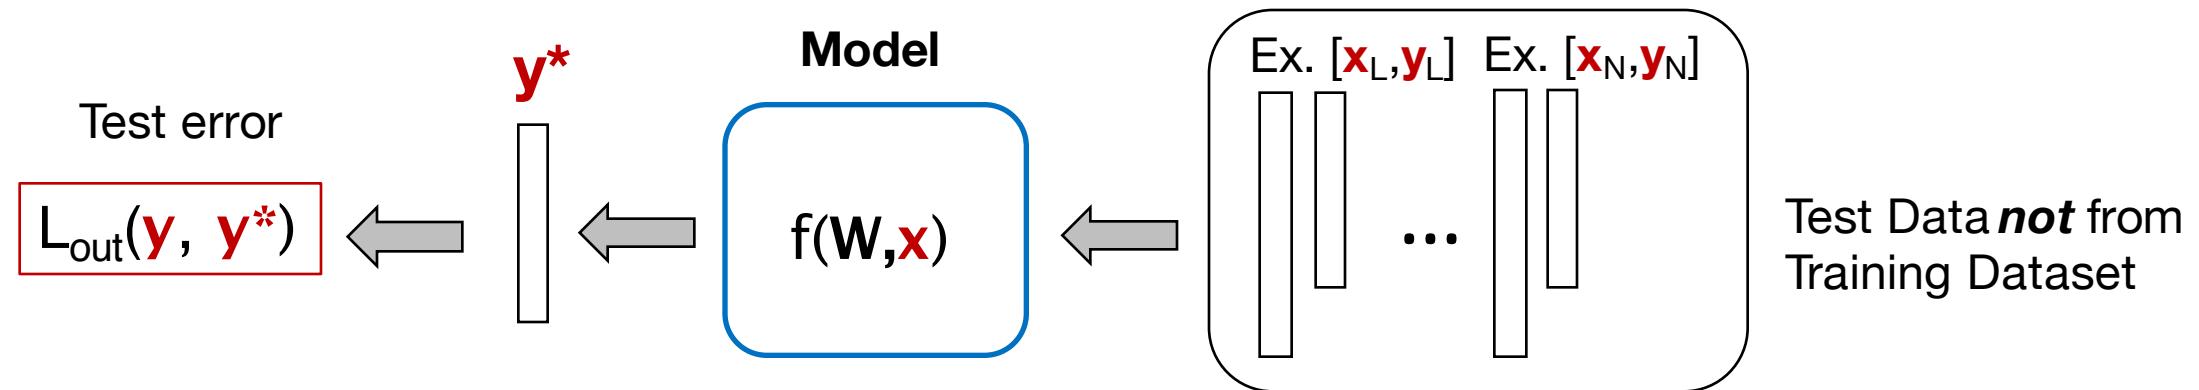
What we need for network training:

- 1. Labeled examples**
- 2. A model and loss function**
- 3. A way to minimize the loss function  $L$**



## Summary of machine learning pipeline:

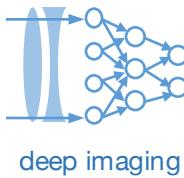
### 2. Network Testing



What we need for network testing:

4. ***Unique* labeled test data**
5. **Evaluation of model error**

# Let's start with a simpler approach: linear regression



$$L = \frac{1}{N} \sum_{i=1}^N L_i(f(x_i, \mathbf{W}), y_i)$$

General linear model:

$$L = \frac{1}{N} \sum_{i=1}^N L_i(\mathbf{W}\mathbf{x}_i, y_i)$$

Assume 1 class =  
1 linear fit

Use MSE error  
model

Where labels  
determined by  
thresholding

$$L = \frac{1}{N} \sum_{i=1}^N L_i(w^T x_i - y_i)$$

$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$

$$f(\mathbf{x}_i) = y_i^* = \text{sgn}(\mathbf{w}^T \mathbf{x}_i)$$

# classes  $\uparrow$   $\mathbf{y}$  =  $\boxed{\mathbf{w}}$   $\mathbf{x}$

1 var.  $\uparrow$   $\mathbf{y}$  =  $\boxed{\mathbf{w}^T}$   $\mathbf{x}$

$$\text{sgn}(x) := \begin{cases} -1 & \text{if } x < 0, \\ 0 & \text{if } x = 0, \\ 1 & \text{if } x > 0. \end{cases}$$

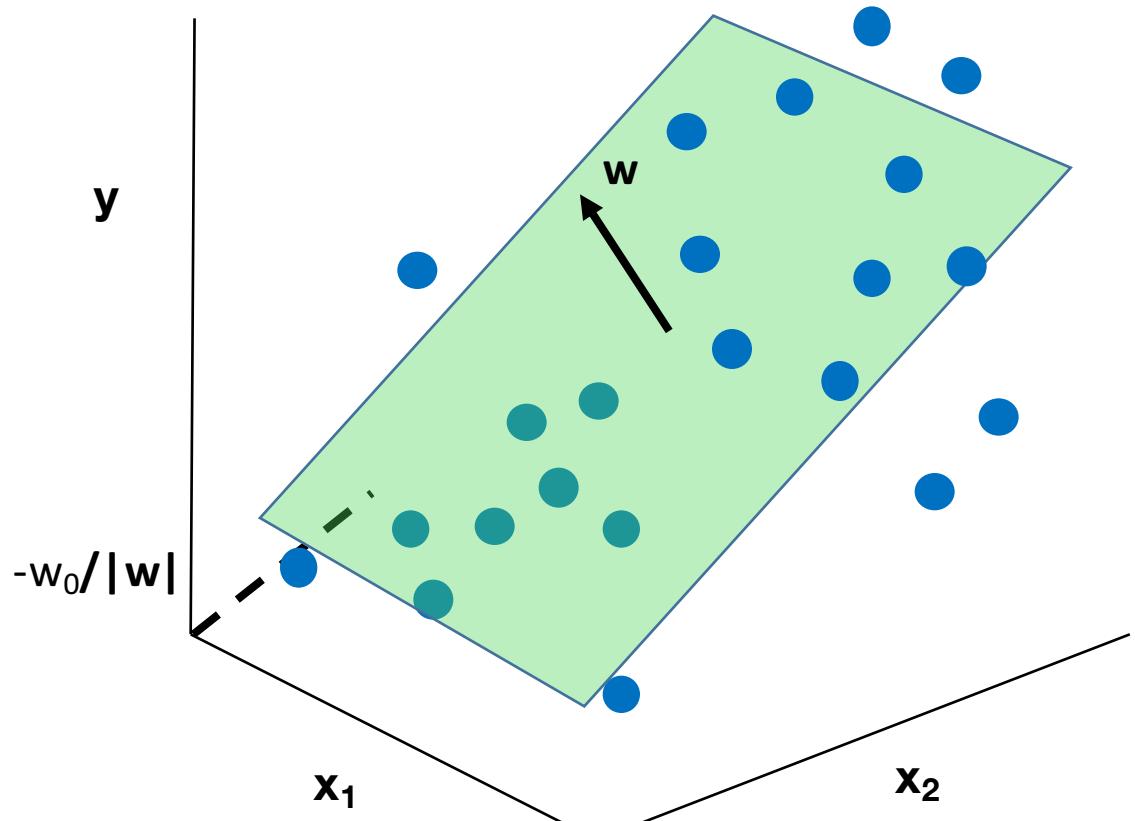
# Why does linear regression with $\text{sgn}()$ achieve classification?

Without  $\text{sgn}()$ : regression for best fit

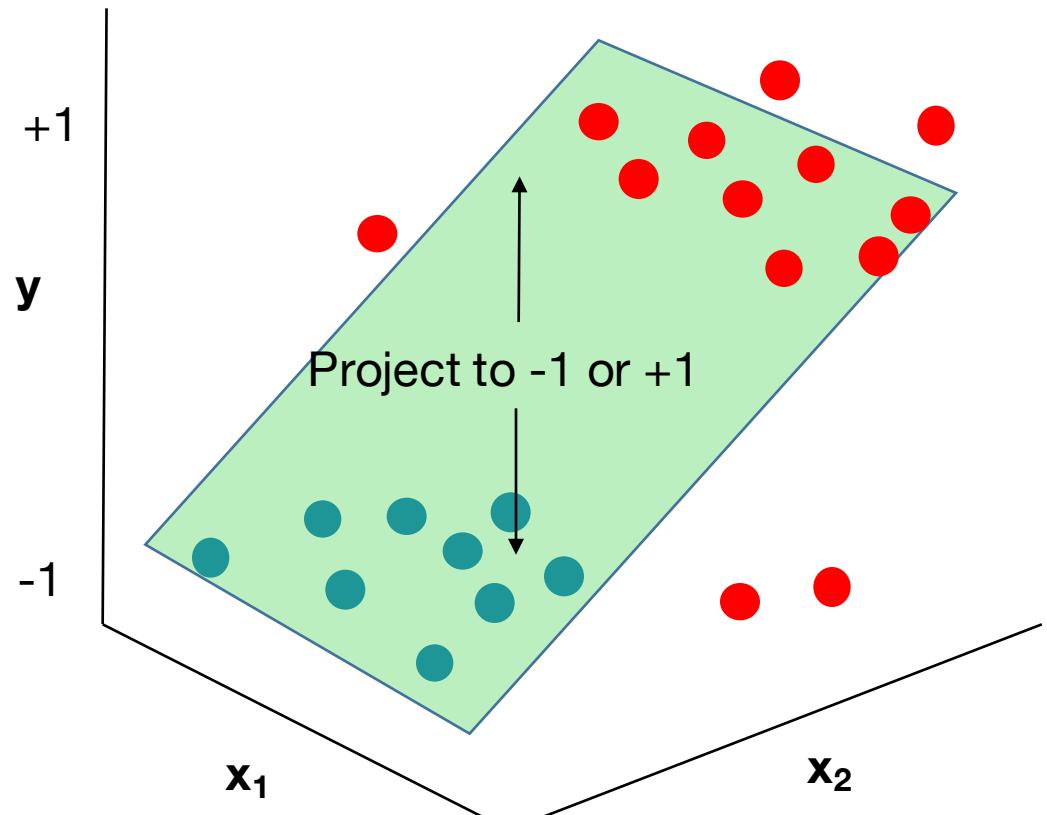
$$f(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i$$

$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$

- If  $y_i$  can be anything, minimizing  $L$  makes  $\mathbf{w}$  the plane of best fit



# Why does linear regression with $\text{sgn}()$ achieve classification?



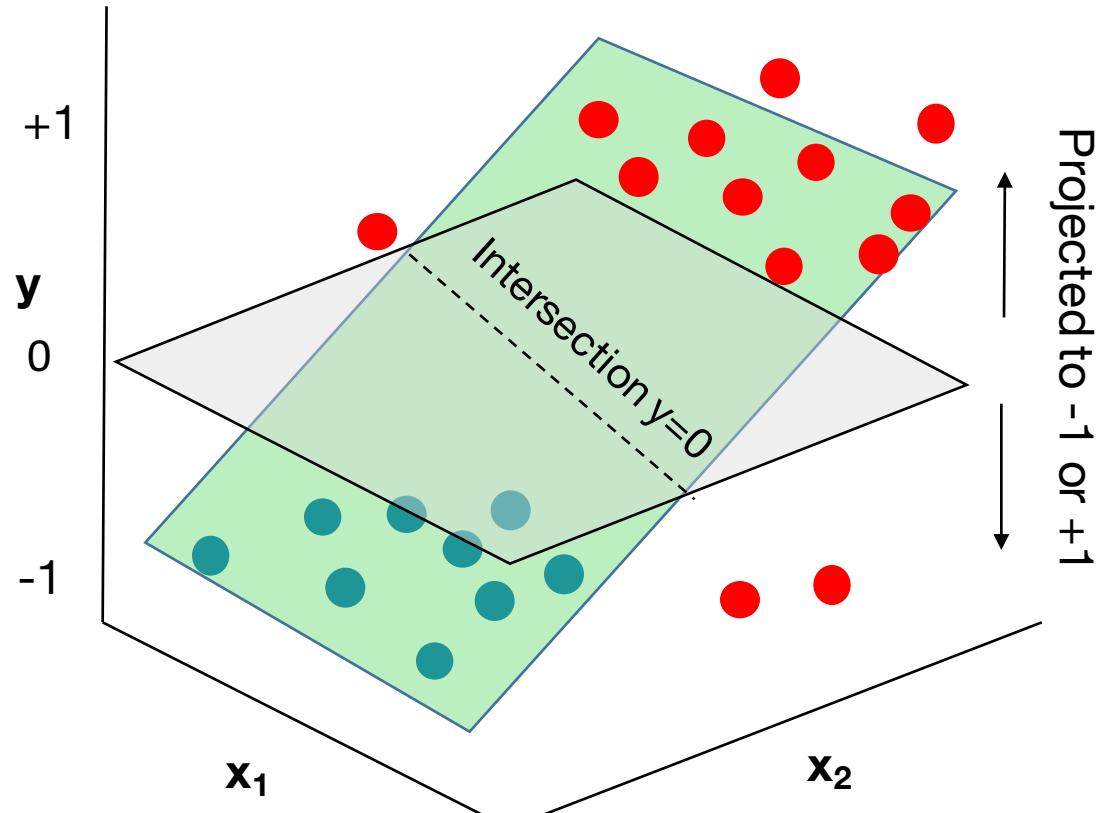
Without  $\text{sgn}()$ : regression for best fit

$$f(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i$$

$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$

- $y_i$  can only be  $-1$  or  $+1$ , which defines its class
- Can still find plane of best fit

# Why does linear regression with $\text{sgn}()$ achieve classification?



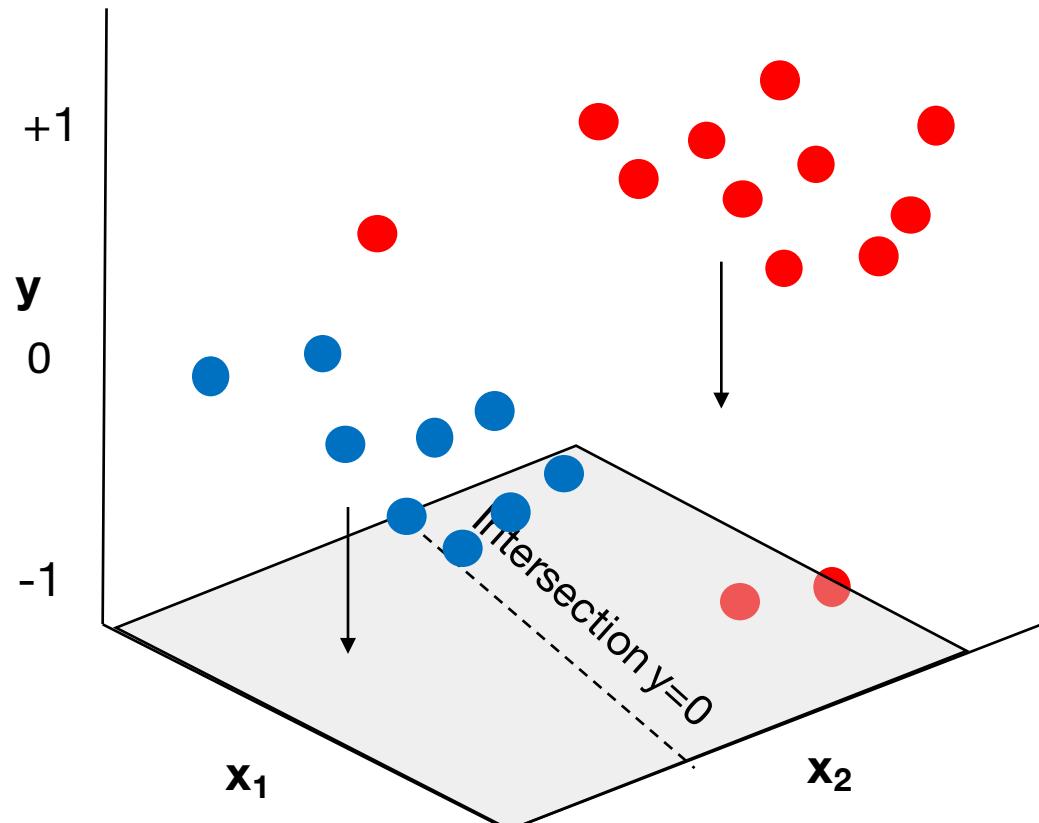
With  $\text{sgn}()$  operation:

$$f(\mathbf{x}_i) = y_i^* = \text{sgn}(\mathbf{w}^T \mathbf{x}_i)$$

$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$

- Anything point to one side of  $y=0$  intersection is class  $+1$ , anything on the other side of intersection is class  $-1$

# Why does linear regression with $\text{sgn}()$ achieve classification?



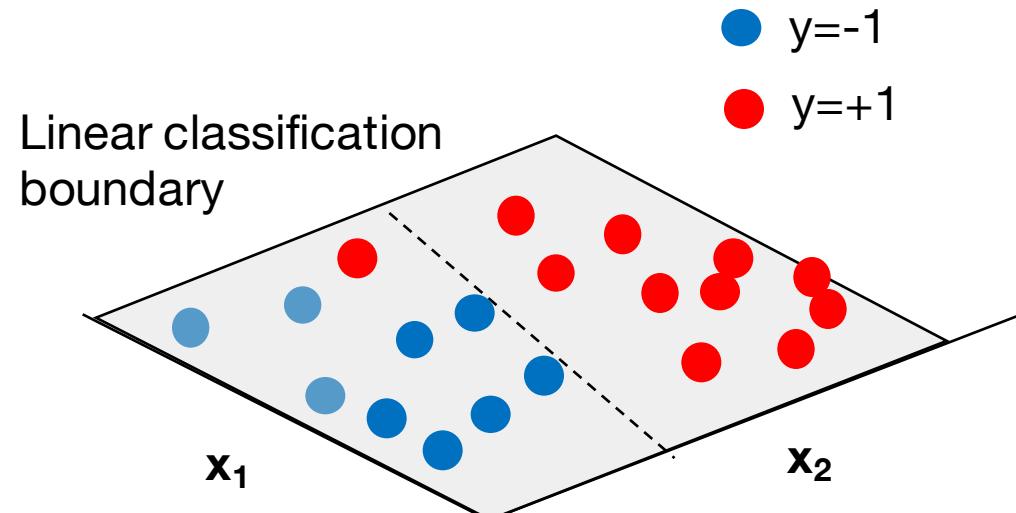
With  $\text{sgn}()$  operation:

$$f(\mathbf{x}_i) = y_i^* = \text{sgn}(\mathbf{w}^T \mathbf{x}_i)$$

$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$

- $y$  axis isn't really needed now & can view this decision boundary in 2D

# Why does linear regression with $\text{sgn}()$ achieve classification?



With  $\text{sgn}()$  operation:

$$f(\mathbf{x}_i) = y_i^* = \text{sgn}(\mathbf{w}^T \mathbf{x}_i)$$

$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$

Sign operation takes linear regression and makes it a classification operation!

# In the rest of this class: solve via gradient descent



With a matrix, compute this for each entry:

$$\frac{dL(W_i)}{dW_i} = \lim_{h \rightarrow 0} \frac{L(W_i + h) - L(W_i)}{h}$$

## Example:

$$W = [1, 2; 3, 4]$$

$$L(W, x, y) = 12.79$$

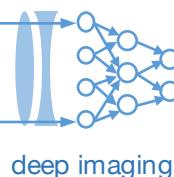
$$W_1 + h = [1.001, 2; 3, 4]$$

$$L(W_1 + h, x, y) = 12.8$$

$$dL(W_1)/dW_1 = 12.8 - 12.79 / .001$$

$$dL(W_1)/dW_1 = 10$$

- Repeat for all entries of  $W$ ,  $dL/dW$  will have  $N \times M$  entries for  $N \times M$  matrix

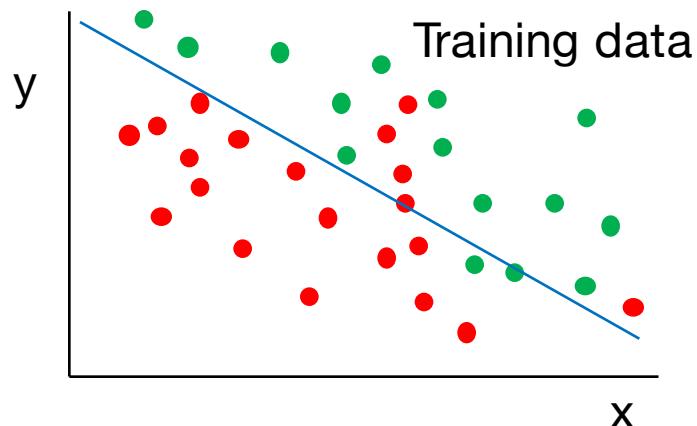
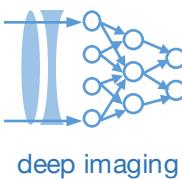


## Steepest descent and the best step size $\epsilon$

1. Evaluate function  $f(\mathbf{x}^{(0)})$  at an initial guess point,  $\mathbf{x}^{(0)}$
2. Compute gradient  $\mathbf{g}^{(0)} = \nabla_{\mathbf{x}} f(\mathbf{x}^{(0)})$
3. Next point  $\mathbf{x}^{(1)} = \mathbf{x}^{(0)} - \epsilon^{(0)} \mathbf{g}^{(0)}$
4. Repeat –  $\mathbf{x}^{(n+1)} = \mathbf{x}^{(n)} - \epsilon^{(n)} \mathbf{g}^{(n)}$ , until  $|\mathbf{x}^{(n+1)} - \mathbf{x}^{(n)}| < \text{threshold } t$

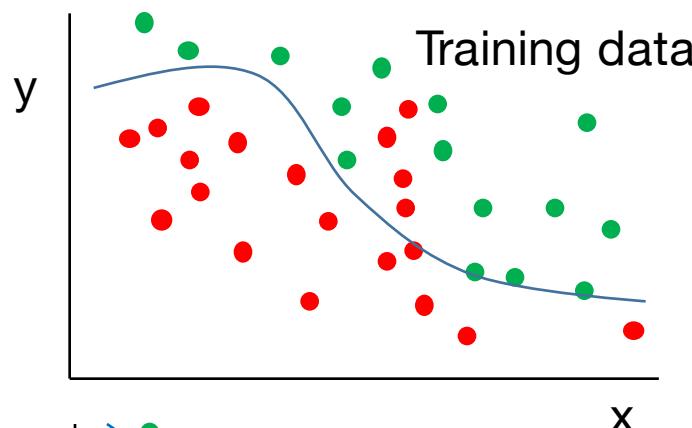
$$L = \frac{1}{N} \sum_{i=1}^N (w^T x_i - y_i)^2$$
$$\nabla L(w) = \frac{2}{N} X^T (Xw - y) = 0$$

```
while previous_step_size > precision and iters < max_iters:  
    prev_x = cur_x  
    cur_x -= gamma * df(prev_x)  
    previous_step_size = abs(cur_x - prev_x)  
    iters+=1
```



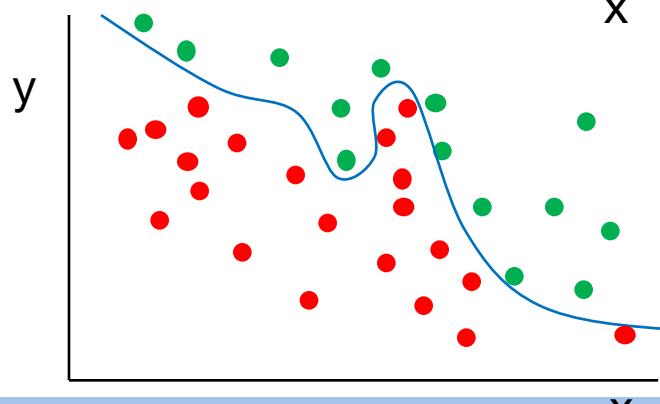
$$f = W_1 x$$

Learned  $f$ : not flexible



$$f = W_2 \max(W_1 x, 0)$$

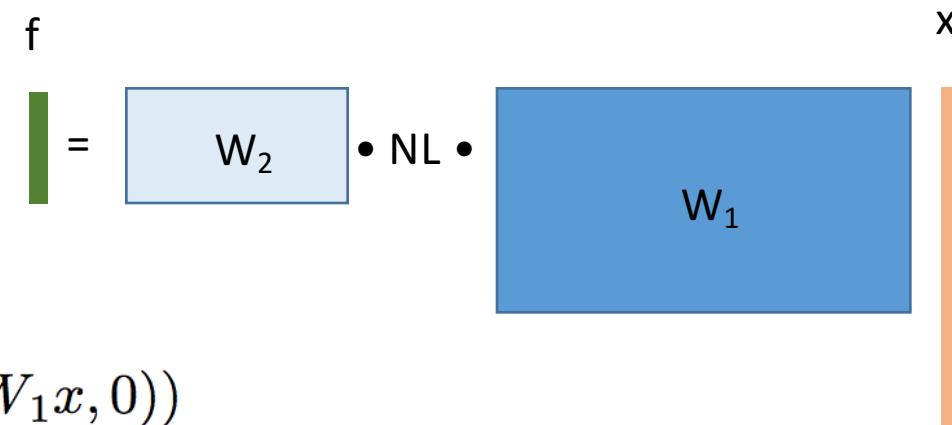
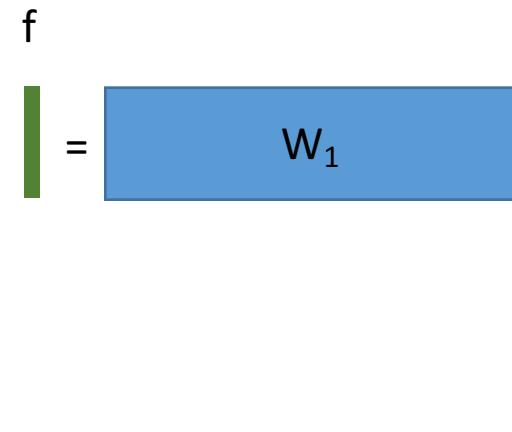
Learned  $f$ : a bit flexible



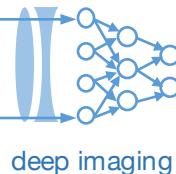
$$f = W_3 \max(0, W_2 \max(W_1 x, 0))$$

Learned  $f$ : more flexible

Does it generalize???



We can keep adding  
these “layers”...



## Before CNN's – understand two competing goals in machine learning

1. Can we make sure the in-sample error  $L_{in}(y, f(x, W))$  is small enough?
  - Appropriate cost function
  - “complex enough” model
  
2. Can we make sure that  $L_{out}(y, f(x, W))$  is close enough to  $L_{in}(y, f(x, W))$ ?
  - Probabilistic analysis says yes!
  - $|L_{in} - L_{out}|$  bounded from above
  - Bound grows with model capacity (bad)
  - Bound shrinks with # of training examples (good)

# Important components of a CNN

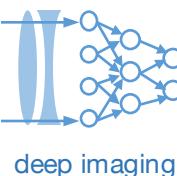
## CNN Architecture

- CONV size, stride, pad, depth
- ReLU & other nonlinearities
- POOL methods
- # of layers, dimensions per layer
- Fully connected layers

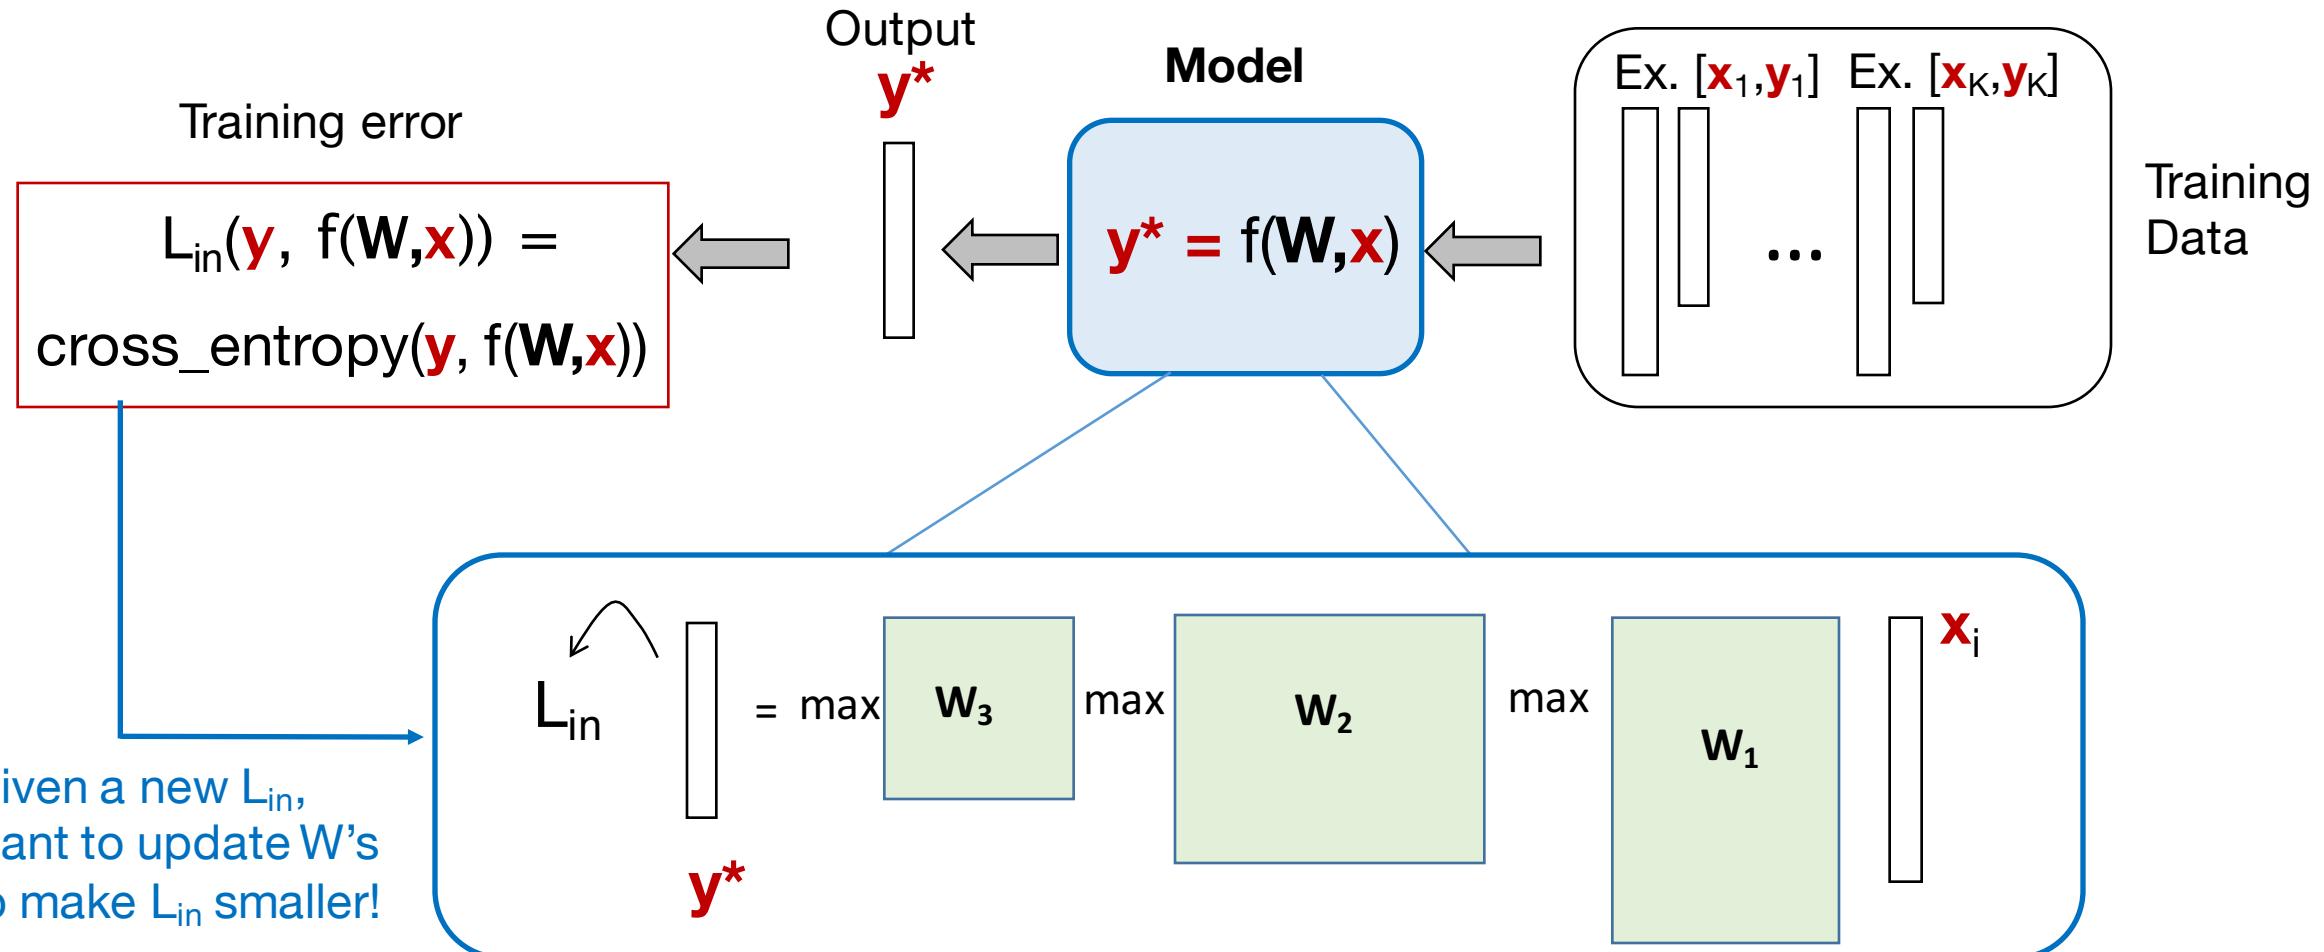
## Loss function & optimization

- Type of loss function
- Regularization
- Gradient descent method
- SGD batch and step size

**Other specifics:** Pre-processing, initialization, dropout, batch normalization, augmentation

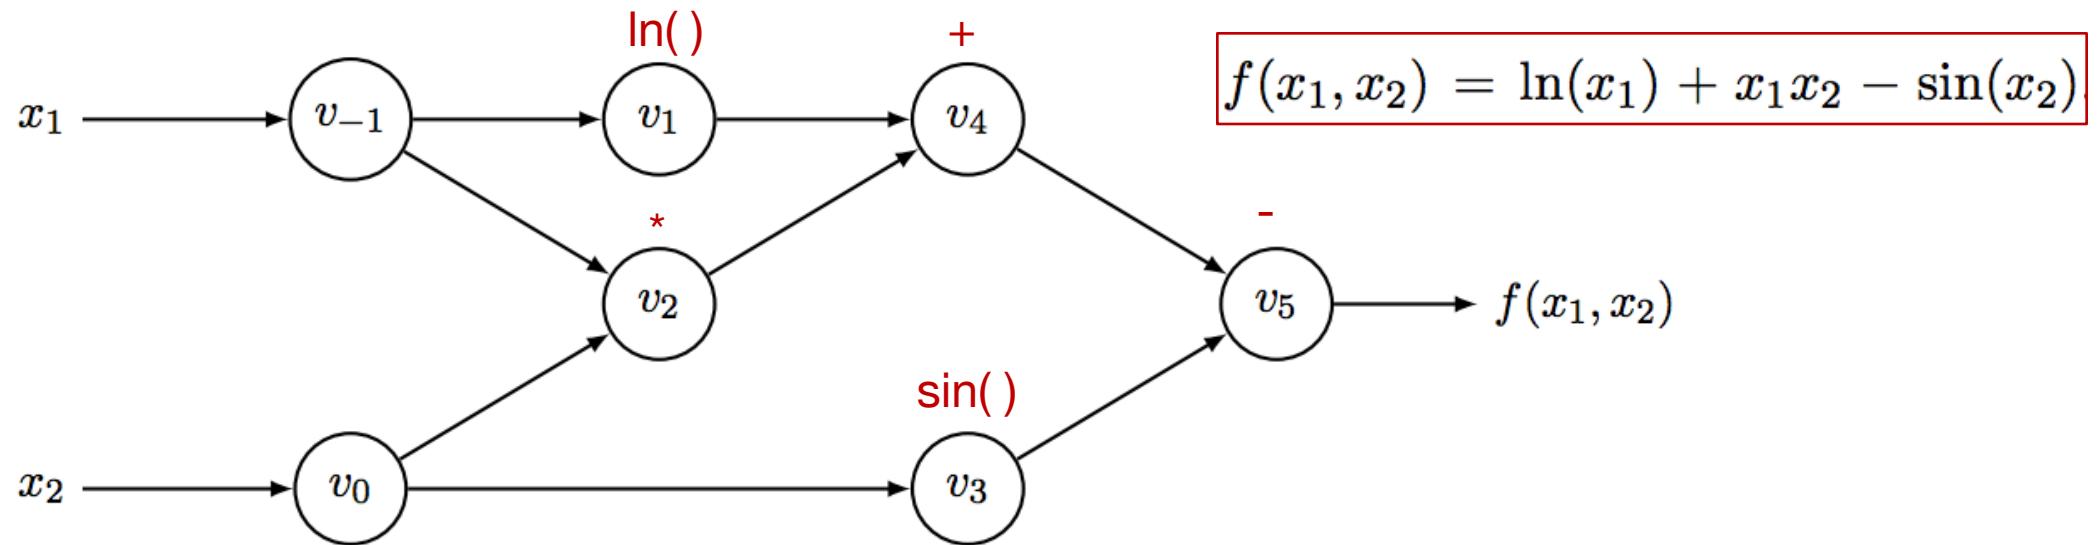


# Our very basic convolutional neural network



Backwards pass uses new  $L_{in}$  to update  $W$ 's – **backpropagation!**

# Automatic differentiation on computational graphs

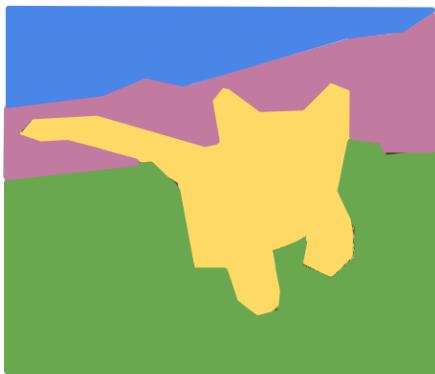


To both determine  $f$  and find  $df/dx_i$ :

- Create graph of local operations
- Compute analytic (symbolic) gradient at each node (unit) in graph
- Use inter-relationships to establish final desired gradient,  $df/dx_1$ 
  - Forward differentiation
  - Backwards differentiation = Backpropagation

# Other Computer Vision Tasks

## Semantic Segmentation



**GRASS, CAT,  
TREE, SKY**

No objects, just pixels

## Classification + Localization



**CAT**

Single Object

## Object Detection



**DOG, DOG, CAT**

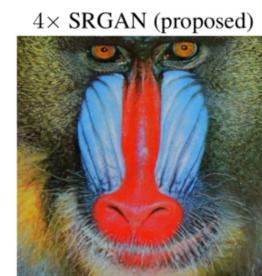
Multiple Object

## Instance Segmentation

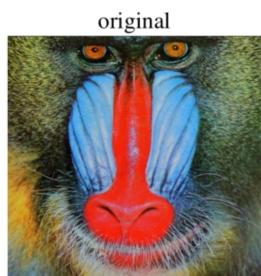


**DOG, DOG, CAT**

## Super-resolution



4× SRGAN (proposed)



original

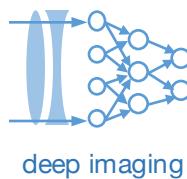
Figure 1: Super-resolved image (left) is almost indistinguishable from original (right). [4× upscaling]

[This image is CC0 public domain](#)

## Post-processing of your results: a few options at different stages

Options to examine your test data after processing:

- ROC curve, Precision-Recall
- Confusion matrix
- Sliding window visualization
- Layer visualizations
- Saliency maps etc.
- tSNE visualization



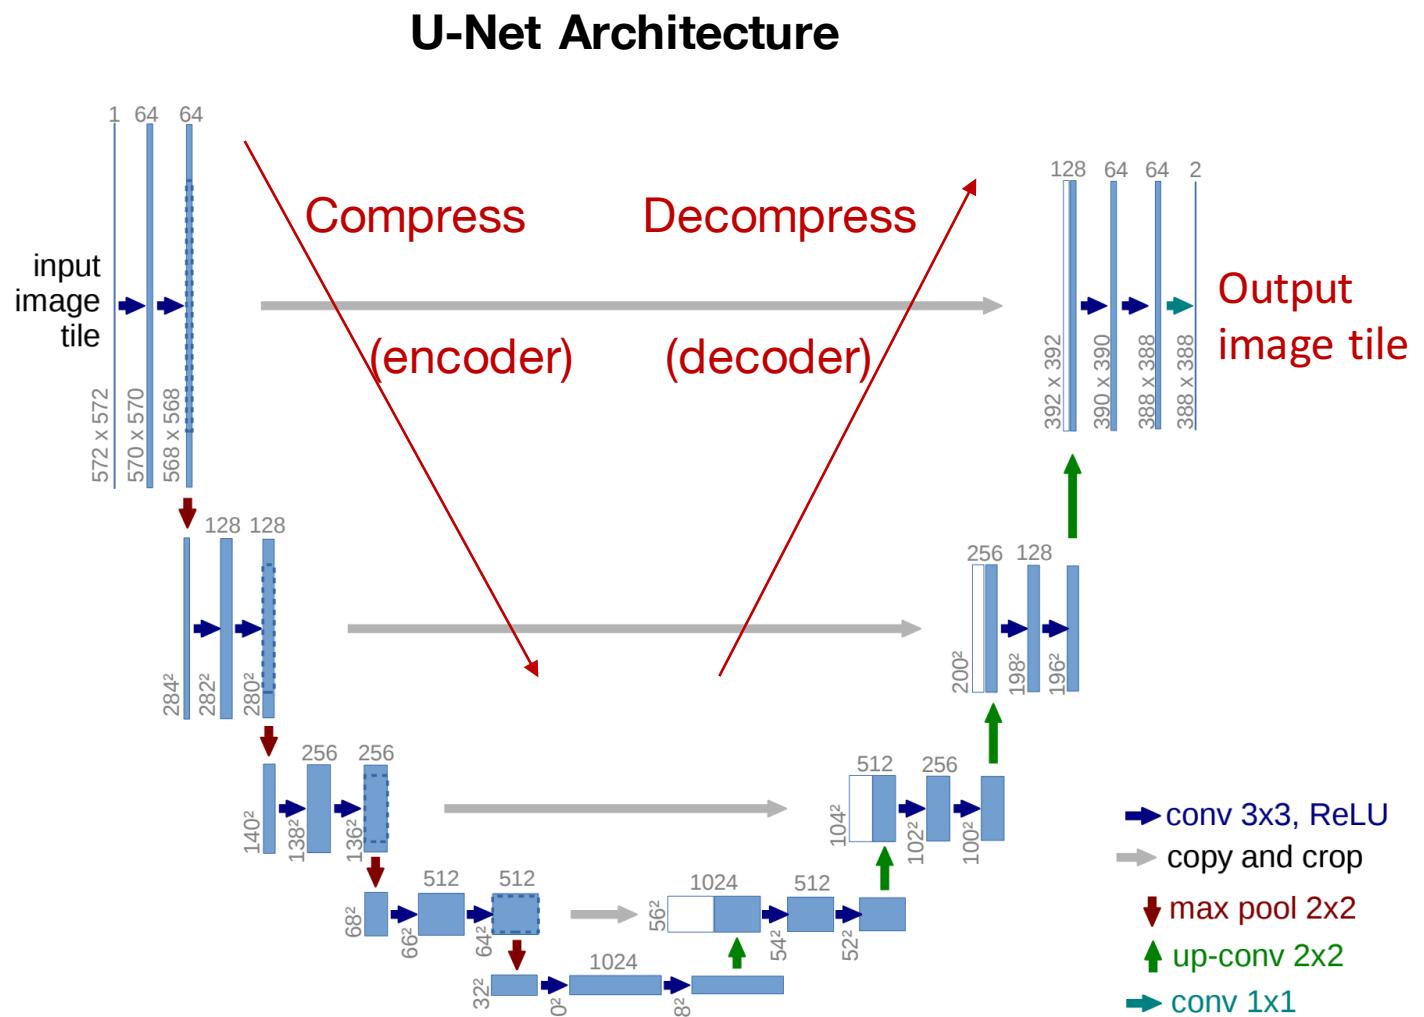
**Instead, compress x-y dimensions of input image**

- Compress spatial features into learned filters
  - Then, decompress learned filters back into same spatial dimensions
  - **Can be an autoencoder**
  - Analogous to image compression
  - A very powerful idea...

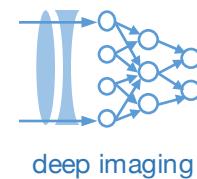
# U-Net: Convolutional Networks for Biomedical Image Segmentation

Olaf Ronneberger, Philipp Fischer, and Thomas Brox

Computer Science Department and BIOSS Centre for Biological Signalling Studies,  
University of Freiburg, Germany  
[ronneber@informatik.uni-freiburg.de](mailto:ronneber@informatik.uni-freiburg.de),  
WWW home page: <http://lmb.informatik.uni-freiburg.de/>



# Bringing together physical and digital image representations



Physical Layers

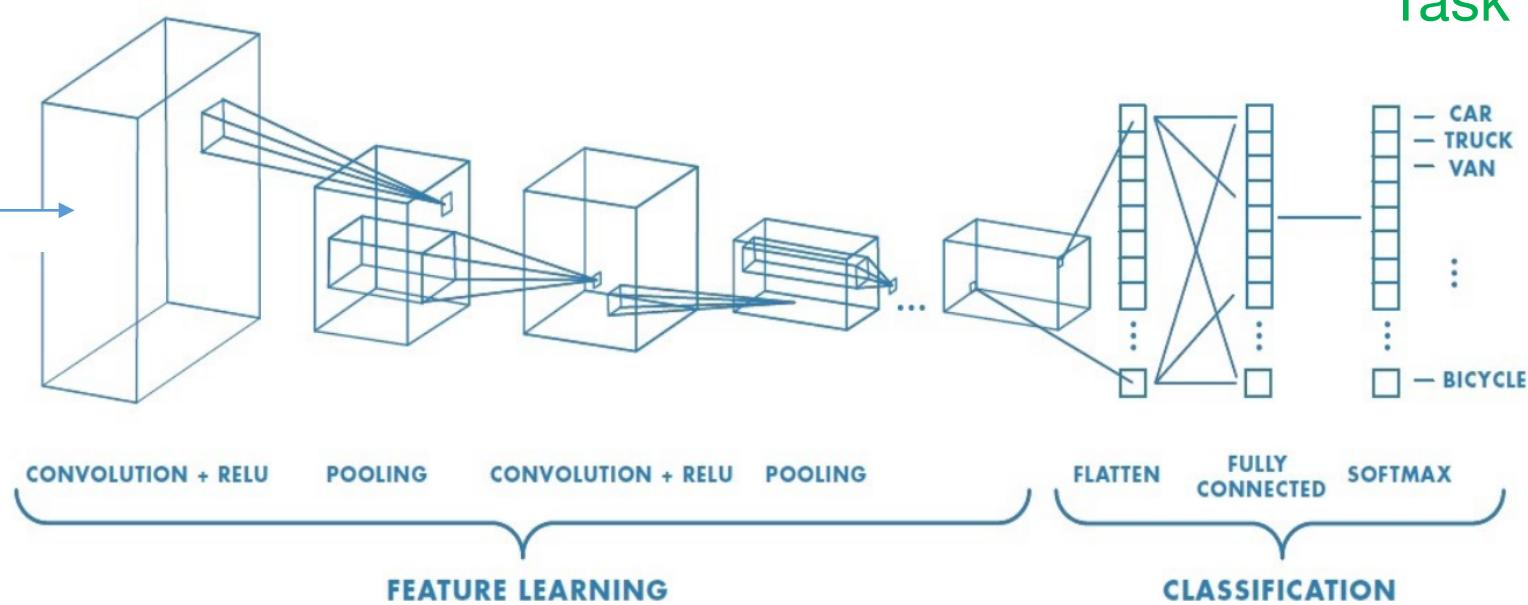
Physical Function  $I_0$

$$f[ ]$$

Digitized  $I_s$

$$I_s = f[I_0]$$

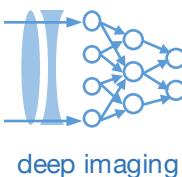
Digital Layers



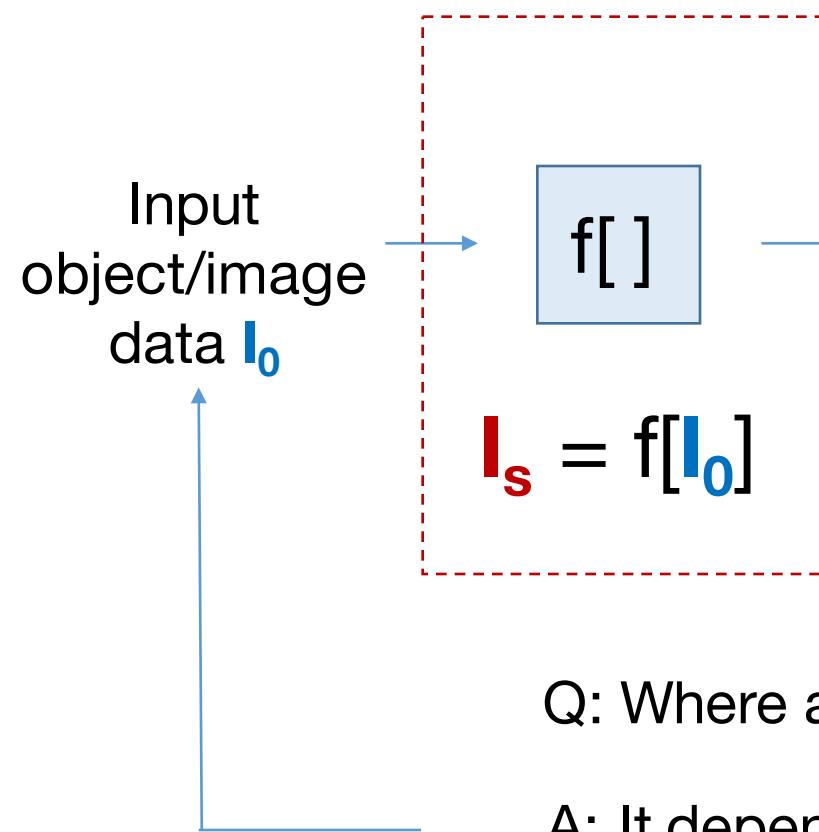
Digital layers

Physical layers

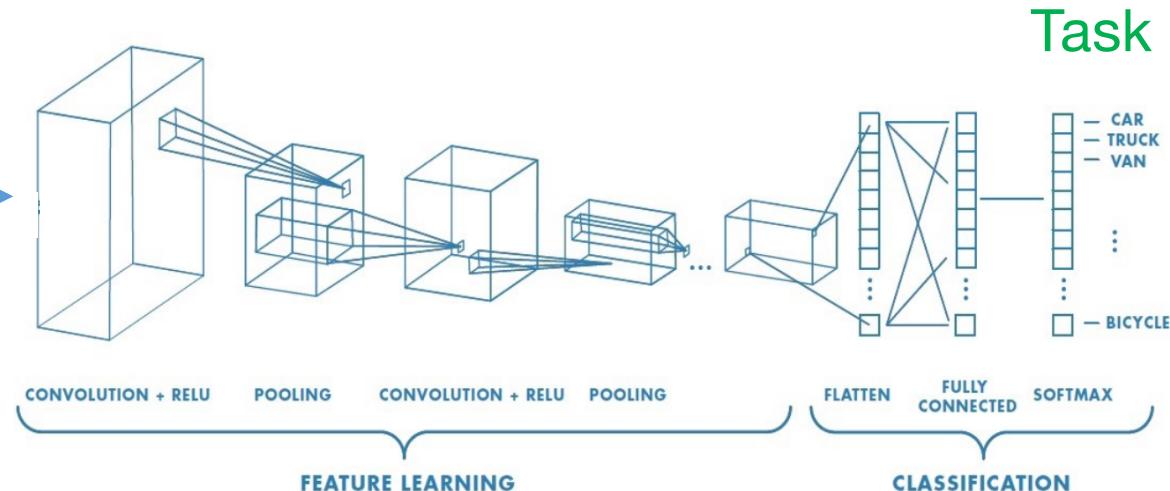
$$\text{Task} = W_n \dots \text{ReLU}[W_1 \text{ReLU}[W_0 f[I_0]] \dots]$$



## Physical Layers



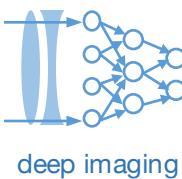
## Digital Layers



Q: Where and how should I implement my physical layer?

A: It depends on your data and implementation

- Situation #1: Fully simulated physical layers
- Situation #2: Experimentally-driven physical layers



## Situation #1: Fully simulated physical layers

Digital Layers

Physical Layers

Input object  
data  $I_0$

$f[ ]$

Digitized

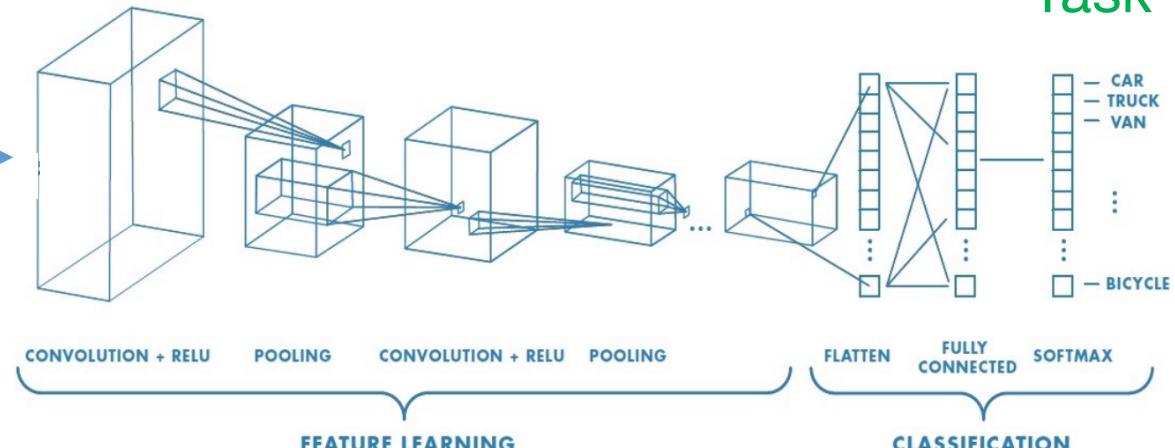
$I_s$

simulation

$$I_s = f[I_0]$$

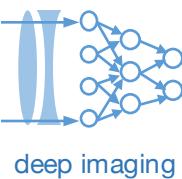
$$I_s = H B S_0$$

$$I_c = |H C S_c|^2$$



Examples:

- Simulate ultrasound samples, and learn illumination/detection geometry to best segment data
- Blur high-resolution labeled images, to learn and ideal blur shape
- Undersample your data first to learn ideal sampling strategy



Example code to achieve this is in Homework #5 and linked on course website:

jupyter physical\_layers\_example (unsaved changes) Logout

File Edit View Insert Cell Kernel Help Not Trusted Python 3

In [2]:

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
tf.__version__
```

Out[2]: '2.0.0'

In [0]:

```
mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_test, x_train = x_test/255.0, x_train/255.0 # normalization
```

In [0]:

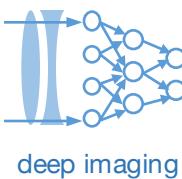
```
# add channel dimension, Tensorflow requests NHWC (or BHWC) format. Number/Batch_size, height, width, channels
# for 1D vectors, it should be Number/Batch_size, length
x_train = x_train[...,None]
x_test = x_test[..., None]
```

In [0]:

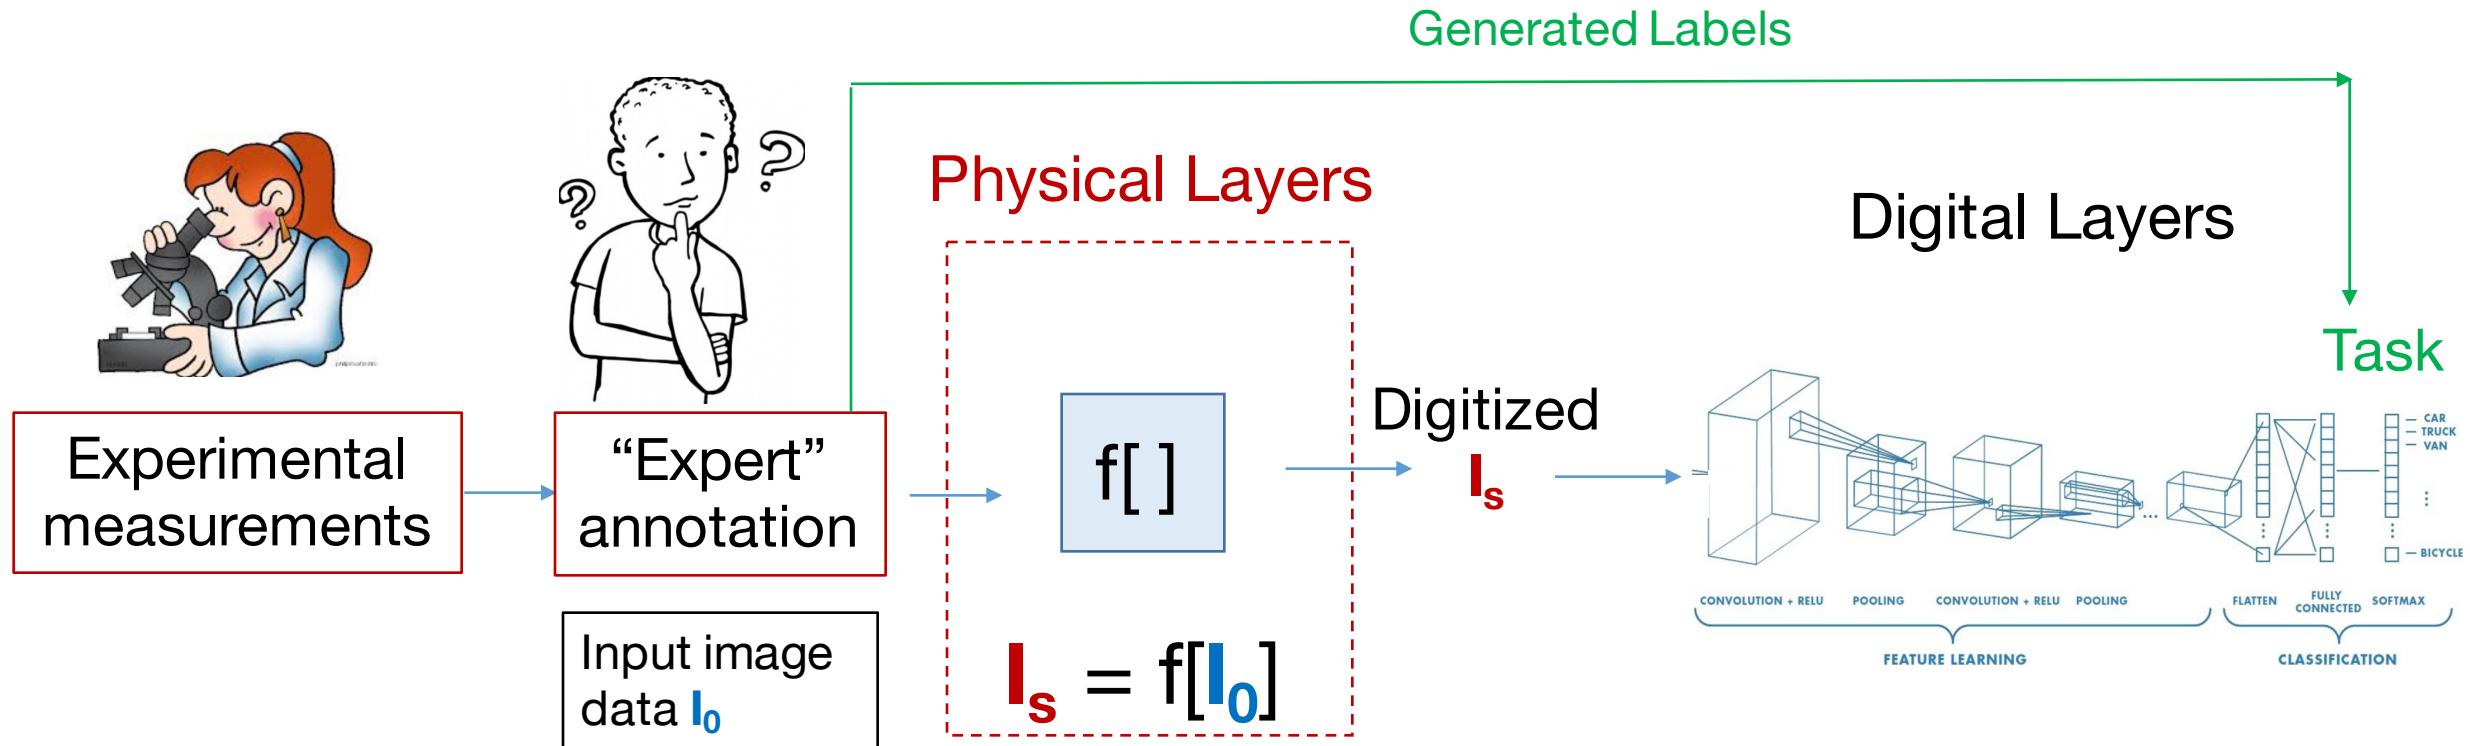
```
# another way to define models
# lets us access intermediate results too

mask_to_multiply = tf.Variable(initial_value = tf.initializers.GlorotNormal()(shape=(28,28,1)), trainable = True)
mask_to_multiply = tf.expand_dims(mask_to_multiply, 0) # add extra axis for batch size

#PHYSICAL LAYERS
input_image = tf.keras.layers.Input(shape = (28,28,1)) # Input layer, shape should be given as HWC
out = tf.math.multiply(input_image, mask_to_multiply) # hadamard product for illumination physical layer
out = tf.keras.layers.Conv2D(filters = 1, kernel_size = 5, strides =1, padding ="same", activation = "relu")(out) # con
```



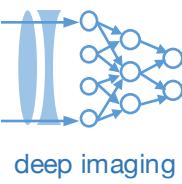
## Situation #2: Experimentally-driven physical layers



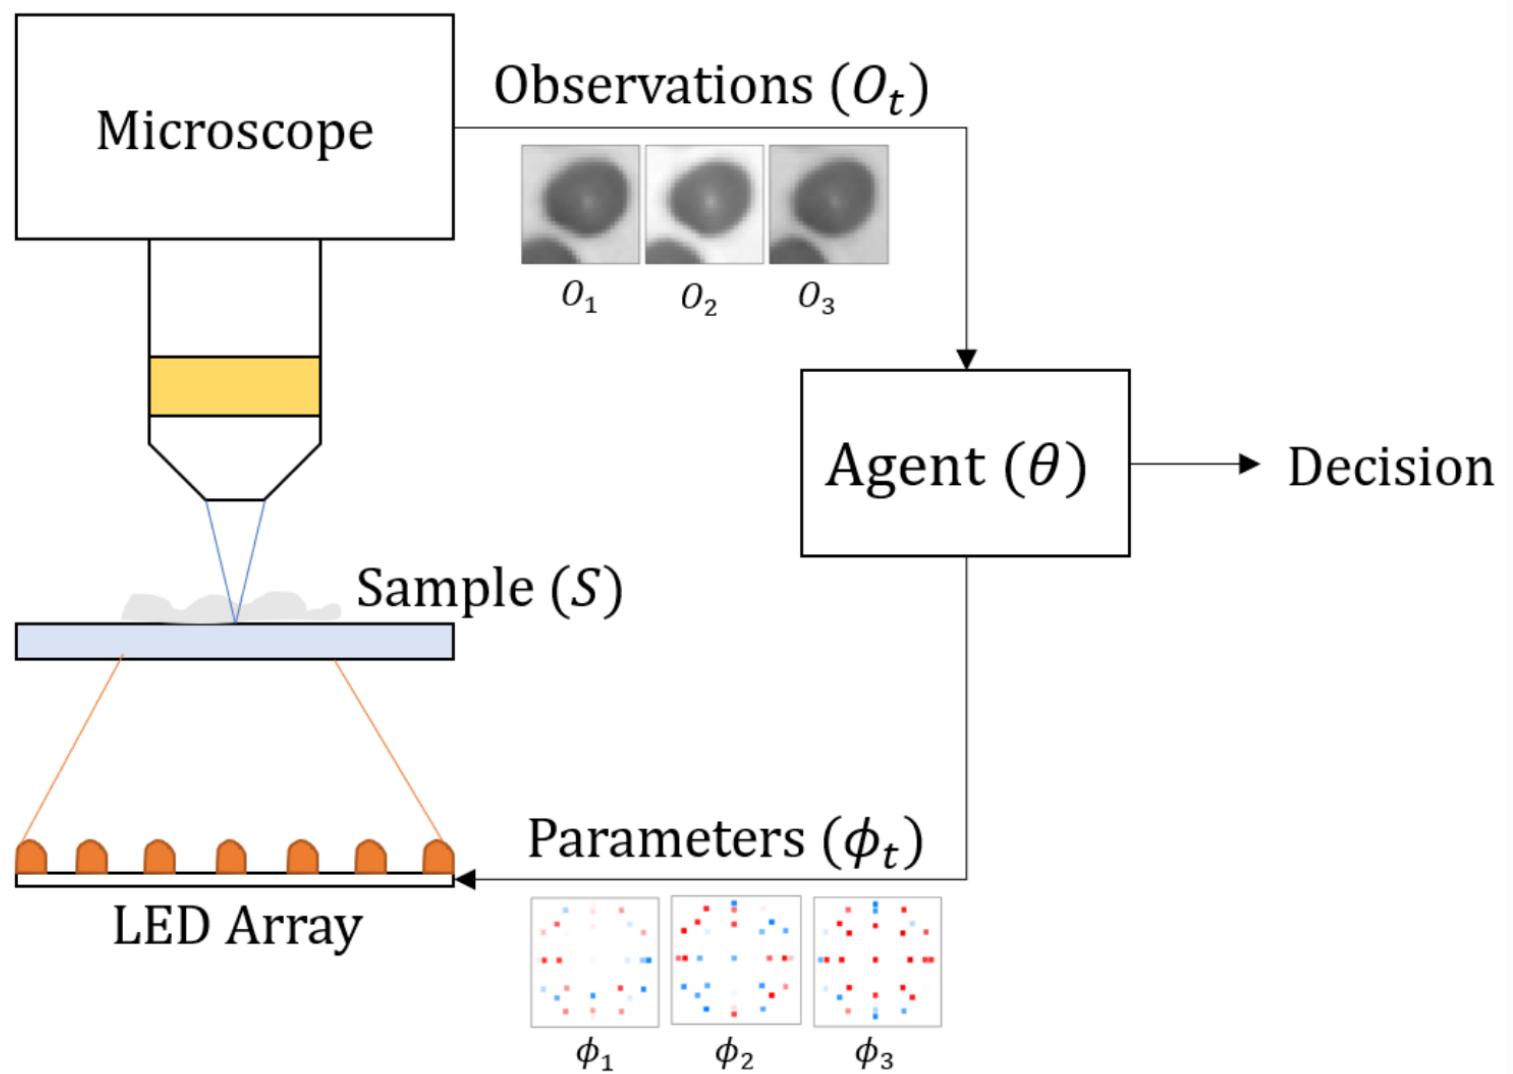
Typically need to use an “insight” to determine function  $f()$  for physical layer

Examples:

- Sum of weighted color channels equals a color filter
- weighted sum of uniquely illuminated images equals a pattern-illum. Image



## How can this be applied to optimized imaging?



## The Machine Learning in Imaging Ethics Questionnaire

Situation 4: In 10 years, you go up to a modified microscope, “the Tissue Scanner 3000”, that has a number of fancy lenses and lights. As a machine learning expert by now, you’re aware that this microscope is optimized for looking at skin lesions. It performs a scan with a particular lighting configuration and reports a score of 98% confident that the lesion is benign, allowing you to look through other examples. It asks if you’d like another scan for additional confidence or a different outcome, at which point the illumination changes and it does some more scanning and reports a 99% confidence level. You can continue with another scan, but...

Are you now comfortable with leaving the office?

Yes:

No: